

Lecture 5: Secure Channel, TLS/SSL & Miscellaneous Cryptography Topics

Part 2:

5.5 Authenticated encryption

5.6 Time-memory tradeoff for dictionary attack (Optional)

5.7 Birthday attack variant

5.8 Other interesting cryptography topics

5.9 Summary of cryptography

5.5 Authenticated Encryption


What is Authenticated Encryption

- ***Authenticated encryption:***
symmetric encryption that returns **both** ciphertext and authentication tag
- It combines **cipher** and **MAC**:
ensures message confidentiality and authenticity
- **Authenticated encryption** process: $AE(K_{AB}, M) = (C, T)$
- **Decryption** process: $AD(K_{AB}, C, T) = M$ **only if T is valid**
- Different **variants/approaches**:
 - Encrypt-and-MAC (E&M)
 - MAC-then-Encrypt
 - Encrypt-then-MAC
 - Specialized authenticated cipher

Encrypt-and-MAC (E&M)

- The sender computes the ciphertext C and tag T **separately**
- It performs **encryption**, e.g. using 2 keys K_{1AB} and K_{2AB} as follows:
 - $C = E(K_{1AB}, M)$
 - $T = \text{MAC}(K_{2AB}, M)$
- It finally sends **(C, T)**

Illustration with 1 key
(Source: Wikipedia)




- It is used in **SSH** (with a strong MAC like HMAC-SHA-256)
- Issue: T may not be random looking, and could leak information

MAC-then-Encrypt (MtE)

- The sender first computes the **tag** $T = \text{MAC}(K_{2AB}, M)$
- It then **generates** the ciphertext $C = E(K_{1AB}, M || T)$
- It finally sends C

Illustration with 1 key
(Source: Wikipedia)




- It is used in **SSL** and **TLS** up to version 1.2:
the latest TLS v1.3 uses authenticated ciphers, e.g. AES-GCM
- Issue: a decryption is still needed on a corrupted message

Encrypt-then-MAC (EtM)

- The sender first **generates** the ciphertext $C = E(K_{1AB}, M)$
- It then computes the **tag** $T = \text{MAC}(K_{2AB}, C)$
- It finally sends **(C, T)**

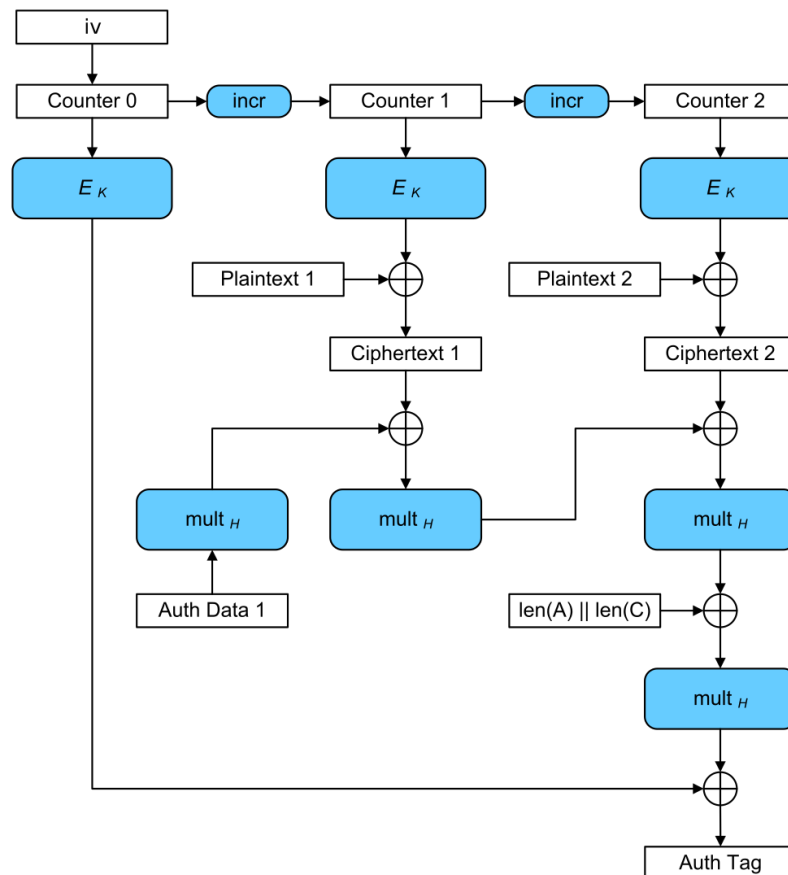
Illustration with 1 key
(Source: Wikipedia)



- It is used in **IPsec**
- Feature: a decryption is **not** performed on a corrupted message

Authenticated Cipher

- It returns an authentication tag **together with** the ciphertext
- An example is **AES-GCM** (AES in the Galois counter mode):
 - One authenticated cipher in TLS version 1.3
 - The most widely used authenticated cipher



Source: Wikipedia

5.6 Time-Memory Tradeoff for Dictionary Attack

Reference:

See the “*Precomputed hash chains*” Section of: http://en.wikipedia.org/wiki/Rainbow_table

The above Wiki page describe “Rainbow table”, which is an improved variant of time-memory tradeoff.

The original basic variant is described in the section “Precomputed hash chain”.

“Inverting” a Hash Digest in Real-Time

- Suppose a hash $H()$ is **collision resistant**: it is also one-way
- Thus, given a digest y , it is **difficult** to find a x s.t. $H(x)=y$
- Suppose we know that x is chosen from a **relatively small set** of *dictionary* D
- For illustration, assume x is a randomly & uniformly chosen **50-bit** message
- Now, even if $H()$ is one-way, given the digest y , it is still feasible to find a x in D s.t. $H(x)=y$
- One method of “inverting”: **exhaustively search** 2^{50} messages in D
- Although feasible, this would take **days** of computing time
- As the attacker, we want to **speed up** the inverting process to support a “real-time” attack

Speeding up Using a Large Table

- Supposed we are allowed to perform a **pre-processing**
- Let's view the 50 bits message as an **integer**
- One naive method is to build a dataset with 2^{50} elements:
 $(H(x), x)$ for $x = 0, 1, 2, \dots, 2^{50}-1$
and store these elements in a data structure **T** that supports a **fast lookup** (e.g., a **hash table** that facilitates a constant-time lookup)
- Now, given a digest **y**, we can **query** the data structure and readily find the associated **x**
- **Issue:** such table is too large: 2^{50} entries = 2^{10} “Tera” entries
- **Solution:** the time-memory tradeoff is a technique that “trades off” time for memory, e.g. a **lower lookup time** for a **higher storage**

Time-Memory Tradeoff (TMT)

- The main idea: use a precomputed **hash-chain**
- (Note: the term “hash-chain” appears in different context and refer to different techniques)
- Define a **reduce function** $R()$ that maps a digest y to a word w in the dictionary D
- For illustration, if D consists of all 50-bit messages, and each digest is 320 bits, then a possible reduce function simply **keeps the first 50 bits** of input:

$$R(b_0b_1\dots b_{320}) = b_0b_1\dots b_{49}$$


- Note that $R()$ is clearly **not** an inverse of $H()$
- Here is a **pre-computed hash chain**, which starts from a randomly-chosen word w_0 in D

$$w_0 \rightarrow y_0 = H(w_0) \rightarrow w_1 = R(y_0) \rightarrow y_1 = H(w_1) \rightarrow w_2 = R(y_1) \rightarrow y_2 = H(w_2)$$

E.g.:

$$\text{“hello”} \rightarrow \text{A0C0...20} \rightarrow \text{“qwert1”} \rightarrow \text{03F0...50} \rightarrow \text{“Pikachu”} \rightarrow \text{77FF...3A}$$


Hash $H()$ and Reduce $R()$: Illustration



Source: <http://kestas.kuliukas.com/RainbowTables/>

Building Pre-Computed Hash Chain Dataset


$$w_0 \rightarrow y_0 = H(w_0) \rightarrow w_1 = R(y_0) \rightarrow y_1 = H(w_1) \rightarrow w_2 = R(y_1) \rightarrow y_2 = H(w_2)$$



- The **dataset-building steps** are as follows:
 1. Select a randomly-chosen word w_0 in D
 2. Create the **hash chain** for w_0 as shown above
 3. Call w_0 the **starting-point**, and y_2 the **ending-point**
 4. Store the pair (w_0, y_2) in the data-structure T
 5. Repeat the process with other randomly-chosen starting points

Querying Pre-Computed Hash Chain Dataset (1/2)

$$w_0 \rightarrow y_0 = H(w_0) \rightarrow w_1 = R(y_0) \rightarrow y_1 = H(w_1) \rightarrow w_2 = R(y_1) \rightarrow y_2 = H(w_2)$$



- Given a digest y , first search for y in the data-structure T
- Suppose y is in T :
 - That is, y is one of the ending-points stored
 - Let's assume that w_0 is the corresponding **starting point**: hence $y = y_2$ in the above chain
 - A pre-image of y is w_2 , but we don't know w_2 at this point
 - Nevertheless, the fact that y_2 is the ending-point implies that w_2 is **within** the chain starting from w_0
 - We can construct the chain $w_0, y_0, w_1, y_1, w_2, y_2$
 - When the process hits y_2 , we have found the required w_2

Querying Pre-Computed Hash Chain Dataset (2/2)

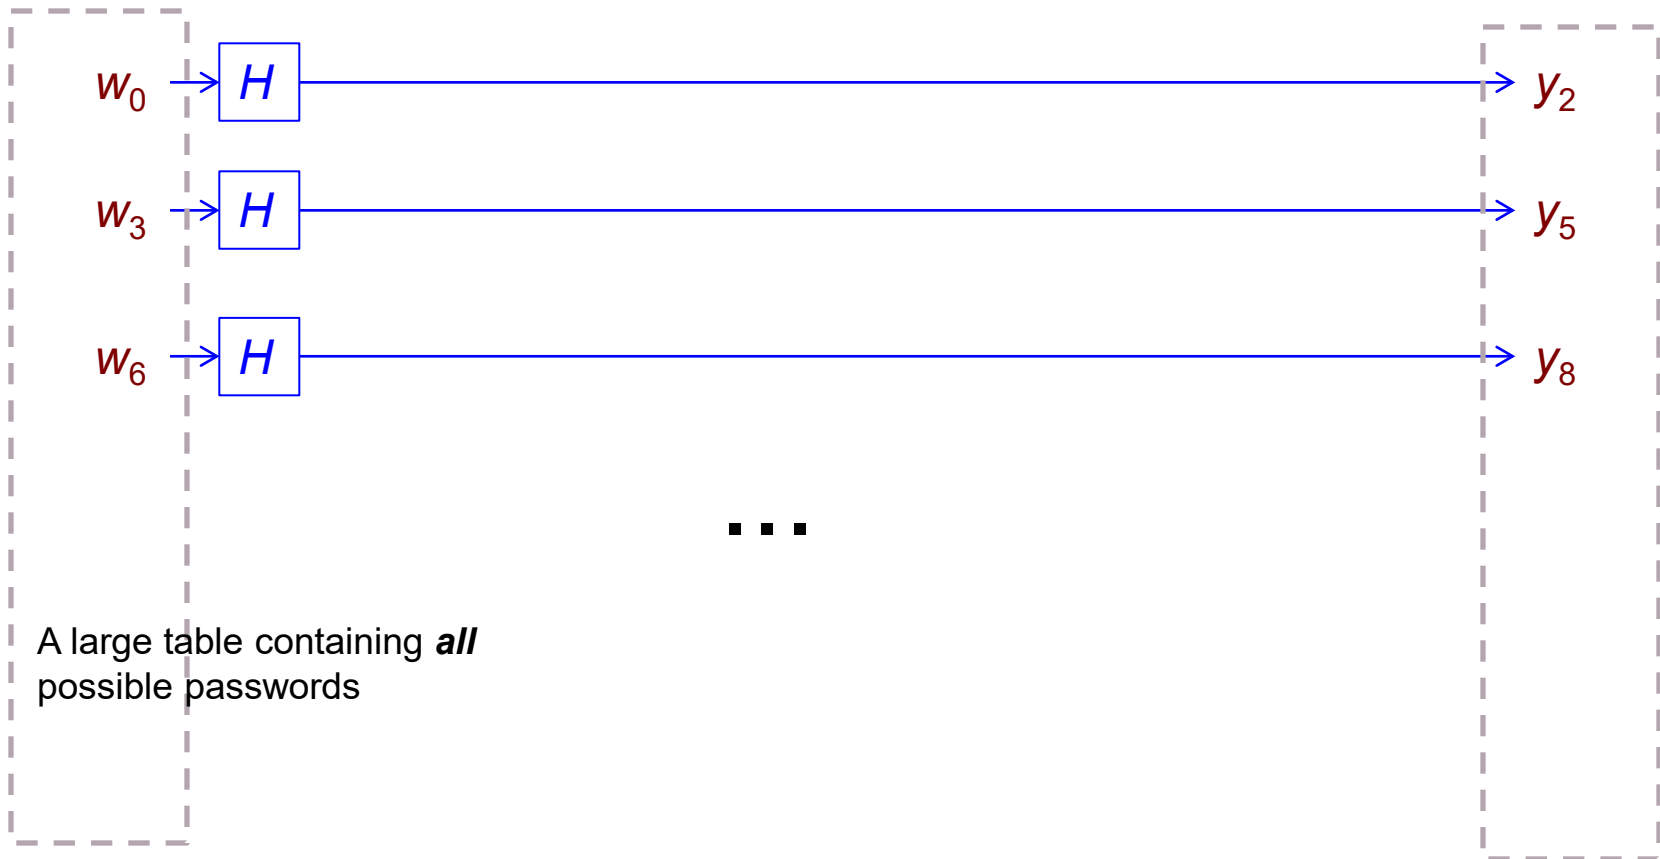
$$w_0 \rightarrow y_0 = H(w_0) \rightarrow w_1 = R(y_0) \rightarrow y_1 = H(w_1) \rightarrow w_2 = R(y_1) \rightarrow y_2 = H(w_2)$$



- Suppose y is **not** in **T** :
 - Compute $y' = H(R(y))$
 - Search the data-structure for y'
 - Suppose y' is in **T** :
 - Let's assume that the starting-point is w_0 (hence $y' = y_2$)
 - With high chances (it's not certain due to possible collisions), $y = y_1$
 - So, a pre-image of y is w_1 , i.e. $H(w_1) = y$
 - At this point, we don't know w_1
 - Constructing the chain from w_0 , and see if w_1 can be found, otherwise skip (*due to a collision issue described next*)
 - If y' is not in **T** , compute $y'' = H(R(y'))$ and repeat this query process


Without Hash Chains (As in Tutorial 1)

Given an input y , find w such that $H(w) = y$
by building a **full** lookup table



With Hash Chains (for Tutorial 1)

Given an input y , find w such that $H(w) = y$
by storing **hash-chains**




...

Analysis of Time and Space Required

- Let's **compare** the required space & time of querying stored hash chains with the naive full table method:
 - **Space:** A reduction of space **by a factor of 3** (why?)
 - **Time per query:** the number of hash operations increases by a factor of 3, but also require 2 reduce operations (why?)
 - **Accuracy:** The chains can contain repetitions (*why?*)
- A **general rule** (where k is a parameter):
we can choose the length of the chain so that the **reduction of space** is a factor of k , with the **increase of search time penalty** by a factor of k
- In our **50-bit** example:
 - The total number of entries in the full “virtual table” is 2^{50}
 - Suppose we choose $k=2^{15}$
 - The hash-chain storage is reduced to 2^{35} entries; whereas the query time increases to 2^{15} hash operations, which can still be computed in real-time

Remark: Collisions due to Hash Function (Rare/Unlikely)

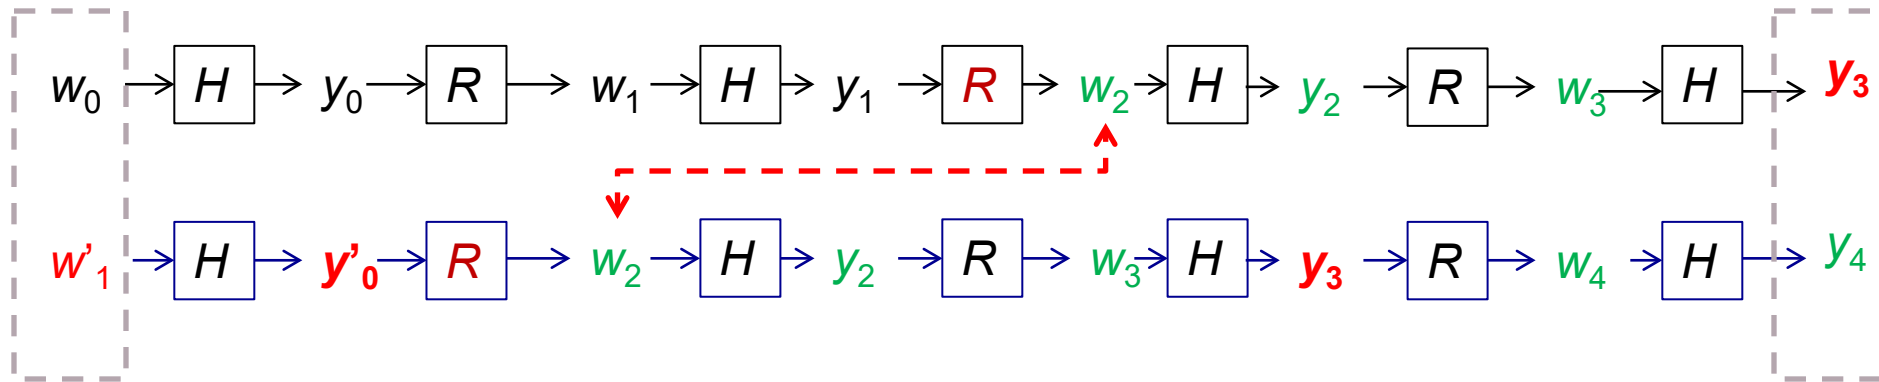
- The following **collision** is due to $H()$:



- It is extremely **unlikely**, since we assume that $H()$ is “secure”, i.e. a strong hash function
- This type of collision can thus be **omitted** in our design consideration

Remark: Collisions due to Reduce Function (Frequent/Likely)

- Collisions due to the **reduce function** may happen **frequently**, i.e. two different digests being mapped to the same word



When given y'_0 , the algorithm is unable to find w'_1 in the first chain since:


The querying algorithm initially performs these steps:

(1) lookup y'_0 , (2) lookup y_2 (3) lookup y_3 (4) search the chain starting from w_0

This causes two issues:

- Inefficiency in **storage**: part of the chain is **duplicated**, e.g. w_2 and w_3 are stored twice
- Inefficiency in **search**: it leads to searches in the **wrong chains**, before hitting the right chain, e.g. for querying y'_0 , the lookup process would transverse **both chains**

Collisions and Hash-Chain Merges: Illustrated

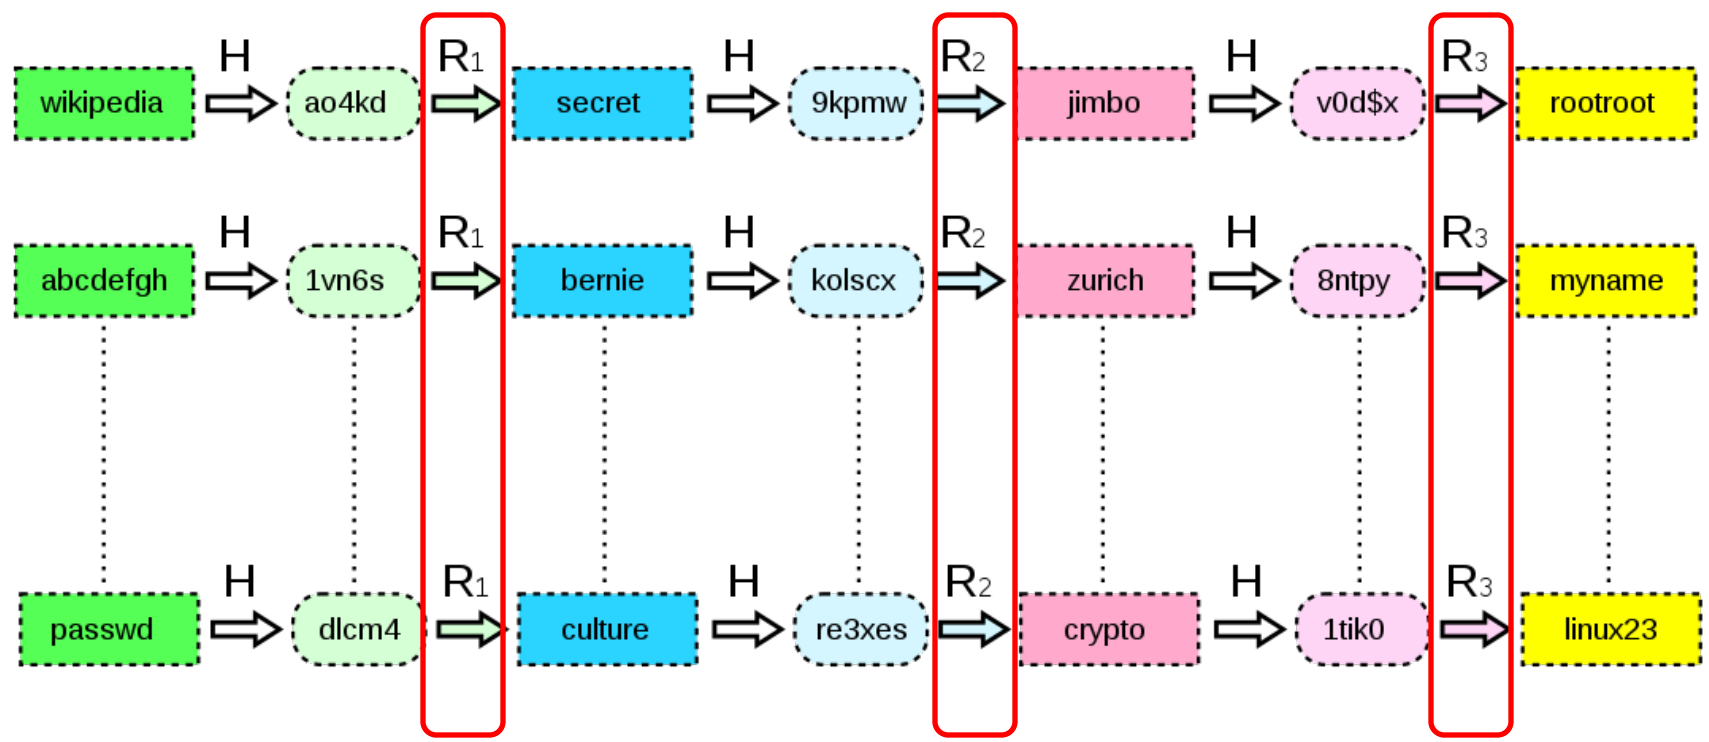


Source: <http://kestas.kuliukas.com/RainbowTables/>

Improved Variant: Rainbow Table


- **Rainbow table** gives a simple but **effective method** to address the collision issue in time-memory trade-off (the method is simple, but its analysis is quite complex)
- Rainbow table utilizes **multiple “reduce” functions**
- Details are not included in this module
- To find out more, see:
 - http://en.wikipedia.org/wiki/Rainbow_table
 - <http://kestas.kuliukas.com/RainbowTables/>
- The original research paper:
P. Oechslin, *Making a Faster Cryptanalytical Time-Memory Trade-Off*, CRYPTO 2003
<http://lasec.epfl.ch/~oechslin/publications/crypto03.pdf>

Improved Variant: Rainbow Table



Source: <https://cyberhoot.com/cybrary/rainbow-tables/>

Improved Variant: Rainbow Table




Source: Wikipedia

5.7 Birthday Attack Variant

Birthday Attack on Hash

- Suppose the digest of a hash is 80 bits: $T = 2^{80}$
- Now, an attacker wants to find a collision
- If the attacker randomly generates 2^{41} messages ($M = 2^{41}$), then $M > 1.17 T^{0.5}$
- Hence, with probability more than 0.5, among the 2^{41} messages, ***two of them give the same digest!***



In general, the probability that a collision occurs $\approx 1 - \exp(-M^2 / (2T))$

A Variant of Birthday Attack

- A variant of birthday attack is shown below
- Let S be a set of k distinct elements, where each element is a n -bit binary string
- Now, let us independently and randomly select m n -bit binary strings
- It can be shown that, the probability that at least one of the randomly chosen strings is in S is (larger than):

$$1 - 2.7^{-km2^{-n}}$$

- Notice that the set S and the set of the generated m strings are *different*
- *How can we visualize this setting and formula?*

5.8 Other Interesting Cryptography Topics

Interesting Types of Encryption

- **Format-preserving encryption:**

- A basic cipher does not care if, for instance, a plaintext is an image
- The ciphertext is **not** a viewable image
- A format-preserving encryption solves this issue: ciphertexts have the **same format** as the plaintexts
- Other possible target **plaintext types**:
IP address, ZIP code, credit card numbers

- **Fully-homomorphic encryption:**

- It enables its user to replace a ciphertext $C = E(K, M)$ with another ciphertext $C' = E(K, F(M))$, where $F()$ is as a function of M , **without ever decrypting** the initial ciphertext C
- Example: M is a text document, $F()$ is a modification of part of the text
- It's very useful for a **cloud provider**:
it doesn't know the plaintext/data, but **can change** the data as requested by the data owner (on the owner's behalf)
- It is still very **slow**: a basic operation needs an unacceptably long time!

5.9 Summary of Cryptography

Cryptography: Summary

- We have covered various cryptography topics in this module
- The **main objectives**:
 - Learn how cryptographic schemes and primitives work
 - Learn how to use them correctly
 - Learn how to reason about their security
- What cryptography **provides**?
 - It provides many useful primitives
 - It serves as the basis for many security mechanisms
- However, cryptography:
 - Is **not** the solution to all security problems:
software vulnerabilities, social engineering attacks, etc.
 - Needs to be implemented and deployed **securely/properly**
 - Is not something you should invent/design yourself

Importance of Crypto?



Source: Wikipedia