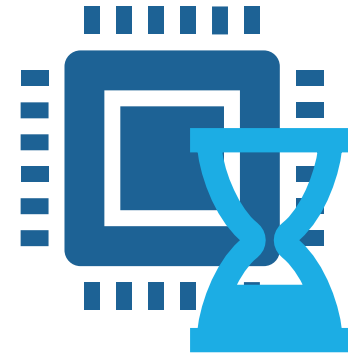


# CS3210

## Parallel Computing

Changes from Monday in Green



## Lab 2

Mon (4pm)

Tues (2pm)

# Admin Updates

---

- **Assignment 1 Part 1 due 11am next Monday**
  - **Late penalty: 10% per day, up to a week**
  - Consider updating your program, measurements and report after this lab
  - **See FAQ for questions about compiler choice and flags**
  - New installed packages: **clang**
  - Request new packages/use of non-standard libraries: drop teaching team a message
- Lab 1 comments to release sometime later this week
  - Will post announcement; look in LumiNUS gradebook

# Admin

## Lab Machine Info

---

- Hostnames are: **soctf-pdc-`<node ID>`**
  - Available: 24/7, at least until end of Assignment 1

Node IDs	001 - 008	009 - 016	018 - 019	020 - 021
CPU	Xeon Silver 4114	Intel Core i7-7700	Xeon Silver 4114	Intel Core i9-9700
# sockets	1	1	2	1
# cores/skt	10	4	10	8
SMT?	Yes	Yes	Yes	No
RAM (GB)	32	32	64 (NUMA)	32
SWAP (GB)	2	2	~9	~14

# Admin

# Roadmap

---

- Please use only the machines assigned to your lab pair
  - Rotate between the machines with your lab partner
- Check if someone else is on the same machine!
  - Outside of this: free-for-all, but please be considerate
- Today's lab
  - Part 1: Introduction to OpenMP programming
  - Part 2: Introduction to Performance Instrumentation

Admin

# Lab 2 Submission Instructions

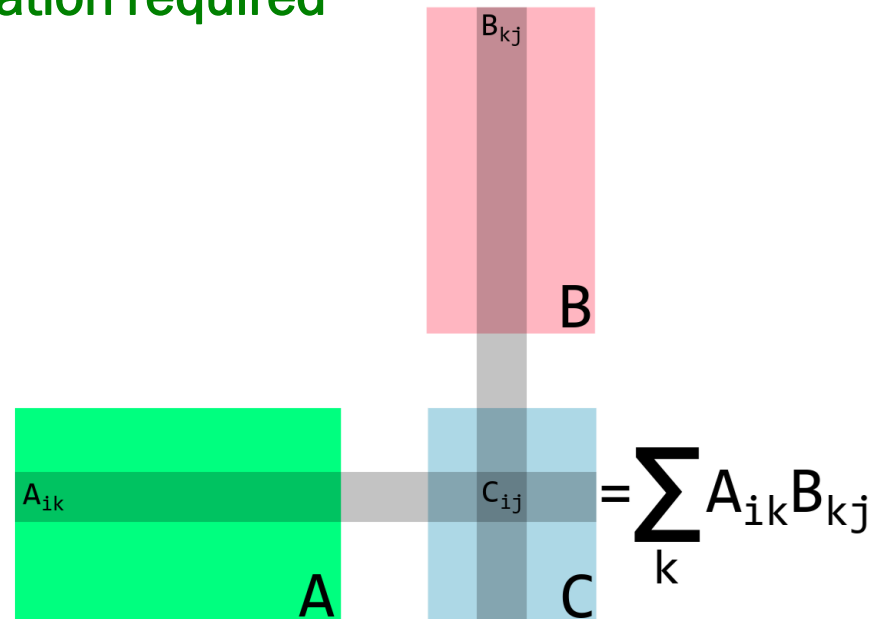
---

- **Due next Tuesday (15 Sep), 11.59pm**
  - Each student submits their own PDF writeup (indep. work)
  - Name it A0123456Z.pdf
  - No need to include your programs in writeup
- Contents
  - Your responses to ex4 to ex6 (**max 1 paragraph each in addition to any calculations performed**)
  - Raw measurements for ex6 with hostnames
  - Graph of results with justifications for observations

## Part 1

# Matrix Multiplication (1)

- Classical example of an embarrassingly parallel problem
  - Trivial to parallelise - what kind of parallelism (task or data)?
  - **Why is it trivial to parallelise?** No synchronization or coordination required

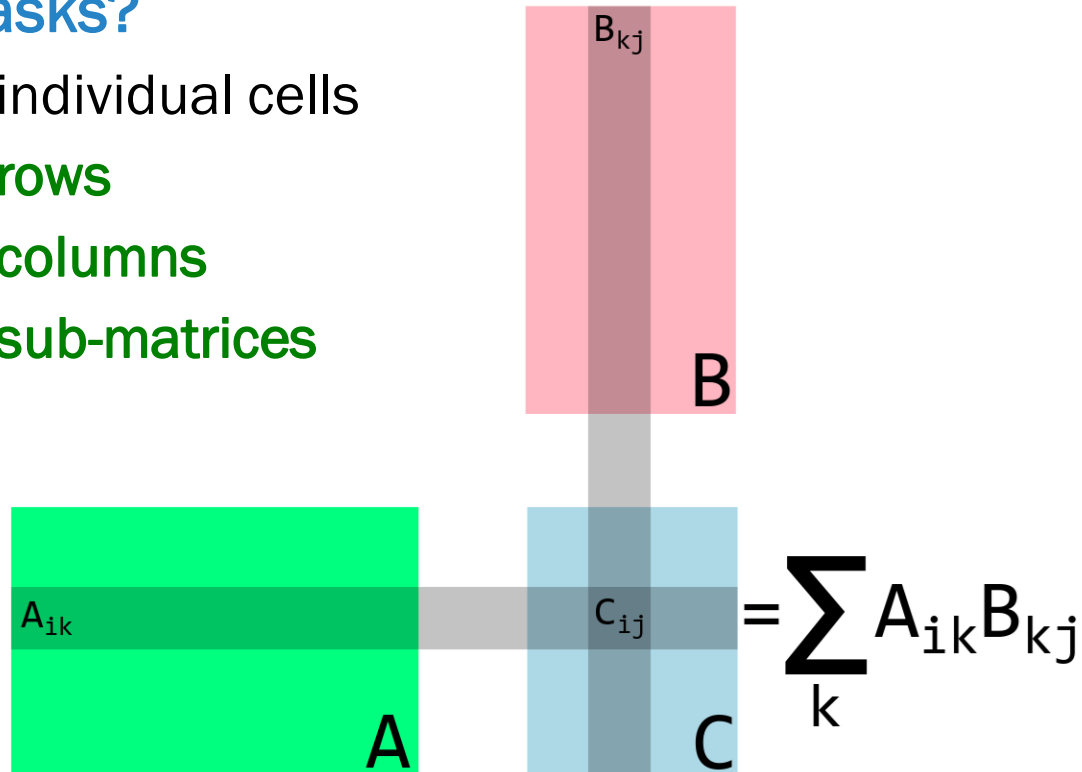


## Part 1

# Matrix Multiplication (2)

- In what ways can we decompose the computation of  $C$  into tasks?

- By individual cells
- By rows
- By columns
- By sub-matrices

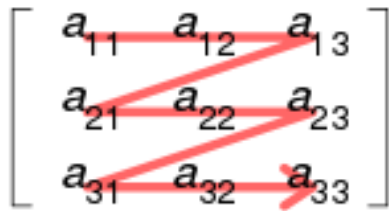


# Part 1

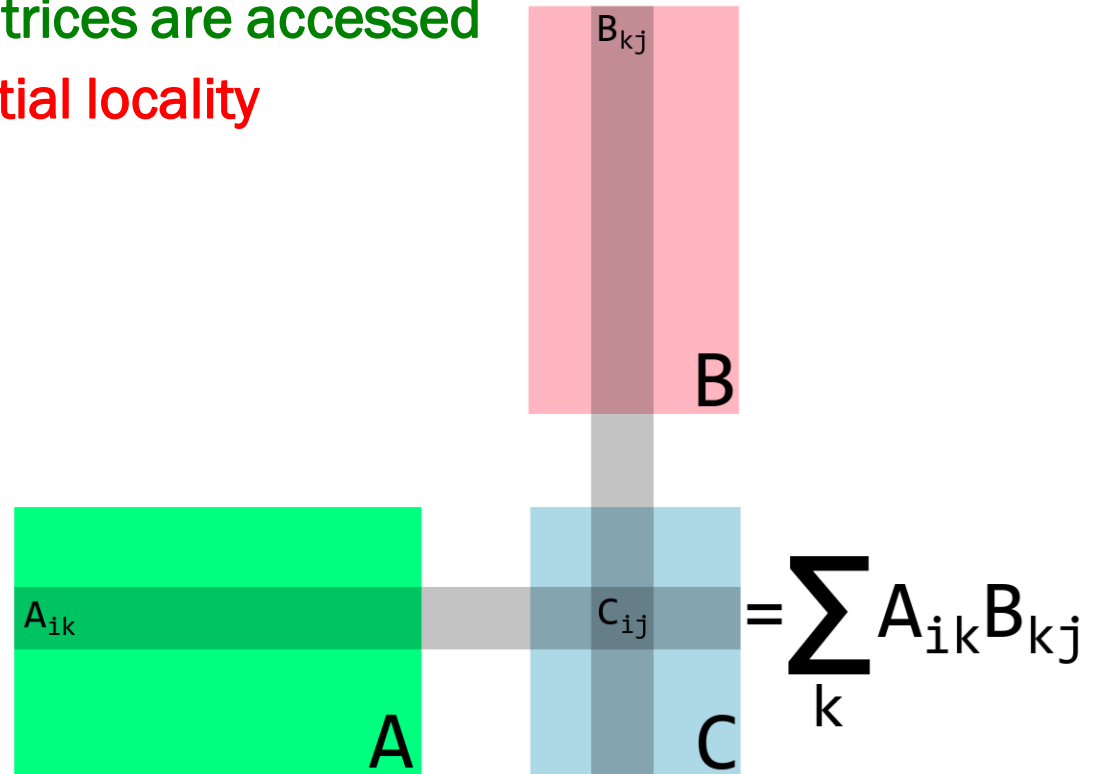
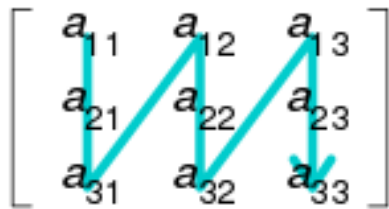
## Matrix Multiplication (3)

- How can we layout the matrices in memory?
  - Different: how matrices are accessed
  - Key principle: spatial locality

Row-major order



Column-major order





# Part 1

## OpenMP (1)

---

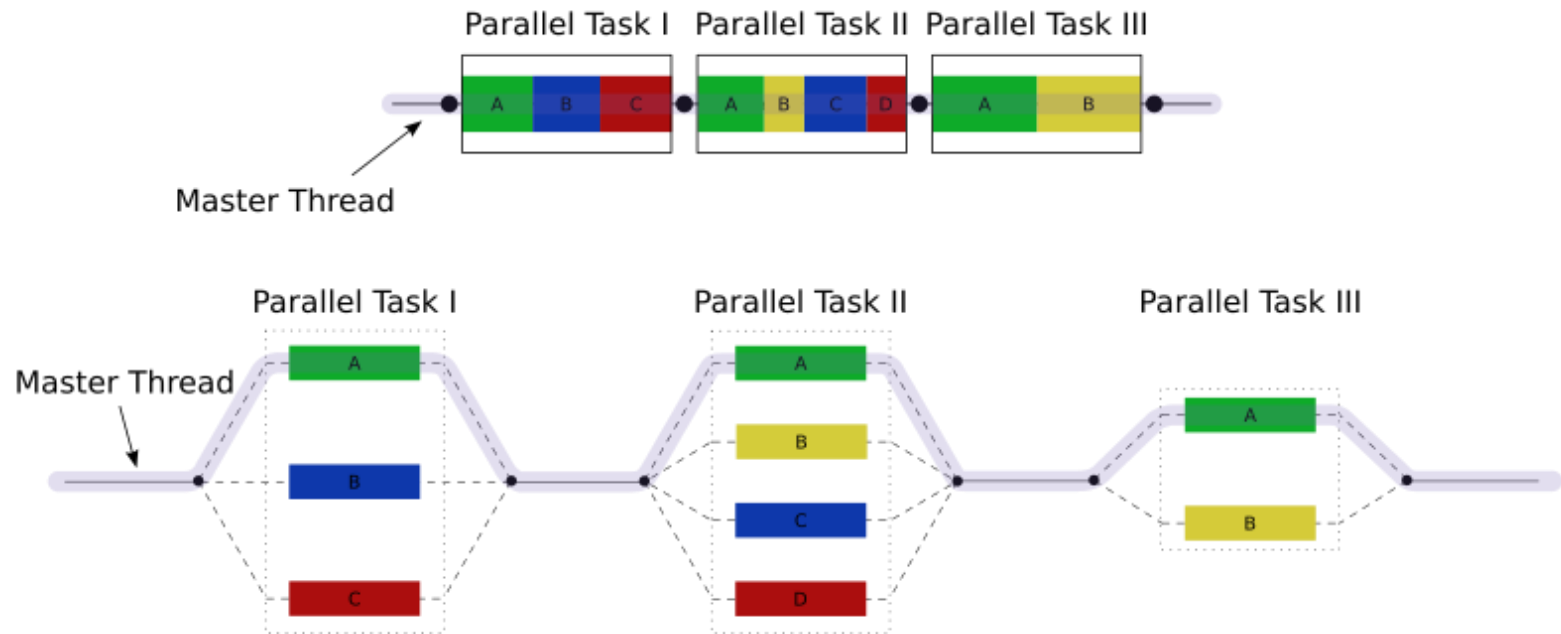


- OpenMP: Open Multi-Processing
  - Set of compiler directives and library routines to specify high-level parallelism in C, C++ and Fortran
  - **Shared-memory programming with threads**
  - Supports both task and data parallelism
  - C header: **#include <omp.h>**
  - Compilation flag: **-fopenmp**
- Directive structure
  - **#pragma omp <directive> [<clause>...] [{ block }]**

# Part 1

## OpenMP (2)

- Structure of OpenMP programs
  - Interleaved sequential sections and parallel regions
  - Each parallel region: one master and multiple worker threads



## Part 1

# OpenMP Directives (1)

---

- Parallel region directive: **parallel**
  - Demarcates a parallel region of code
  - Current thread becomes master thread (with thread ID **0**) at start of this block and forks additional worker threads
  - **Default number of worker threads forked is 1 less than the number of logical cores on that machine**
  - Can change this with **omp\_set\_num\_threads(int)** or by setting the environment variable **OMP\_NUM\_THREADS**

```
#pragma omp parallel
{
    structured block
}
```

## Part 1

# OpenMP Directives (2)

---

- Loop distribution directive: **for**
  - Can only be specified inside a parallel region
  - Can be combined as **#pragma omp parallel for**
  - **Divides iterations of a for loop amongst existing threads**
  - Additional clauses allow you to specify order and number of iterations assigned to each thread, e.g. **schedule** clause

```
#pragma omp parallel
{
    #pragma omp for [clause...]
    for (...; ...; ...) {
        loop body
    }
}
```

# Part 1

## OpenMP Directives (3)

---

- Loop distribution directive: **for**
  - Four types of scheduling: **static**, **dynamic**, **guided**, **runtime**
  - If the **schedule** clause is not specified - default is implementation-defined

### STATIC



### DYNAMIC



### GUIDED A



### GUIDED B



# Part 1

## OpenMP Directives (4)

---

- Loop distribution directive: **for**
  - What is the expected output of this section of OpenMP code compiled with GCC and run on a 4-core machine?

```
#pragma omp parallel for
for (int i = 0; i < 12; i++) {
    printf("Thread %d - iteration %d", omp_get_thread_num(), i);
}
```

- **Should see: each thread assigned a contiguous “chunk” of 3 loop iterations**
- **Are the messages from the same thread printed in order? Yes**
- **Are the messages between threads printed in order? No**

## Part 1

# OpenMP Directives (5)

---

- Task distribution directive: **sections**
  - Can only be specified inside a parallel region
  - Specify code sections with **#pragma omp section**; each a task that will be assigned to available threads, one at a time

```
#pragma omp parallel {  
    #pragma omp sections {  
        #pragma omp section {  
            structured block  
        }  
  
        #pragma omp section {  
            structured block  
        }  
    }  
}
```

## Part 1

# OpenMP Directives (6)

---

- Barrier synchronization directive: **barrier**
  - Some directives (e.g. **for**) have an implied barrier at the end
  - Blocks threads that arrive early until all threads reach the barrier

```
#pragma omp barrier
```

- Atomic operation directive: **atomic**
  - Specifies a memory location must be updated atomically

```
#pragma omp atomic  
  
statement expression
```



## Part 1

# OpenMP Directives (7)

---

- Master section directive: **master**
  - Specifies a region that must only be executed by the master thread

```
#pragma omp master
{
    structured block
}
```

- Critical section directive: **critical**
  - Specifies a critical section

```
#pragma omp critical
{
    structured block
}
```

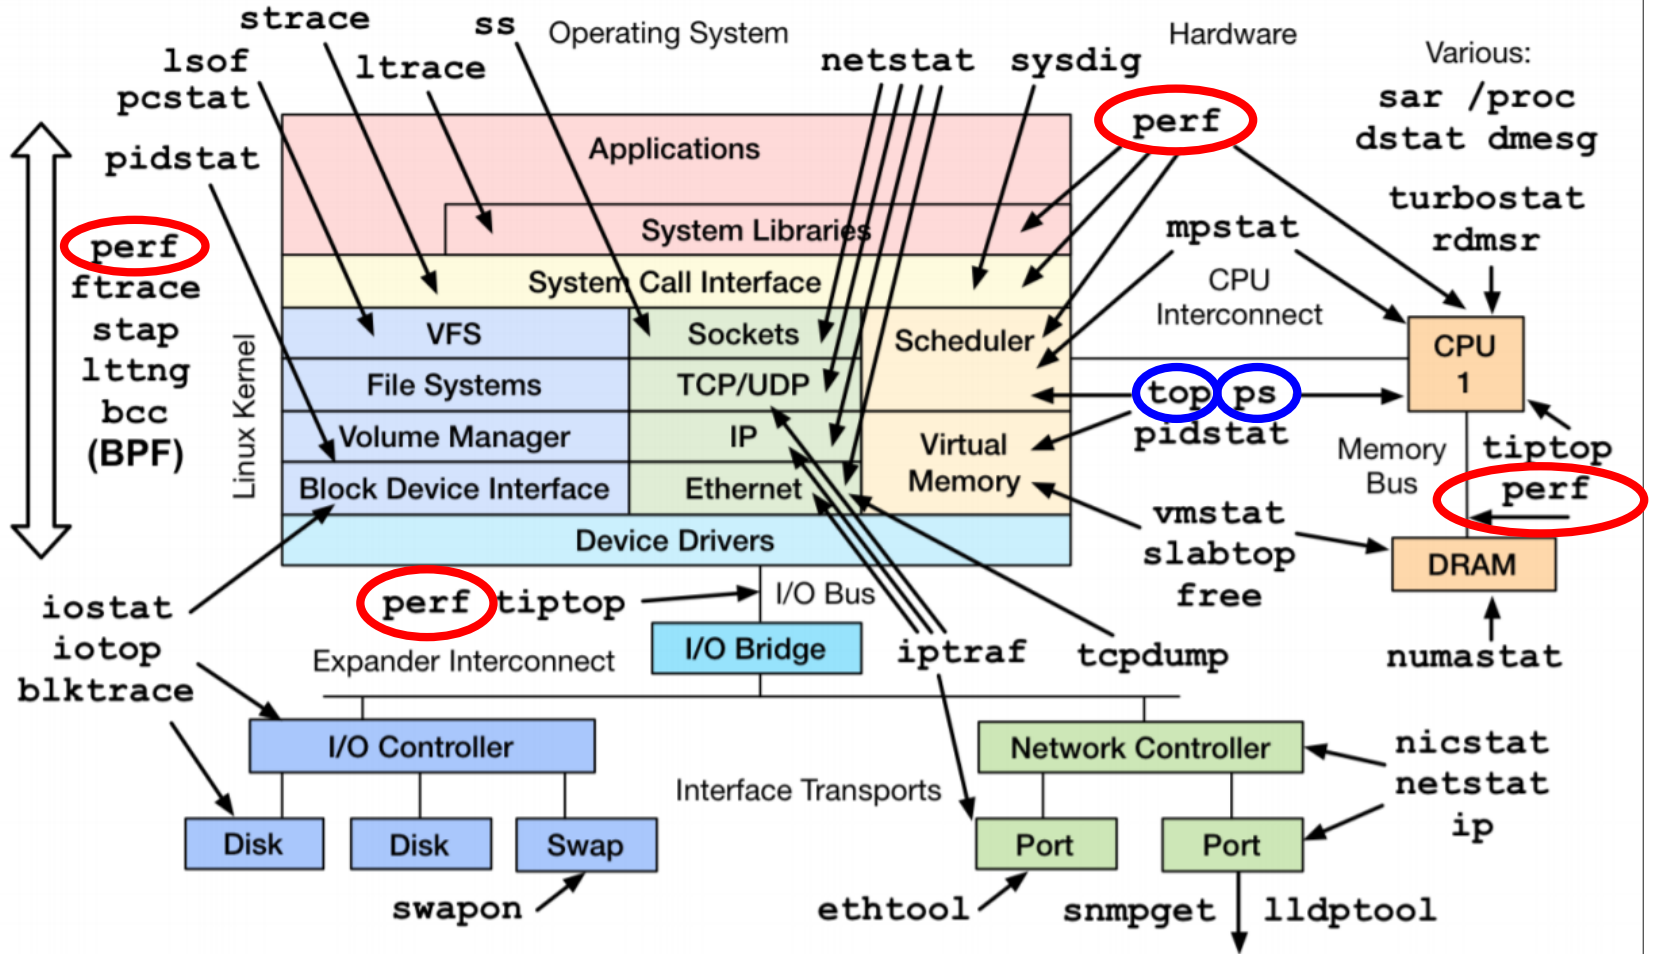
## Part 2

# perf Utility

---

- The **perf** utility is one of the most popular performance profiling tools on Linux
  - User space utility; not available by default
  - Ubuntu: install **linux-tools-common** package with the **apt** package manager with **sudo apt install <package>**
    - You may be prompted to install other required packages - do so
  - Profiles entire system during execution of a shell command
  - Part of a kernel-based subsystem that exposes performance counters in hardware and software and tracepoints (e.g. for interrupts, scheduling, etc.)
  - Command: **perf stat <command> <args>...**

# Linux Performance Observability Tools



## Part 2

# perf Utility Output

- Sample output: **perf stat <program> <args>...**

```
ubuntu-16-04@ubuntu1604-Aspire-4752:~/Desktop$ sudo perf stat ./beh 100000000
[5681]: Step 0
[5682]: Step 0
[5681]: Step 1
[5682]: Step 1
[5681]: Step 2
[5682]: Step 2
[5681]: Step 3
[5682]: Step 3
[5681]: Step 4
[5682]: Step 4
[5682] Child Done!
[5681] Parent Done!

Performance counter stats for './beh 100000000':

 3238.831006 task-clock (msec)    #    1.994 CPUs utilized
           6 context-switches    #    0.002 K/sec
           0 cpu-migrations      #    0.000 K/sec
          87 page-faults         #    0.027 K/sec
 7.077,783.766 cycles             #    2.185 GHz                   (83.21%)
 6,071,779,673 stalled-cycles-frontend #   85.79% frontend cycles idle (83.31%)
 17,203,110 stalled-cycles-backend  #    0.24% backend cycles idle  (66.76%)
 4,004,748,213 instructions       #    0.57 insn per cycle        (83.43%)
                               #   1.52 stalled cycles per insn (83.45%)
 999,786,370 branches             # 308.687 M/sec                 (83.28%)
   26,545 branch-misses          #    0.00% of all branches
 1.624490259 seconds time elapsed
```

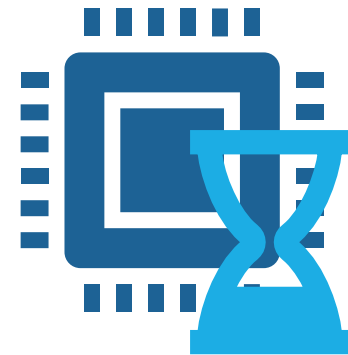
Observe the  
number of  
context-  
switches

Do you think  
this is real,  
kernel or  
user time?

# CS3210

## Parallel Computing

Thank you! Any questions?



### Lab 2

Mon (4pm)

Tues (2pm)

**[bit.ly/cs3210-t04-qn](https://bit.ly/cs3210-t04-qn)**