

CS3210

Parallel Computing

Changes from Monday in Green



Tut 2

Mon (4pm)

Tues (2pm)

Admin Updates

- **Assignment 1 part 1 due yesterday morning, 11am**
 - Late penalty: 10% per day, up to a week
 - If you have submitted, download your submission and check its contents are correct
- Lab 1 comments released in LumiNUS gradebook
 - Full credit awarded if you followed submission instructions
 - For clarifications: email me or drop me a message
- End-tutorial Quiz 2 later

Q1

Task Parallelism

- Parbegin-parend pattern
 - Demarcates a parallel region in code
 - A set of threads are spawned at **parbegin**
 - Threads execute subsequent statements until **parend**, where they are joined/destroyed
 - Statements only executed in parallel with another statement if specified with **parallel** keyword, e.g. **A parallel B**
 - A sub-region is demarcated with **do (...) end**

```

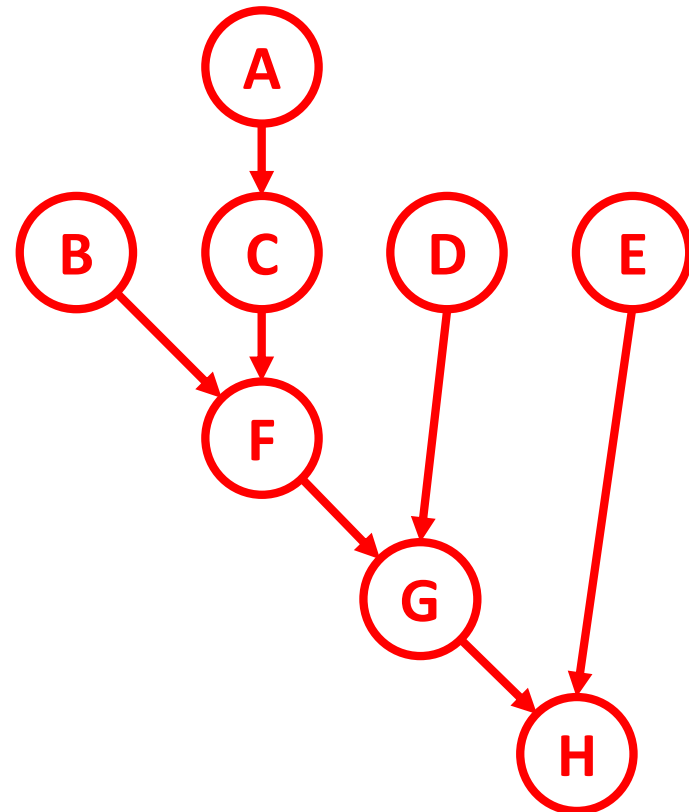
parbegin
do
  parbegin
    do
      parbegin
        do
          A
          C
        end
        parallel
        B
      parend
      F
    end
    parallel
    D
  parend
  G
end
parallel
E
parend
H

```

Q1(a)

Task Parallelism

- Sketch out the task dependence graph for code fragment 1

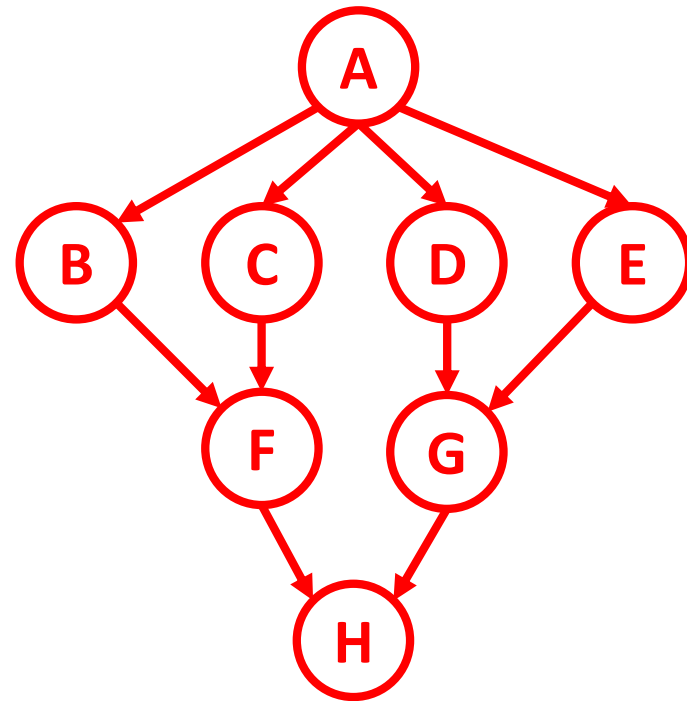


```
A
parbegin
do
  parbegin
    B
    parallel
    C
  parend
  F
end
parallel
do
  parbegin
    D
    parallel
    E
  parend
  G
end
parend
H
```

Q1(a)

Task Parallelism

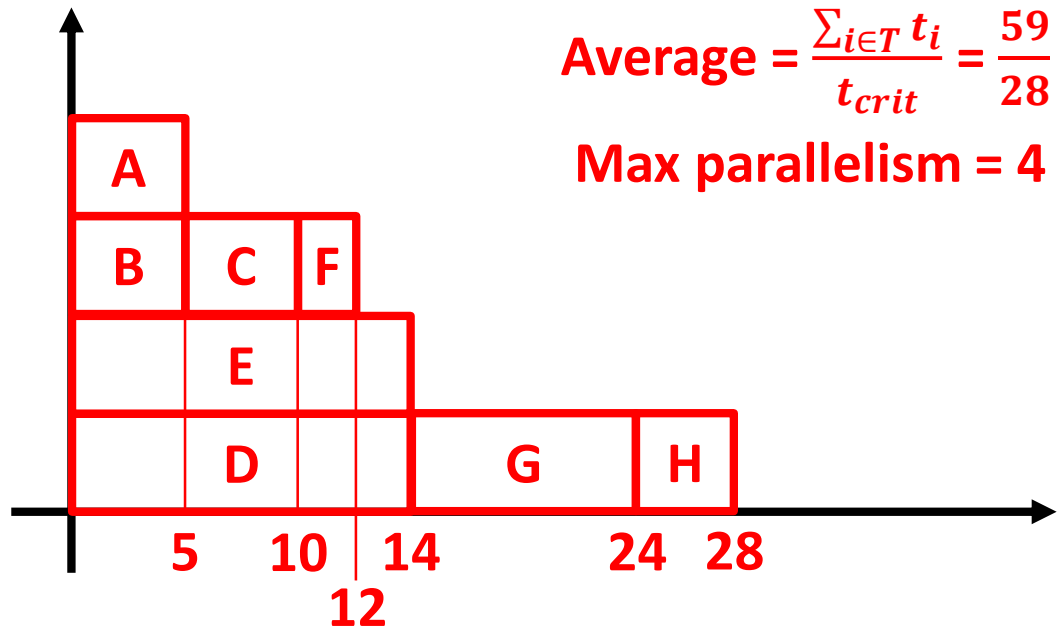
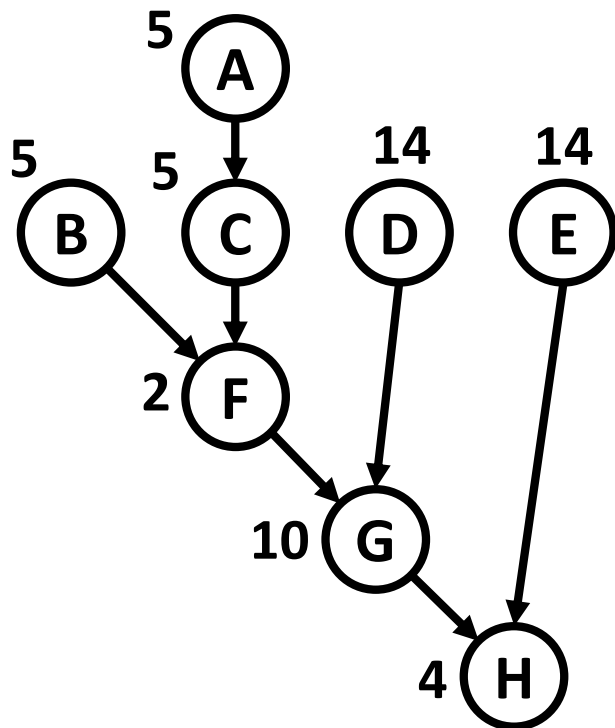
- Sketch out the task dependence graph for code fragment 2



Q1(b)

Task Parallelism

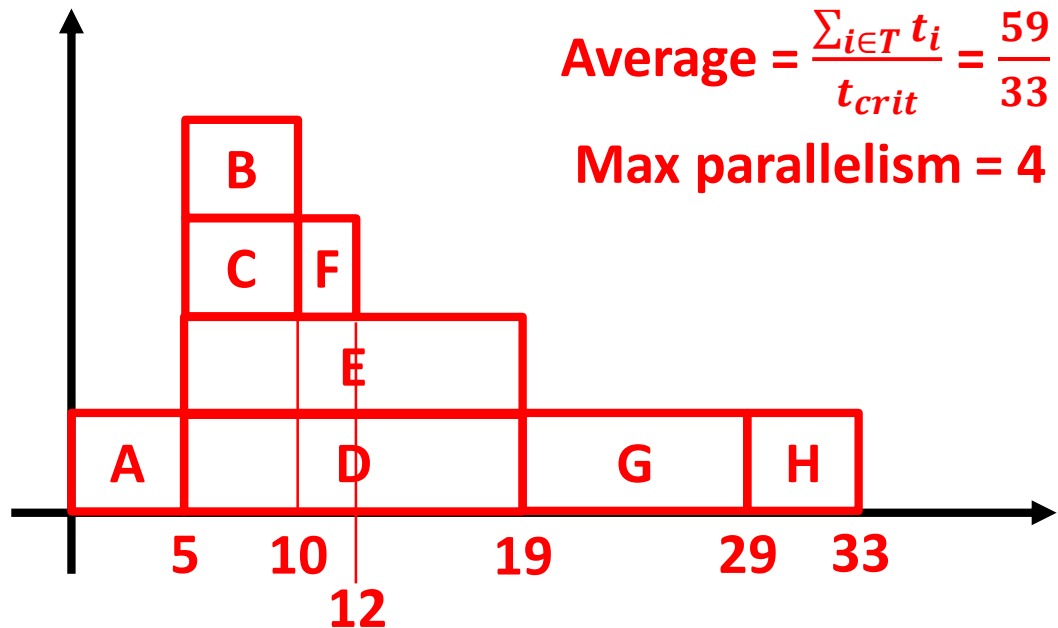
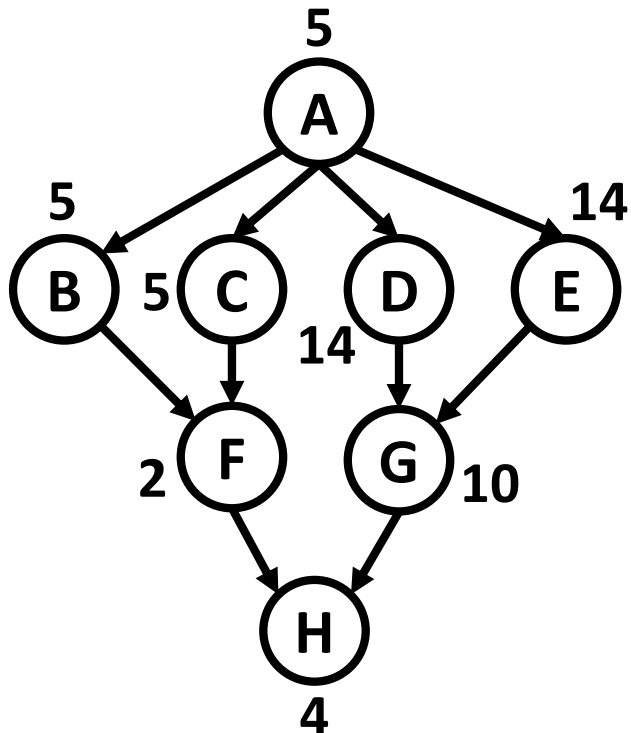
- Plot the degree of task parallelism over time for fragment 1, and derive its average and max parallelism



Q1(b)

Task Parallelism

- Plot the degree of task parallelism over time for fragment 2, and derive its average and max parallelism



Q1(c)

Task Parallelism

- What is the maximum achievable speedup for both fragments 1 and 2, given *infinite resources*?
 - The execution at each point cannot utilize more resources than the number of parallel tasks
 - **Therefore, this is equal to average degree of task parallelism**

Code fragment 1

Critical path $t_{crit} = 28$

$$\text{Average} = \frac{\sum_{i \in T} t_i}{t_{crit}} = \frac{59}{28}$$

Max parallelism = 4

Code fragment 2

Critical path $t_{crit} = 33$

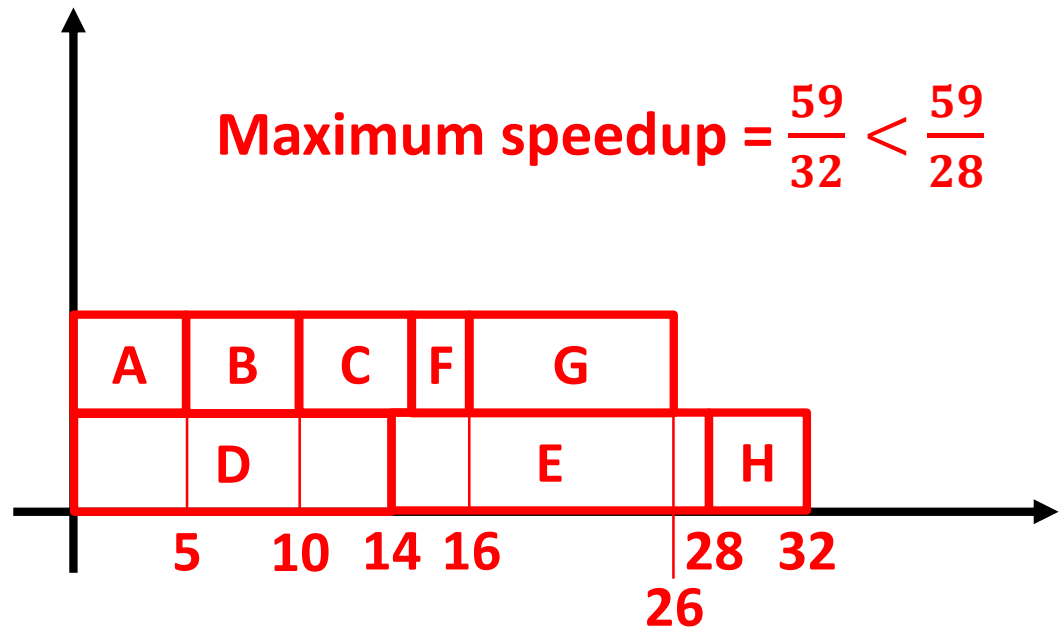
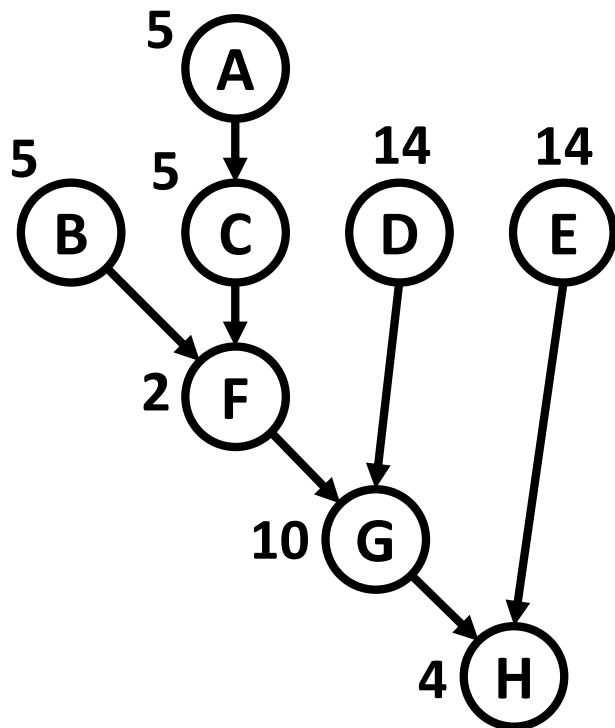
$$\text{Average} = \frac{\sum_{i \in T} t_i}{t_{crit}} = \frac{59}{33}$$

Max parallelism = 4

Q1(d)

Task Parallelism

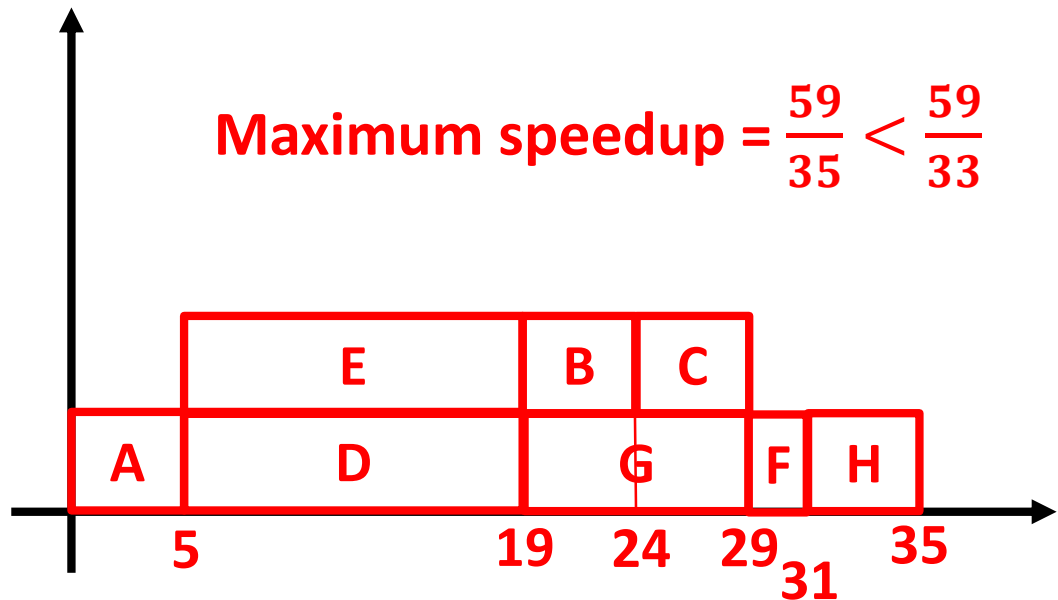
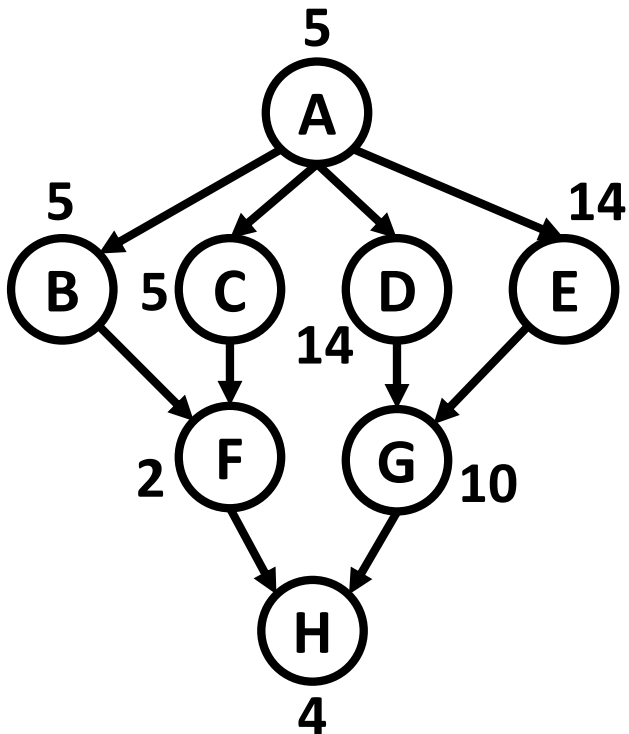
- If we only have 2 processing elements, what then is the maximum achievable speedup for fragment 1?



Q1(d)

Task Parallelism

- If we only have 2 processing elements, what then is the maximum achievable speedup for fragment 2?



Q2(a)

Parallel Programming Models

- Consider the matrix-vector multiplication problem
 - What type of parallelism does this problem entail (data or task)?
 - **Data parallelism**

$$\begin{array}{ccc} A & \times & x \\ \left[\begin{array}{c} \\ \\ \end{array} \right] & \times & \left[\begin{array}{c} \\ \\ \end{array} \right] & = & y \\ & & & = & \left[\begin{array}{c} \\ \\ \end{array} \right] \end{array}$$

$m \times n$ matrix
(m rows,
 n columns)

$n \times 1$ matrix
(n -dimensional
vector)

m -dimensional
vector

Q2(b)

Parallel Programming Models

- What type of architecture from Flynn's taxonomy would you choose to solve it with? Explain.
 - **SIMD (single-instruction multiple-data)**
- What parallel programming pattern would you employ to solve the matrix-vector multiplication problem?
 - Patterns: fork-join, parbegin-parend, SPMD/SIMD, master-worker, client-server, pipelining, task pool, producer-consumer
 - **Master-worker pattern/SIMD**

Q3

Performance of Parallel Systems

- Consider two processors P1 and P2 with the same ISA
 - P1 has a clock rate of 6 GHz, P2 has a clock rate of 2 GHz
 - Three types of instructions, with different CPI (cycles per instruction) for each on each processor
 - Three compilers C1, C2, C3 that generate programs with same total instructions (but different proportions)

Instructn Type	CPI (P1)	CPI (P2)	C1	C2	C3
FP arithmetic	4	2	30%	30%	50%
Int ADD/SUB	6	4	50%	20%	20%
Int MUL	8	3	20%	50%	30%

Q3(a)

Performance of Parallel Systems

- If C1 is used for both processors, how much faster is P1 than P2?

Instructn Type	CPI (P1)	CPI (P2)	C1	C2	C3
FP arithmetic	4	2	30%	30%	50%
Int ADD/SUB	6	4	50%	20%	20%
Int MUL	8	3	20%	50%	30%

- Average CPI $\overline{CPI} = \sum_{i=1}^n P_i \times CPI_i$
- Time taken $T = I \times \overline{CPI} \times T_{cyc} = \frac{N_{cyc}}{f}$
- % difference = $\frac{T_2 - T_1}{T_2} = 39.6\%$

Q3(b)

Performance of Parallel Systems

- If C2 is used for both processors, how much faster is P1 than P2?

Instructn Type	CPI (P1)	CPI (P2)	C1	C2	C3
FP arithmetic	4	2	30%	30%	50%
Int ADD/SUB	6	4	50%	20%	20%
Int MUL	8	3	20%	50%	30%

- Average CPI $\overline{CPI} = \sum_{i=1}^n P_i \times CPI_i$
- Time taken $T = I \times \overline{CPI} \times T_{cyc} = \frac{N_{cyc}}{f}$
- % difference = $\frac{T_2 - T_1}{T_2} = 26.4\%$

Q3(c,d)

Performance of Parallel Systems

- Which of the three compilers is best for each of P1 and P2? **(Hint: do we need to consider the clock frequency?)**

Instructn Type	CPI (P1)	CPI (P2)	C1	C2	C3
FP arithmetic	4	2	30%	30%	50%
Int ADD/SUB	6	4	50%	20%	20%
Int MUL	8	3	20%	50%	30%

Compiler	Weighted CPI (P1)	Weighted CPI (P2)
C1		
C2		
C3		

Q4(a)

Parallel Performance Scaling

- Consider a program that for problem size N has a:
 - Sequential portion: with N instructions
 - Parallel portion: with N^2 instructions
- Suppose $N = 100$; what is the speedup with 10 and 100 processors respectively? What is the upper limit?
 - Use Amdahl's law for p processors and sequential fraction f

$$S_p(N) = \frac{1}{f + \frac{1-f}{p}}$$

Q4(a)

Parallel Performance Scaling

- Suppose $N = 100$; what is the speedup with 10 and 100 processors respectively? What is the upper limit?

$$f = \frac{I_{seq}}{I_{total}} = \frac{N}{N + N^2} = \frac{1}{1 + N} = \frac{1}{101}$$

$$S_{10}(100) = \frac{1}{\frac{1}{101} + \frac{100}{101(10)}} = 9.181$$

$$S_{100}(100) = \frac{1}{\frac{1}{101} + \frac{100}{101(100)}} = 50.5$$

$$S_{\infty}(100) = \frac{1}{\frac{1}{101} + 0} = \frac{1}{f} = 101$$

Q4(b)

Parallel Performance Scaling

- Assume each instruction can be executed in 2 cycles on a 1 GHz processor
 - Calculate the time needed to run the program with $N = 100$ on a single core

$$I_{N=100} = N + N^2 = 100(101) = 10100$$

$$t_{N=100} = \frac{2 \times I_{N=100}}{f} = \frac{20200}{10^9} = 0.0000202s$$

Q4(b)

Parallel Performance Scaling

- Given the same amount of time, what is the problem size we can solve with 10 and 100 processors?

$$0.0000202 = t(N, p) = \left(N + \frac{N^2}{p} \right) \times \frac{2}{10^9}$$

$\frac{N^2}{p \times 10^9} + \frac{N}{10^9} - 0.0000101 = 0$

$p = 10 \Rightarrow N \approx 312$

$p = 100 \Rightarrow N \approx 956$

CPI

f

I_{worst}
(on processor that executes the sequential part)

- What is the speedup for these problem sizes and processor counts? 😊

Admin

Tutorial Quiz

- Go to Luminus > Quiz
 - 1% of your grade
 - 8 minutes for 4 MCQ questions
 - **Don't be stressed - full credit for $\geq 60\%$ correct (you should be able to do this after the tutorial, well, hopefully)**
- 10 minutes now to attempt it

CS3210

Parallel Computing

Thank you! Any questions?



Tut 2

Mon (4pm)

Tues (2pm)