


CS3210

Parallel Computing

Changes from Monday in Green



Tut 3

Mon (4pm)

Tues (2pm)

Admin Updates

- **Assignment 2 released last weekend - due 19 Oct, 11am**
 - Topic: CUDA parallelisation of hash-based proof-of-work
 - Weightage: 10%
 - **Please start on this early - there are much fewer reserved Compute Cluster nodes than lab machines**
- Assignment 1 still being graded (no ETA)
 - Late submissions for Part 2 are still accepted
- End-tutorial Quiz 3 later

Review

Memory Coherence (I)

- Cache coherence problem arises on multi-processor systems where each processor has a private cache
 - Multiple copies of the same data exists in different caches
 - **A write by one processor to a location may not become visible
→ other processors may keep reading stale data**
- Why not unify all private caches into a shared cache?
 - Increased latency from larger cache size
 - Central bandwidth bottleneck and negative interference (cache conflicts/thrashing)

Review

Memory Coherence (II)

- Goal: reading a memory location should return the last value written by any processor
- Three properties
 - **Program order**: preserves sequential order of each processor
 - **Write propagation**: writes by one processor must become visible to all others
 - **Write serialization**: all writes to a location must be seen in the same order by all processors

Review

Coherence and Consistency

- Coherence: defines requirements for observed behavior of reads and writes to the same memory location
 - All processors must agree on order of reads/writes to X
- Consistency: defines behavior (constraints) of reads and writes to different memory locations, as observed by other processors
 - Coherence only guarantees that writes to address X will eventually propagate to SMs
 - Consistency deals with when writes to X propagate to other processors, relative to reads and writes to other addresses

Review

Memory Operation Orderings

- A program defines a sequence of loads and stores: a program order
- Four types of memory operation orderings
 - $W \rightarrow R$: write to X commits before subsequent read from Y
 - $R \rightarrow R$: read from X commits before subsequent read from Y
 - $R \rightarrow W$: read to X commits before subsequent write to Y
 - $W \rightarrow W$: write to X commits before subsequent write to Y

Review

Consistency Models (I)

- **Sequential consistency: preserves all four orderings**
 - Effect of each memory operation must be visible to all processors prior to next operation on any processor
- **Relaxed consistency: allow relaxation of ordering of memory operations in absence of data dependencies**
 - Data dependency arises when one operation writes to X, which is then read by a later operation
 - Goal: hide memory latencies to improve performance

Review

Consistency Models (II)

- Summary table (weak consistency models not covered)

Property	SC	RC: TSO	RC: PC	RC: PSO
Preserves data dep?	Yes	Yes	Yes	Yes
Preserve $W \rightarrow R$	Yes	No	No	No
Preserve $R \rightarrow R$	Yes	Yes	Yes	Yes
Preserve $R \rightarrow W$	Yes	Yes	Yes	Yes
Preserve $W \rightarrow W$	Yes	Yes	Yes	No
Write propagation	Yes	Yes	No	Yes

Return value of any write (even from other processor) before write is propagated or serialised



Q1

Memory Consistency

- Consider this shared memory program on 3 processors

P1

```
[ 1] X = 1  
[ 2] Y = 1  
[ 3] while (T == 0);  
[ 4] print A
```

P2

```
[ 5] while (Y == 0);  
[ 6] while (Z == 0);  
[ 7] T = X  
[ 8] A = X  
[ 9] B = X
```

P3

```
[10] Z = 2  
[11] X = 3  
[12] while (T == 0);  
[13] print B
```

- Goal of this program: ensure **print A** and **print B** print same value
 - All variables are initialized to 0

P1

```
[ 1] X = 1  
[ 2] Y = 1  
[ 3] while (T == 0);  
[ 4] print A
```

P2

```
[ 5] while (Y == 0);  
[ 6] while (Z == 0);  
[ 7] T = X  
[ 8] A = X  
[ 9] B = X
```

P3

```
[10] Z = 2  
[11] X = 3  
[12] while (T == 0);  
[13] print B
```

Q1(a)

Memory Consistency

- Under the sequential consistency model, what are the possible values of A , B , T , X after the program finishes?
 - Note this is different from the values printed by P1 and P3!
- What is preserved under sequential consistency?
 - All four types of memory orderings
 - Write propagation and serialization (implied)

P1

```
[ 1] X = 1  
[ 2] Y = 1  
[ 3] while (T == 0);  
[ 4] print A
```

P2

```
[ 5] while (Y == 0);  
[ 6] while (Z == 0);  
[ 7] T = X  
[ 8] A = X  
[ 9] B = X
```


P3

```
[10] Z = 2  
[11] X = 3  
[12] while (T == 0);  
[13] print B
```

Q1(a)

Memory Consistency

- What can be inferred from this instruction dependency graph?



P1

```
[ 1] X = 1
[ 2] Y = 1
[ 3] while (T == 0);
[ 4] print A
```

P2

```
[ 5] while (Y == 0);
[ 6] while (Z == 0);
[ 7] T = X
[ 8] A = X
[ 9] B = X
```

P3

```
[10] Z = 2
[11] X = 3
[12] while (T == 0);
[13] print B
```

Q1(a)

Memory Consistency

- Answer

A	B	T	X	Possible execution scenario
1	1	1	1	10, 11, 1, 2, 5 - 9, 3, 4, 12, 13
1	1	1	3	1, 2, 10, 5 - 9, 3, 4, 11, 12, 13
1	3	1	3	1, 2, 10, 5 - 8, 3, 4, 11, 9, 12, 13
3	3	1	3	1, 2, 10, 5 - 7, 11, 8, 9, 3, 4, 12, 13
3	3	3	3	1, 2, 10, 11, 5 - 9, 3, 4, 12, 13

- Prove the following are impossible:

A	B	T	X	Reason
1	3	1	1	B = 3 implies final value of X = 3
3	1	3	1	A = 3 implies B = 3 by prog. order

P1


```
[ 1] X = 1
[ 2] Y = 1
[ 3] while (T == 0);
[ 4] print A
```

P2

```
[ 5] while (Y == 0);
[ 6] while (Z == 0);
[ 7] T = X
[ 8] A = X
[ 9] B = X
```

P3

```
[10] Z = 2
[11] X = 3
[12] while (T == 0);
[13] print B
```



Q1(b)

Memory Consistency

- The code fails its goal even under sequential consistency - give one such execution scenario


➤ General idea: ensure (11) gets executed after (8) but before (9)

- Sample execution


➤ (1) → (2) → (10) → (5) → (6) → (7) → (8) → (3) → (4) [A = 1] → (11) → (9) → (12) → (13) [B = 3]

P1
[1] X = 1
[2] Y = 1
[3] while (T == 0);
[4] **print A**

P2
[5] while (Y == 0);
[6] while (Z == 0);
[7] T = X
[8] A = X
[9] B = X



P3
[10] Z = 2
[11] X = 3
[12] while (T == 0);
[13] **print B**



Q1(c)

Memory Consistency

- Rewrite the code such that the code achieves its goal under sequential consistency
 - General idea: prevent (11) from executing in between (8) and (9)
- Answer
 - Swap (10) and (11) so that X = 3 is set first before Z = 2
 - Move (7) after (8) and (9) to ensure both P1 and P3 can only print after both A and B are modified to the same value of X

P1

```
[ 1] X = 1
[ 2] Y = 1
[ 3] while (T == 0);
[ 4] print A
```

P2

```
[ 5] while (Y == 0);
[ 6] while (Z == 0);
[ 8] A = X
[ 9] B = X
[ 7] T = X
```

P3

```
[11] X = 3
[10] Z = 2
[12] while (T == 0);
[13] print B
```

Q1(d)

Memory Consistency

- Would this modified program fail under relaxed consistency? If so, which particular models?
 - What instruction orderings are relaxed by each of the three RC models: TSO, PC, PSO?
- Answer
 - Does not fail under TSO and PC
 - Fails under PSO, as it allows $W \rightarrow W$ relaxation (can you give a sample execution scenario?)

Q2

Relaxed Consistency

- Consider this shared memory program on 3 processors

P1

```
[ 1] X = 5  
[ 2] Y = 2  
[ 3] print Y
```

P2

```
[ 4] while (X == 0);  
[ 5] Y = 6  
[ 6] print Z
```

P3

```
[ 7] while (Y == 0);  
[ 8] Z = 5  
[ 9] print X
```

- This program was run on systems that utilize different *relaxed consistency models*, namely TSO, PC and PSO
 - All variables are initialized to 0

P1
[1] X = 5
[2] Y = 2
[3] **print Y**


P2
[4] while (X == 0);
[5] Y = 6
[6] **print Z**

P3
[7] while (Y == 0);
[8] Z = 5
[9] **print X**

Q2(a)

Relaxed Consistency

- What can be inferred from the instruction dependency graph?



- Observe (7) can execute once execution of one of (3) or (5) becomes visible to P3

```
[ 1] X = P1 5
[ 2] Y = 2
[ 3] print Y
```

```
[ 4] while (X == 0);
[ 5] Y = 6
[ 6] print Z
```

```
[ 7] while (Y == 0);
[ 8] Z = 5
[ 9] print X
```

Q2(a)

Relaxed Consistency

- In one run, the following output was observed
 - $X = 5, Y = 2, Z = 0$
- Under which RC model(s) is this output valid, and why?
 - This is already valid under sequential consistency! (Can you give the execution scenario?)
 - Since SC is more stringent than all RC models, it is valid under all RC models (TSO, PC, PSO)

P1
[1] X = 5
[2] Y = 2
[3] print Y

P2
[4] while (X == 0);
[5] Y = 6
[6] print Z

P3
[7] while (Y == 0);
[8] Z = 5
[9] print X

Q2(b)

Relaxed Consistency

- In one run, the following output was observed
 - X = 0, Y = 2, Z = 5
- Under which RC model(s) is this output valid, and why?
 - Only PC and PSO (invalid under TSO)
 - Notice X = 0 is printed by P3 ⇒ Y has been set by at least one of P1 or P2
 - W → W not relaxed: Y set ⇒ X set by P1 ⇒ X = 5 not seen by P3 (PC)
 - W → W relaxed: then Y = 2 set before X = 5 ⇒ lets P3 print 0 (PSO)

Q3

Parallel Programming Models

- Consider the problem of matrix-vector multiplication on a GPU with CUDA: $A \cdot b = c$
 - A is an $N \times N$ matrix and b, c are both vectors of size N
 - i.e. $m = n$ below

$$\begin{array}{ccc} A & \times & x \\ \left[\begin{array}{c} \\ \\ \end{array} \right] & \times & \left[\begin{array}{c} \\ \\ \end{array} \right] & = & y \\ & & & = & \left[\begin{array}{c} \\ \\ \end{array} \right] \end{array}$$

$m \times n$ matrix
(m rows,
 n columns)

$n \times 1$ matrix
(n -dimensional
vector)

m -dimensional
vector

Q3(a)

Parallel Programming Models

- How would you classify the memory model for CUDA: shared, distributed or distributed-shared? Explain with reference to the CUDA memory hierarchy.
 - **Distributed-shared memory model**
 - **Different types of memory accessible at different scope**
 - Distributed (thread-private): registers, local memory
 - Shared (block/kernel level): shared memory, global memory

Q3(b)

Parallel Programming Models

- A kernel uses 64 registers and is launched with each block containing 512 threads. The kernel requires little local and shared memory and is run on a GPU device of Compute Capability 6.x with 65536 registers
- How many blocks and warps can reside on the GPU concurrently?

➤ **Blocks:** $N_{blocks} = \frac{N_{registers}}{N_{registers/block}} = \frac{65536}{512 \times 64} = 2$

➤ **Warps:** $N_{warps} = N_{blocks} \times N_{warps/block} = 2 \times \frac{512}{32} = 32$

Q3(c)

Parallel Programming Models

- Discuss how you would implement matrix-vector multiplication in CUDA, detailing
 - Types of memory the kernel will use
 - Grid dimensions, block dimensions (and thus number of threads in a block)
 - Kernel code
 - Code that invokes the kernel
- In what ways can we decompose the problem into tasks to be assigned to CUDA threads?

Q3(c)

Parallel Programming Models

- **Considerations**

- **Decomposition** into independent tasks: each task as a multiplication/addition, row of A, or multiple rows from A
- **Dimensionality**: for N tasks, 1D grid of $B = \frac{N}{T}$ blocks, each 1D block containing T threads, depends on number of tasks
- $T = 32i$ for some $i \in \mathbb{Z}^+$, for **even division into warps**
- Kernel code - just a for loop, then write result to one cell of c
- **Bijective mapping** from threads to a row/chunk of rows in A
- Global memory: slow but cached in L1D\$, **coalesced txns**
- Shared memory: fast but poor performance if **bank conflicts** arise (especially for common vector b) → need **strided access**

Q3(d)

Parallel Programming Models

- A CUDA implementation of the matrix-vector multiplication stores matrix A and vectors b, c all in shared memory. However, the size of the data exceeds the available shared memory
- How would you change your implementation to continue to use shared memory for large matrix sizes?
 - Only bring in sections of the required rows of the matrix A into shared memory, i.e. stream matrix A from global memory
 - Alternatively, launch kernel with more blocks of smaller size

Admin


Tutorial Quiz

- Go to Luminus > Quiz
 - 1% of your grade
 - 8 minutes for 4 MCQ questions
 - **Don't be stressed - full credit for $\geq 50\%$ correct (you should be able to do this after the tutorial, well, hopefully)**
- 10 minutes now to attempt it

CS3210

Parallel Computing

Thank you! Any questions?



Tut 3

Mon (4pm)

Tues (2pm)