

# Lab 5

## Distributed-Memory Programming using MPI

CS3210 – 2020/21 Semester 1

### Learning Outcomes

1. Learn to use collective communication
2. Learn how to create, destroy and manage new MPI communicators
3. Learn how to arrange MPI processes into a Cartesian virtual topology

Lab 4 provided you with basic knowledge on MPI programming. This lab aims to provide a more detailed coverage of MPI library calls.

## Part 1: Collective Communication

MPI provides collective communication functions which must involve (be invoked by) all processes in the scope of a communicator. By default, all processes are members of the global communicator `MPI_COMM_WORLD`.



### **Important:**

It is the programmer's responsibility to ensure that all processes within a communicator participate in any collective operations! Failure to do so may result in a deadlock.

There are three types of collective communication:

1. Synchronization communication
2. Data movement operations
3. Collective computation (data movement with reduction operations)

### Part 1.1: Synchronization Communication

There is only one collective synchronization operation in MPI - a barrier:



```
int MPI_Barrier(MPI_Comm comm)
```

All processes in the MPI communicator `comm` will block until all of them reach the barrier. Failure to call this function from all the processes will result in deadlock.

If you are interested, there is a non-blocking variant `MPI_Ibarrier` which returns immediately, independent of whether other processes have called `MPI_Ibarrier`. The barrier semantics are then only enforced at the corresponding completion operation (`MPI_Test` or `MPI_Wait`) with the `MPI_Request` provided to `MPI_Ibarrier`.

## Part 1.2: Data Movement Operations

The data movement (distribution) operations provided by MPI are:



```
-----  
/* MPI_Bcast - broadcasts (sends) a message from the process with rank  
   root to all other processes in the group */  
  
int MPI_Bcast(void *buffer, int count, MPI_Datatype datatype,  
              int root, MPI_Comm comm)  
-----  
  
/* MPI_Scatter - sends data from one process to all processes in  
   a communicator */  
  
int MPI_Scatter(const void *sendbuf, int sendcount, MPI_Datatype sendtype,  
               void *recvbuf, int recvcount, MPI_Datatype recvtype,  
               int root, MPI_Comm comm)  
-----  
  
/* MPI_Gather - gathers data from a group of processes into one root  
   process */  
  
int MPI_Gather(const void *sendbuf, int sendcount, MPI_Datatype sendtype,  
              void *recvbuf, int recvcount, MPI_Datatype recvtype,  
              int root, MPI_Comm comm)  
-----  
  
/* MPI_Allgather - gathers data from a group of processes into every  
   process of that group */  
  
int MPI_Allgather(const void *sendbuf, int sendcount,  
                  MPI_Datatype sendtype, void *recvbuf, int recvcount,  
                  MPI_Datatype recvtype, MPI_Comm comm)  
-----  
  
/* MPI_Alltoall - each process in a group performs a scatter operation,  
   sending a distinct message to all the processes in the group in order  
   by their rank */  
  
int MPI_Alltoall(const void *sendbuf, int sendcount,  
                 MPI_Datatype sendtype, void *recvbuf, int recvcount,  
                 MPI_Datatype recvtype, MPI_Comm comm)  
-----
```



The MPI-3 standard introduced Non-blocking Collective Data Movement operations such as `MPI_Ibroadcast` or `MPI_Iscatter`. Additionally, some of the operations above possess a variant that allows you to specify a varying number of data items to sent to each process, such as `MPI_Scatterv` or `MPI_Allgatherv`. To learn more, check the [OpenMPI documentation](#) for details of the non-blocking and variable variants of the above operations.

## Part 1.3: Collective Computation

The MPI functions that enable collective computations are:



```
-----  
/* MPI_Reduce - reduces values on all processes within a group; the  
   reduction operation must be one of the following:  
   MPI_MAX maximum | MPI_MIN minimum | MPI_SUM sum | MPI_PROD product |  
   MPI_LAND logical AND | MPI_BAND bit-wise AND | MPI_LOR logical OR |  
   MPI_BOR bit-wise OR | MPI_LXOR logical XOR | MPI_BXOR bit-wise XOR |  
   MPI_MAXLOC max value and location | MPI_MINLOC min value and location  
*/  
  
int MPI_Reduce(const void *sendbuf, void *recvbuf, int count,  
               MPI_Datatype datatype, MPI_Op op, int root, MPI_Comm comm)  
-----  
  
/* MPI_Allreduce - applies a reduction operation and places the result in  
   all processes in the communicator (this is equivalent to an MPI_Reduce  
   followed by an MPI_Bcast) */  
  
int MPI_Allreduce(const void *sendbuf, void *recvbuf, int count,  
                  MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)  
-----  
  
/* MPI_Reduce_scatter - first performs an element-wise reduction on a  
   vector across all processes in the group, then splits the result vector  
   into disjoint segments to distribute across the processes (this is  
   equivalent to an MPI_Reduce followed by an MPI_scatter) */  
  
int MPI_Reduce_scatter(const void *sendbuf, void *recvbuf,  
                       const int recvcounts[], MPI_Datatype datatype,  
                       MPI_Op op, MPI_Comm comm)  
-----
```



### Exercise 1

Compile the program `col_comm.c` and run it. Explore the code and output to understand how scatter works in this example. What can we do if we only want to scatter the data only to a subset of the processes of the MPI communicator?

## Part 2: Managing Communicators

One of the major disadvantages of using collective communication is that we must involve all the processes in an MPI communicator. To overcome this, MPI allows us to create custom communicators, add / remove processes to / from communicators and destroy communicators as needed. **Note that a communicator comprises a set of processes (a `MPI_Group`) with an associated context.**

The MPI functions for communicator management are:



```
-----  
/* MPI_Comm_group - returns the group associated with a communicator */  
  
int MPI_Comm_group(MPI_Comm comm, MPI_Group *group)  
-----  
/* MPI_Group_incl - produces a group by reordering an existing group and  
   taking only listed members */  
  
int MPI_Group_incl(MPI_Group group, int n, const int ranks[],  
                   MPI_Group *newgroup)  
-----  
/* MPI_Comm_create - creates a new communicator with a group of  
   processes */  
  
int MPI_Comm_create(MPI_Comm comm, MPI_Group group, MPI_Comm *newcomm)  
-----  
/* MPI_Group_rank - returns the rank of the calling process in the given  
   group */  
  
int MPI_Group_rank(MPI_Group group, int *rank)  
-----  
/* MPI_Comm_rank - returns the rank of the calling process in the given  
   communicator */  
  
int MPI_Comm_rank(MPI_Comm comm, int *rank)
```



### Exercise 2

Compile the program `new_comm.c` and run it. What does the program do?

## Part 3: Cartesian Virtual Topologies

In the context of MPI, a virtual topology describes a mapping/ordering of MPI processes into a geometric space. Generally, there are two main types of topologies supported by MPI, namely, (i) Cartesian and (ii) graph. In this lab, we only introduce the Cartesian virtual topology.

Depending on the problem, the MPI processes may access data in a regular structured pattern in Cartesian space. In these cases, it is useful to arrange the logical MPI processes into a Cartesian virtual topology to facilitate programming and communication. **Remember that there may be no relation between the physical organisation (layout) of the parallel machines and the MPI virtual topology.** Additionally, the MPI topology must be configured and managed by the programmer.

MPI provides three functions to help us manage a Cartesian topology:



```
-----  
/* MPI_Cart_create - makes a new communicator to which Cartesian topology  
   information has been attached */  
  
int MPI_Cart_create(MPI_Comm comm_old, int ndims, const int dims[],  
                   const int periods[], int reorder, MPI_Comm *comm_cart)  
-----  
  
/* MPI_Cart_coords - determines process coordinates in the Cartesian  
   topology, given its rank in the group */  
  
int MPI_Cart_coords(MPI_Comm comm, int rank, int maxdims, int coords[])  
-----  
  
/* MPI_Cart_shift - returns the shifted source and destination ranks,  
   given a shift direction and amount */  
  
int MPI_Cart_shift(MPI_Comm comm, int direction, int disp,  
                  int *rank_source, int *rank_dest)  
-----
```



### Exercise 3

cart.c is an example implementation of a Cartesian topology. Study the program and understand how it works. Compile cart.c and run it. What does the program do?



More information and further reading:

- LLNL MPI Tutorial: <https://computing.llnl.gov/tutorials/mpi/>
- OpenMPI FAQ: <https://www.open-mpi.org/faq/?category=running>
- Another MPI Tutorial: <http://mpitutorial.com/>
- Short Video on Cartesian Topology: <https://youtu.be/dyV0dlC0y7w>
- Advanced Parallel Programming with MPI: [https://htor.inf.ethz.ch/teaching/mpi\\_tutorials/ppopp13/2013-02-24-ppopp-mpi-advanced.pdf](https://htor.inf.ethz.ch/teaching/mpi_tutorials/ppopp13/2013-02-24-ppopp-mpi-advanced.pdf)