CS3210
Introduction

Lecture 1

Outline

- ## Parallel Computing
  - ❑ Motivation
  - ❑ Serial vs Parallel Computing
  - ❑ Serial vs Parallel Processing Units

- ## Basic of Parallel Computing
  - ❑ Decompose, Schedule and Map
  - ❑ Performance
  - ❑ Challenges

- ## Landscape of Parallel Computing
  - ❑ Applications
  - ❑ Systems

Cheetah   vs   Hyena

- Weigh 60 kg
- Fastest land animal
  - Top speed of 120 km/h
  - 0-100 km/h in 3 seconds
- Hunt alone
  - Usually over in 1 minute
- 50% of losing its catch to other predators

- Weigh 50 kg
- Top speed of 65 km/h

- Hunt in group of 4-40
- Can take down large animals like buffalo, rhino, lion

Observations

- Optimized on speed only

- Single Processor:
  - Failed to meet Moore's Law since early 2000

- Optimized to **cooperate** and **work in parallel**

- Parallel Computer:
  - A collection of processing elements that cooperate to solve problems quickly
  - Huge amount of computation power

Parallel Computing - Motivation

- In the first issue of CACM (1958), **Saul Gorn** notes that:

> "We know that the so-called parallel computers are somewhat faster than the serial ones, and that no matter how fast we make **access and arithmetic** serially, we can do correspondingly better in parallel. However access and arithmetic speeds seem to be approaching a definite limit ... Does this mean that digital speeds are reaching a limit, that digital computation of multi-variate partial differential systems must accept it, and that **predicting tomorrow's weather must require more than one day**? Not if we have truly-parallel machines, with say, a thousand instruction registers."

Hitting the Wall

- '70s: "Supercomputers" for Scientific Computing

- Early '90s: Databases

- Inflection point in 2004: Intel hits the Power Density Wall

> "obtaining more computing power by stamping multiple processors on a single chip rather than straining to increase the speed of a single processor."
>
> John Markoff, New York Times, May 17, 2004

How do you make your program run faster?

- **Answer *before 2004*:**
  - Just wait 6 months and buy a new machine!
- **Answer *after 2004*:**
  - Write parallel software

Parallel Computing - Challenges

■ Gorn also recognized the challenges that would be posed to the system designer by the new parallel computers:

"But visualize what it would be like to program for such a machine! If a thousand subroutines were in progress simultaneously, one must program the recognition that they have all finished, if one is to use their results together. The programmer must not only synchronize his subroutines, but schedule all his machine units, unless he is willing to have most of them **sitting idle most of the time**. Not only would **programming techniques be vastly different from the serial ones** in the serial languages we now use, but they would be humanly impossible without machine intervention."
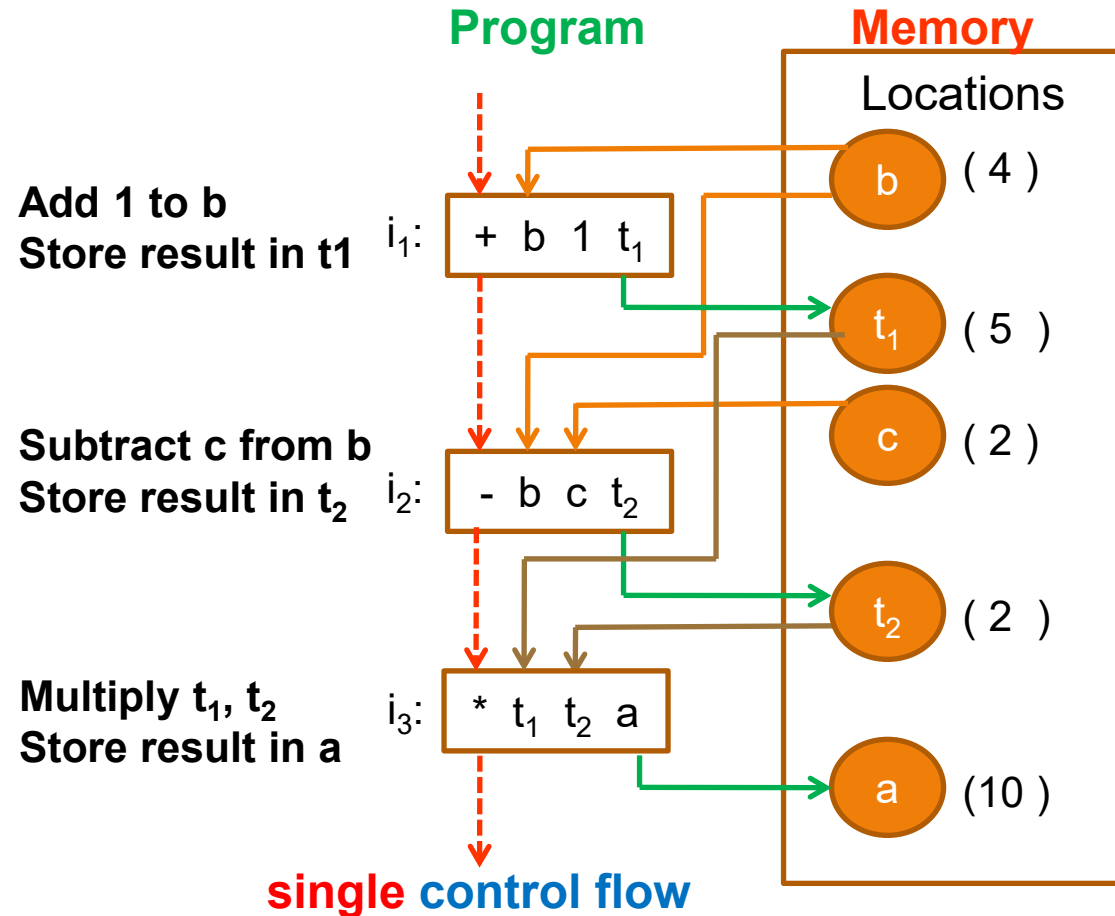
Serial Computing
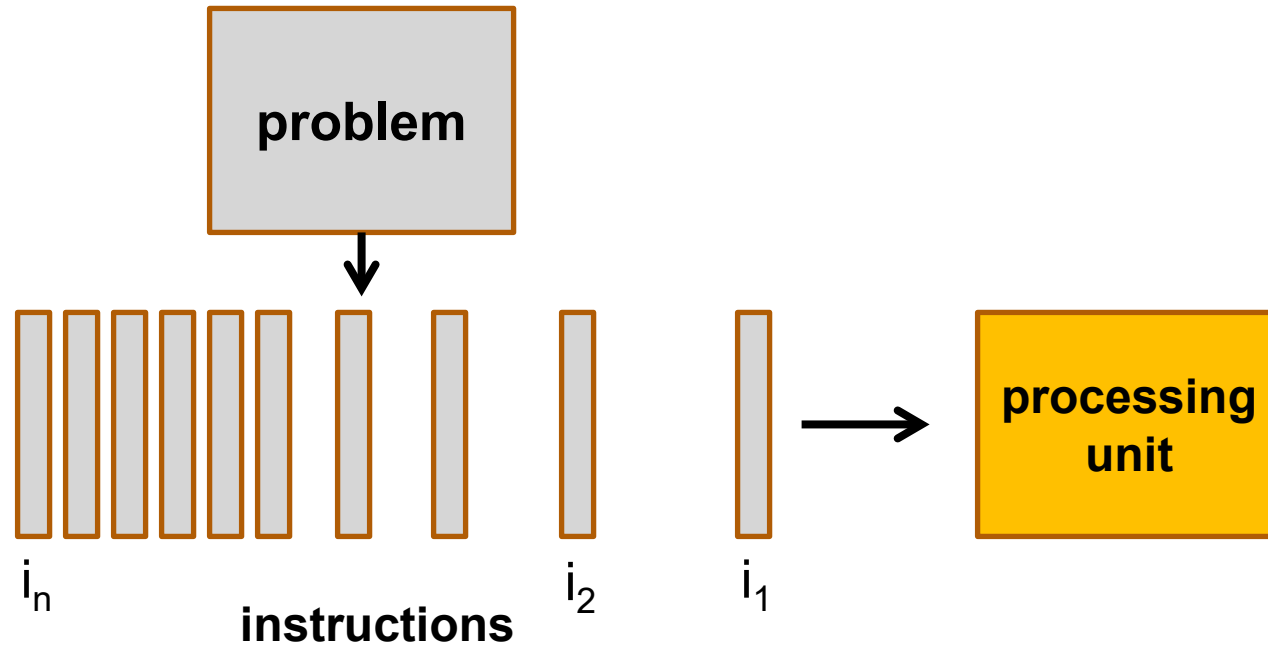
- Traditionally, a problem is divided into a discrete series of instructions
  - Instructions are executed one after another
  - Only one instruction executed at any moment in time

Example: Evaluate `a = (b+1) * (b-c)`

Program

Memory

Locations

Add 1 to b
Store result in t1

i_1: `+ b 1 t₁`

b (4)

t₁ (5)

Subtract c from b
Store result in t₂

i_2: `- b c t₂`

c (2)

t₂ (2)

Multiply t₁, t₂
Store result in a

i_3: `* t₁ t₂ a`

a (10)

single control flow

Sequential Computation Model

Processing Unit: von Neumann Model
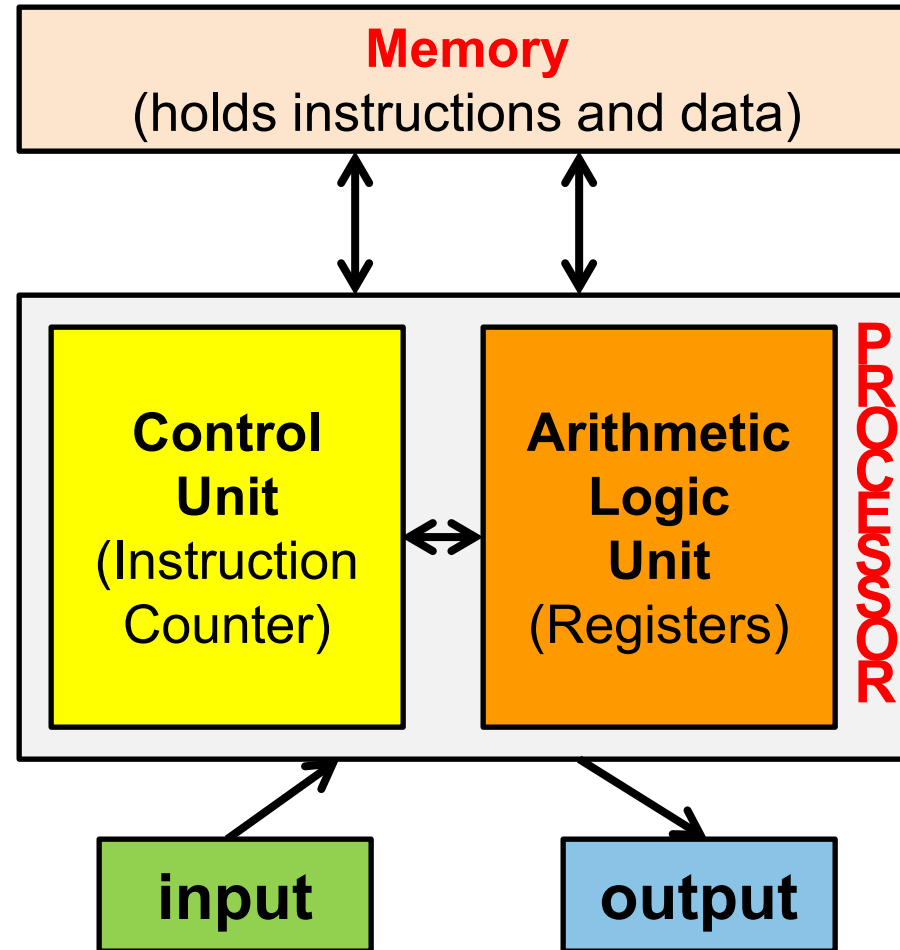
von Neumann Computation Model

- Processor:
  - ❏ Performs instructions
- Memory:
  - ❏ Stores both the instructions and data in cells with unique addresses
- Control Scheme:
  - ❏ Fetches instruction one after another from memory
  - ❏ Shuttles data between memory and processor
- Memory wall
  - ❏ ➔ disparity between processor (< 1 nanosecond or 10^{-9} sec) and memory speed (100-1,000 nanoseconds)

Ways to improve performance

1. **Work hard**   ➢ higher clock frequency

2. **Work smart**   ➢ pipelining, superscalar processor

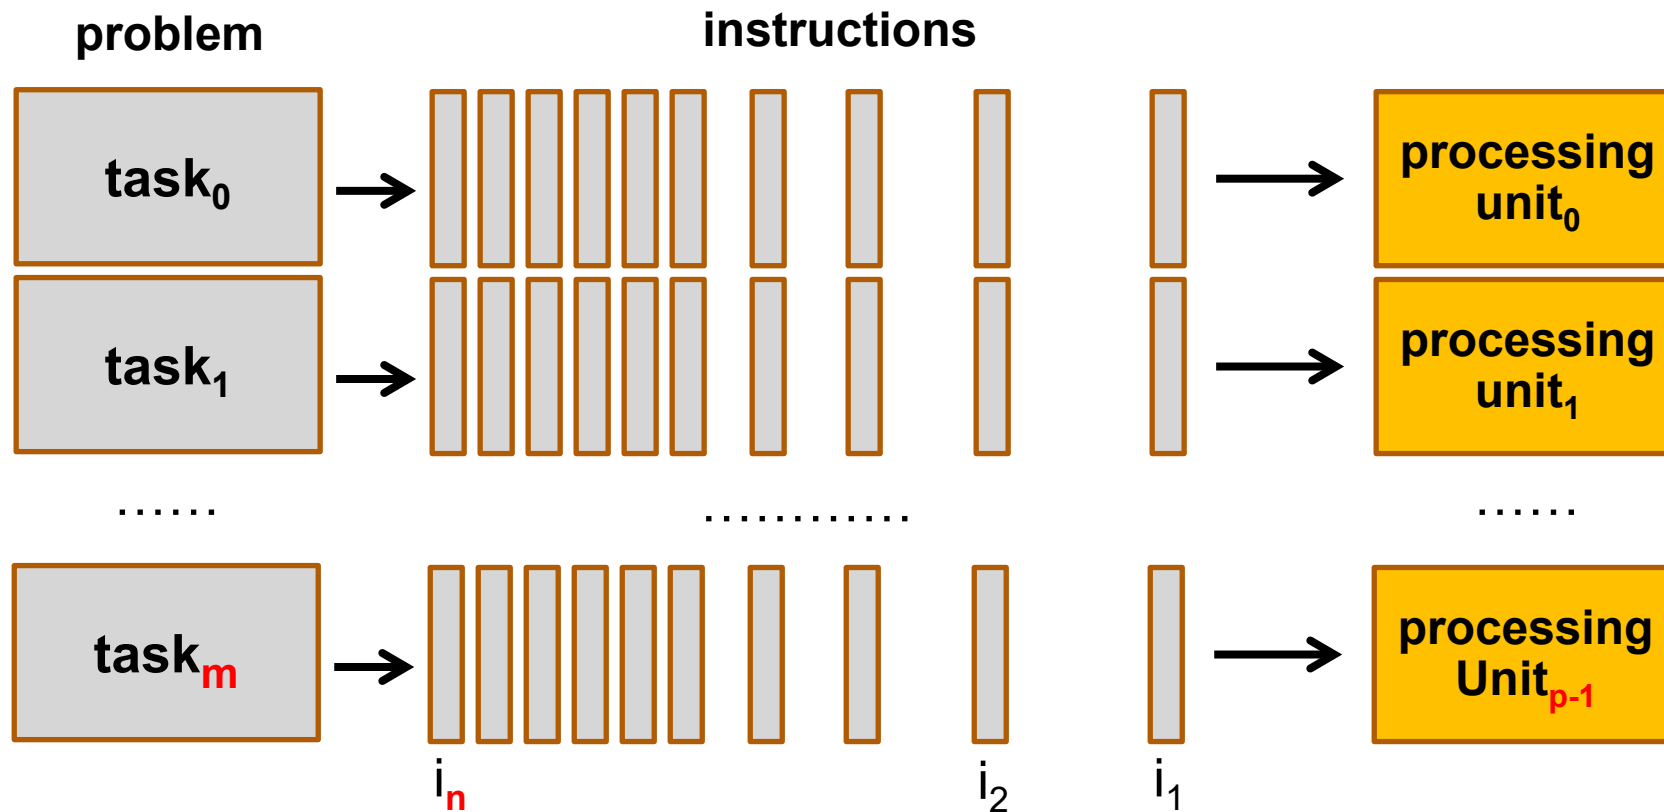3. **Get help**   ➢ replication – multicore, cluster, ..

Parallel Computing

- **Simultaneous use** of **multiple processing units** to solve a problem fast / solve a larger problem

- **Processing units** could be:
    i. A single processor with multiple cores
    ii. A single computer with multiple processors
    iii. A number of computers connected by a network
    iv. Combinations of the above

- Ideally, a problem (application) is partitioned into sufficient number of independent parts for execution on parallel processing elements
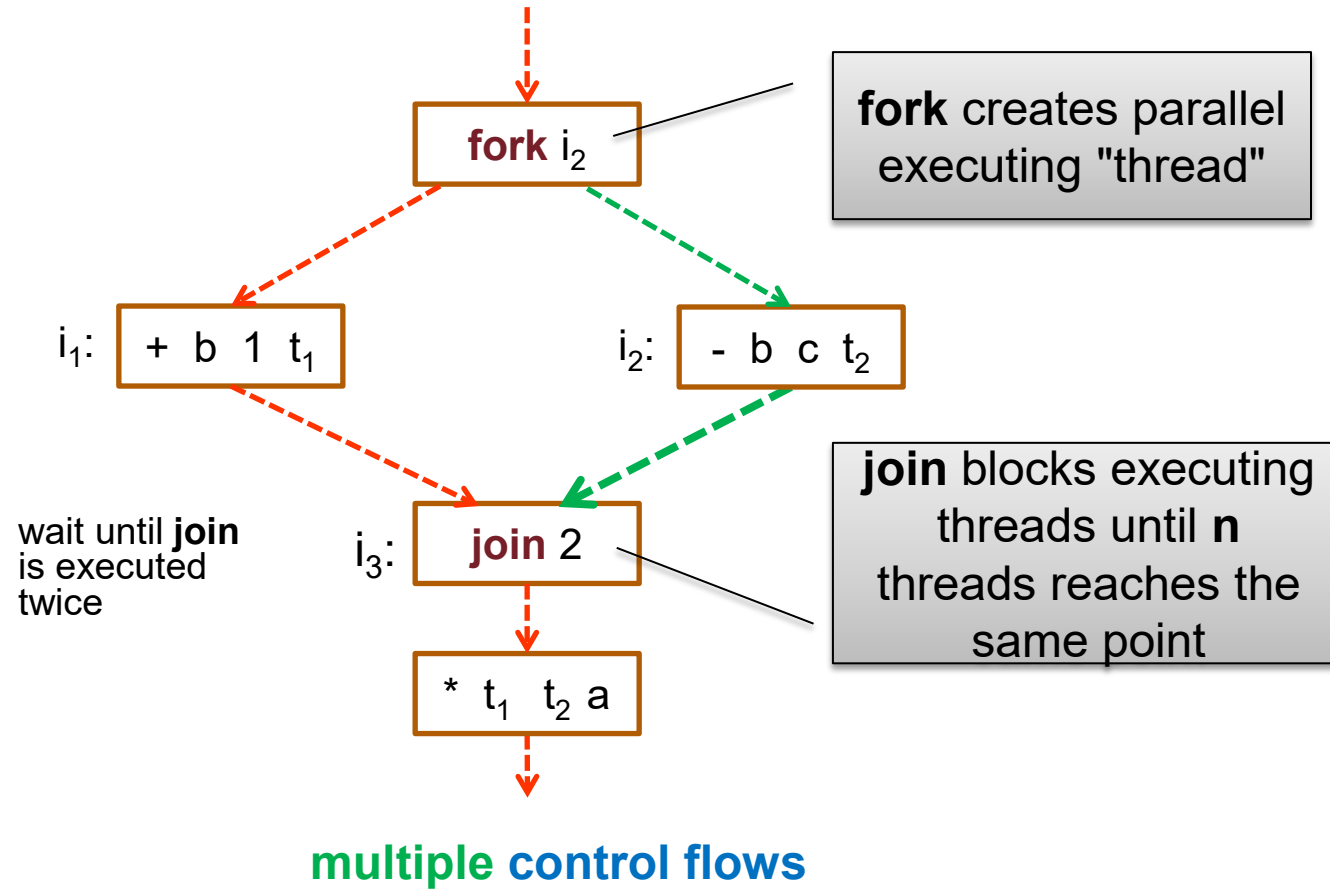
Parallel Computing

problem

instructions

$task_0$ → ... → $processing\ unit_0$

$task_1$ → ... → $processing\ unit_1$

....

$task_m$ → ... → $processing\ Unit_{p-1}$

i_n ... i_2 i_1

1. A problem is divided into **m** discrete parts (tasks) that can be solved concurrently
2. Each part is further broken down to a series of instructions (**i**)
3. **Instructions from each part execute in parallel on different processing units (p)**

Example: Evaluate $a = (b+1) * (b-c)$

fork i_2

fork creates parallel executing "thread"

i_1: `+ b 1 t_1`

i_2: `- b c t_2`

wait until **join** is executed twice

i_3: join 2

join blocks executing threads until **n** threads reaches the same point

`* t_1 t_2 a`

multiple control flows

Parallel Computing (Shared-Memory)

Why Parallel Computing

- Primary reasons:
    1. Overcome limits of serial computing
    2. Save (wall-clock) time
    3. Solve larger problems

- Other reasons:
    1. Take advantage of non-local resources
    2. Cost saving – use multiple "cheaper (commoditized)" computing resources
    3. Overcome memory constraints

- Performance - exploits a large collection of processing units that can communicate and cooperate to solve large problems fast

BASIC OF PARALLEL COMPUTING

Computational Model Attributes

[Operation Mechanism]

- The *primitive units* of computation or basic actions of the computer (data types and operations defined by the instruction set)

[Data Mechanism]

- The definition of *address spaces* available to the computation (how data are accessed and stored)

[Control Mechanism]

- The *schedulable units* of computation (rules for partitioning and scheduling the problem for computation using the primitive units)

[Communication Mechanism]

- The modes and patterns of **communication** among processing elements working in parallel, so that they can exchange needed information

[Synchronization Mechanism]

- Mechanism to ensure needed information arrives at the right time

Basics of Parallel Computing

Application **Problem**

↓ decompose

Tasks

↓ schedule & map

Physical Cores & **Processors**
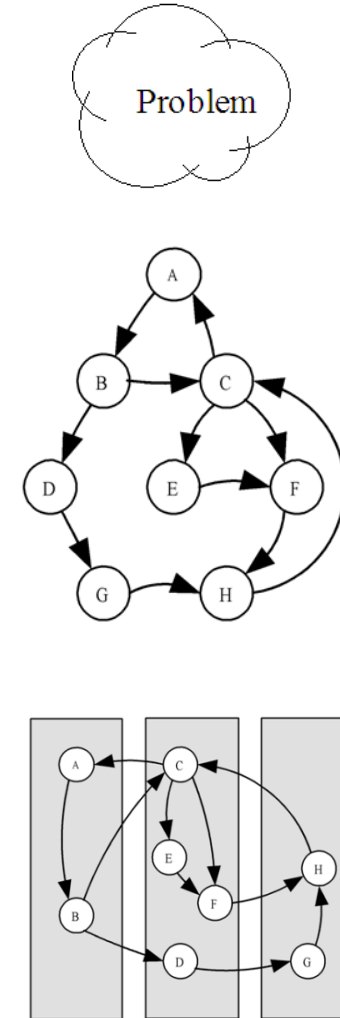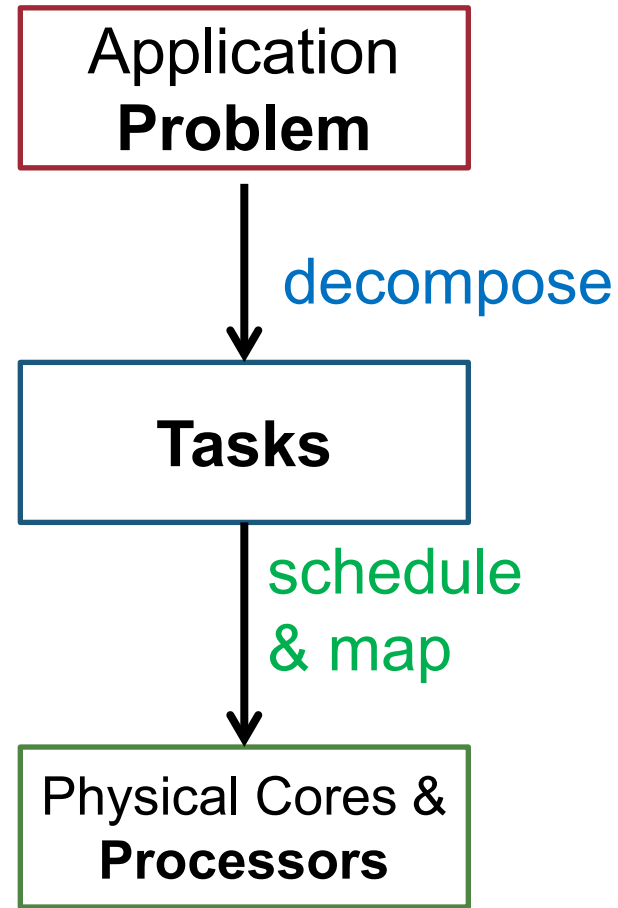
Decomposition

- One problem → many possible *decompositions*

- Complicated and laborious process
  - Potential parallelism of an application dictates how it should be split into **tasks**
  - Size of tasks is called *granularity* – can choose different task sizes
  - Defining the tasks is challenging and difficult to automate (**why?**)

Tasks

- Tasks are coded in a parallel programming language

- **Scheduling**: Assignment of tasks to processes or threads
  - Dictates the task execution order
  - Manually defined? Static? Dynamic?

- **Mapping:** Assignment of processes/threads to physical cores/processors for execution

- Tasks may depend on each other resulting in **data** or **control dependencies**
  - ➔ **impose execution order** of parallel tasks

Dependences and Coordination

- Dependences among tasks impose constraints on scheduling

- To execute correctly, processes and threads need *synchronization* and *coordination*

  ➔ depends on information exchanges between processes and threads and in turns, depends on the hardware memory organization

- Memory organizations: **shared-memory** (threads) and **distributed-memory** (processes)

Example – Serial Solution

- Find the sum for **n** numbers:

```
sum = 0;
for (i = 0; i < n; i++) {
    x = Compute_next_value(. . .);
    sum += x;
}
```

Serial Code

Example – Parallel Solution (1/3)

- Suppose we have **p cores** (p < n)
  - Each core can perform a partial sum of **n/p values**

```
my_sum = 0;
my_start = ……;
my_end = ……;
for (my_i = my_start; my_i < my_end; my_i++){
    my_x = Compute_next_value(. . .);
    my_sum += my_x;
}
```

  - Each core uses its own private variables and executes this block of code independently

Example – Parallel Solution (2/3)

- Upon completion, each core contains:
  - the partial sum in my_sum

- E.g. If n (=24), and p = 8
- Values generated:

1	4	3	9	2	8	5	1	1	5	3	7	2	5	0	4	1	8	6	5	1	2	3	9

- Partial sum calculated in each core:

Core	0	1	2	3	4	5	6	7
my_sum	8	19	7	15	7	13	12	14

Example – Parallel Solution (3/3)

- Next, a core designated as the master core adds up all values of my_sum to form the global sum (8 + 19 + 7 + 15 + 7 + 13 + 12 + 14 = 95)
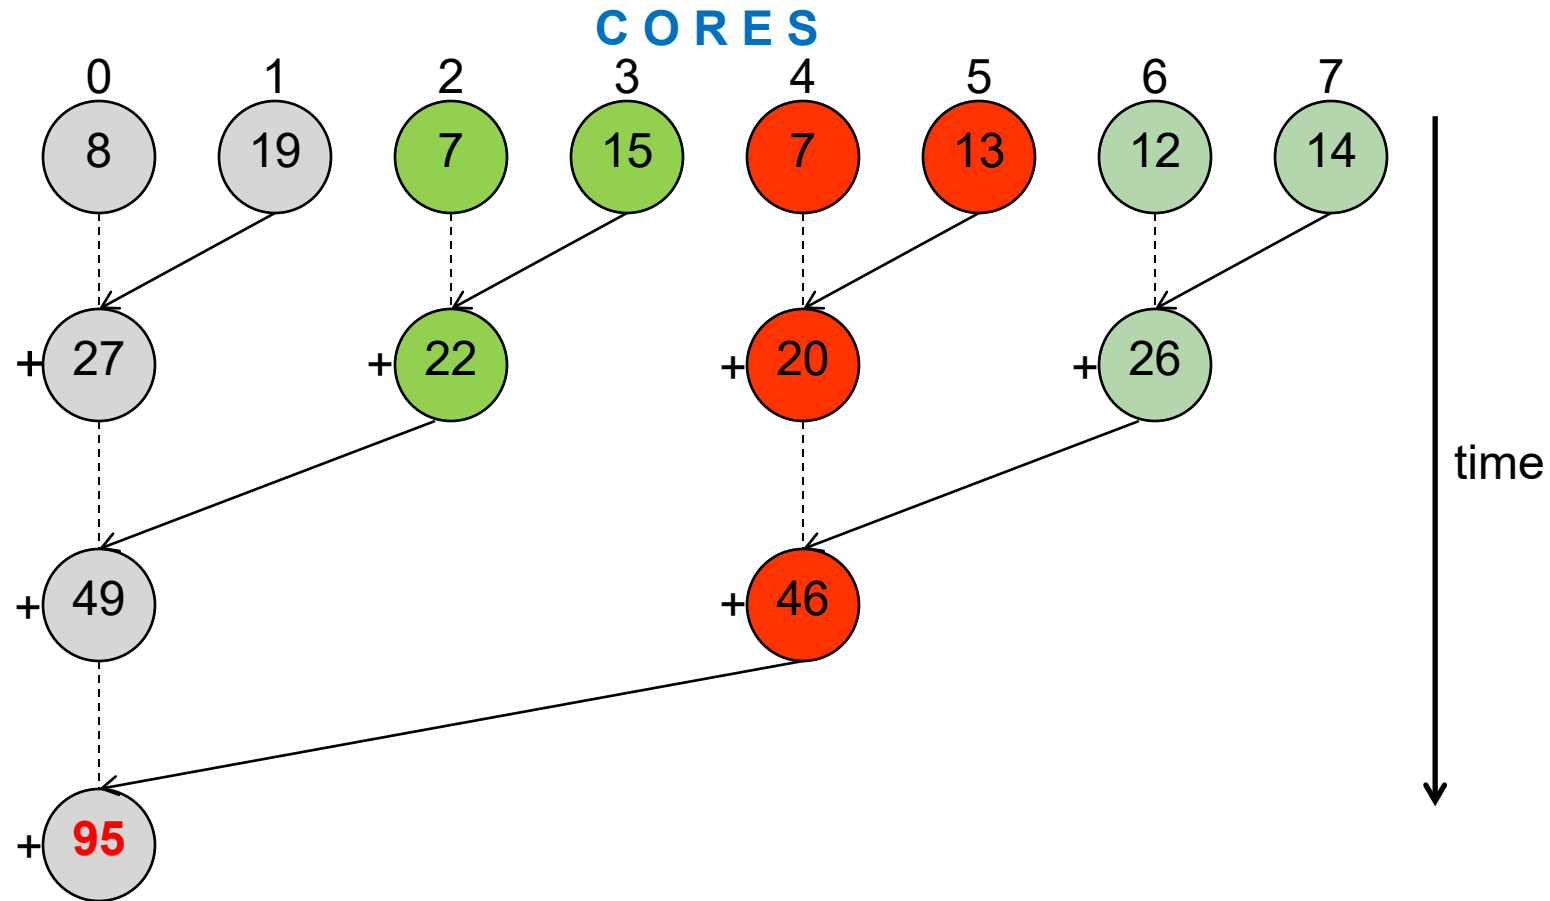
```
if (I'm the master core){
    sum = my_x;
    for each core other than myself {
        receive value from other core
        sum += value;
    }
} else {                    //not master core
    send my_x to the master;
}
```

Parallel Code

Core	0	1	2	3	4	5	6	7
my_sum	95	19	7	15	7	13	12	14

Better Parallel Algorithm

- Share the work of the global summing too:

Comparison: Two Parallel Solutions

- **8 cores**
  - 1st version = 7 ~~additions~~
  - 2nd version = 3 ~~additions~~

  - factor of 2 improvement

- **1000 cores**
  - 1st version = 999 ~~additions~~
  - 2nd version = 10 ~~additions~~

  - factor of 10 improvement

Parallel Performance

- Two perspectives:
  - Execution time versus Throughput

- Parallel execution time = computation time + parallelization overheads

- Parallelization overheads:
  - distribution of work (tasks) to processors
  - information exchange or synchronization
  - idle time
  - etc

Parallel Computing is HARD

> "…when we start talking about parallelism and ease of use of truly parallel computers, we're talking about a problem that's as hard as any that computer science has faced. …
>
> I would be panicked if I were in industry."
>
> John Hennessy, President, Stanford University, 2007

Compiler Challenges

- ## Heterogeneous processors
 - ❑ Increase in the design space for code optimization

- ## Auto-tuners:
  - ❑ Optimizing code at runtime

- ## Software-controlled memory management

Parallel Programming Challenges

- Finding enough parallelism (Amdahl's Law!)

- Granularity

- Locality

- Load balance

- Coordination and synchronization

- Debugging

- Performance modeling / Monitoring

LANDSCAPE OF PARALLEL COMPUTING

Classical Science

Nature

Observation

Physical
Experimentation

Theory
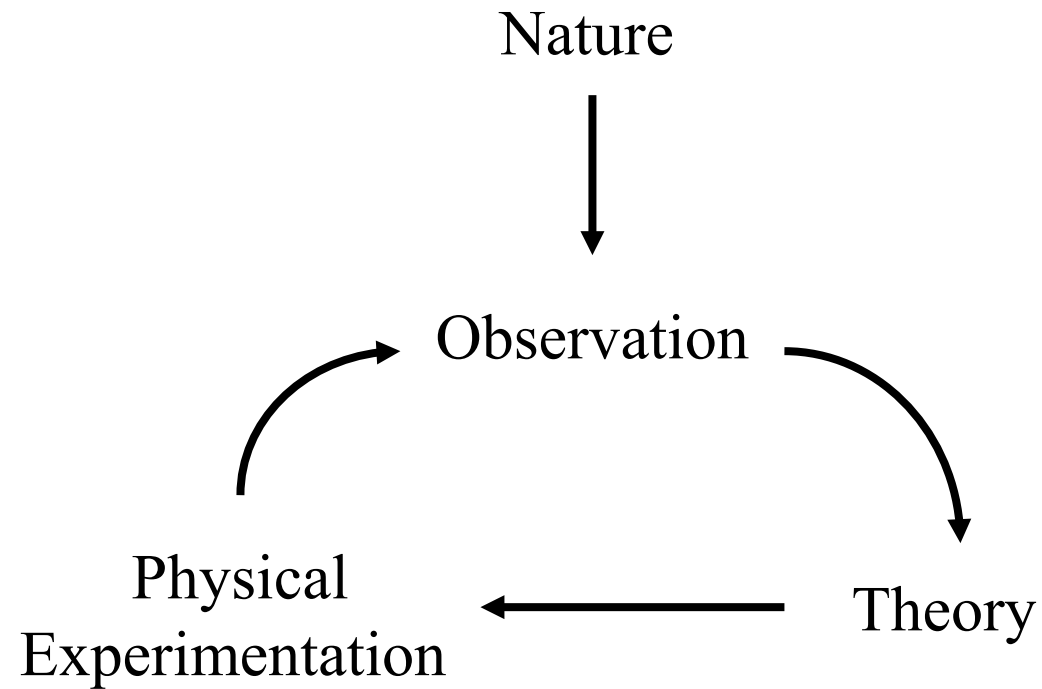
Modern Scientific Method

Science in the 21st century -
Parallel & Distributed Computing
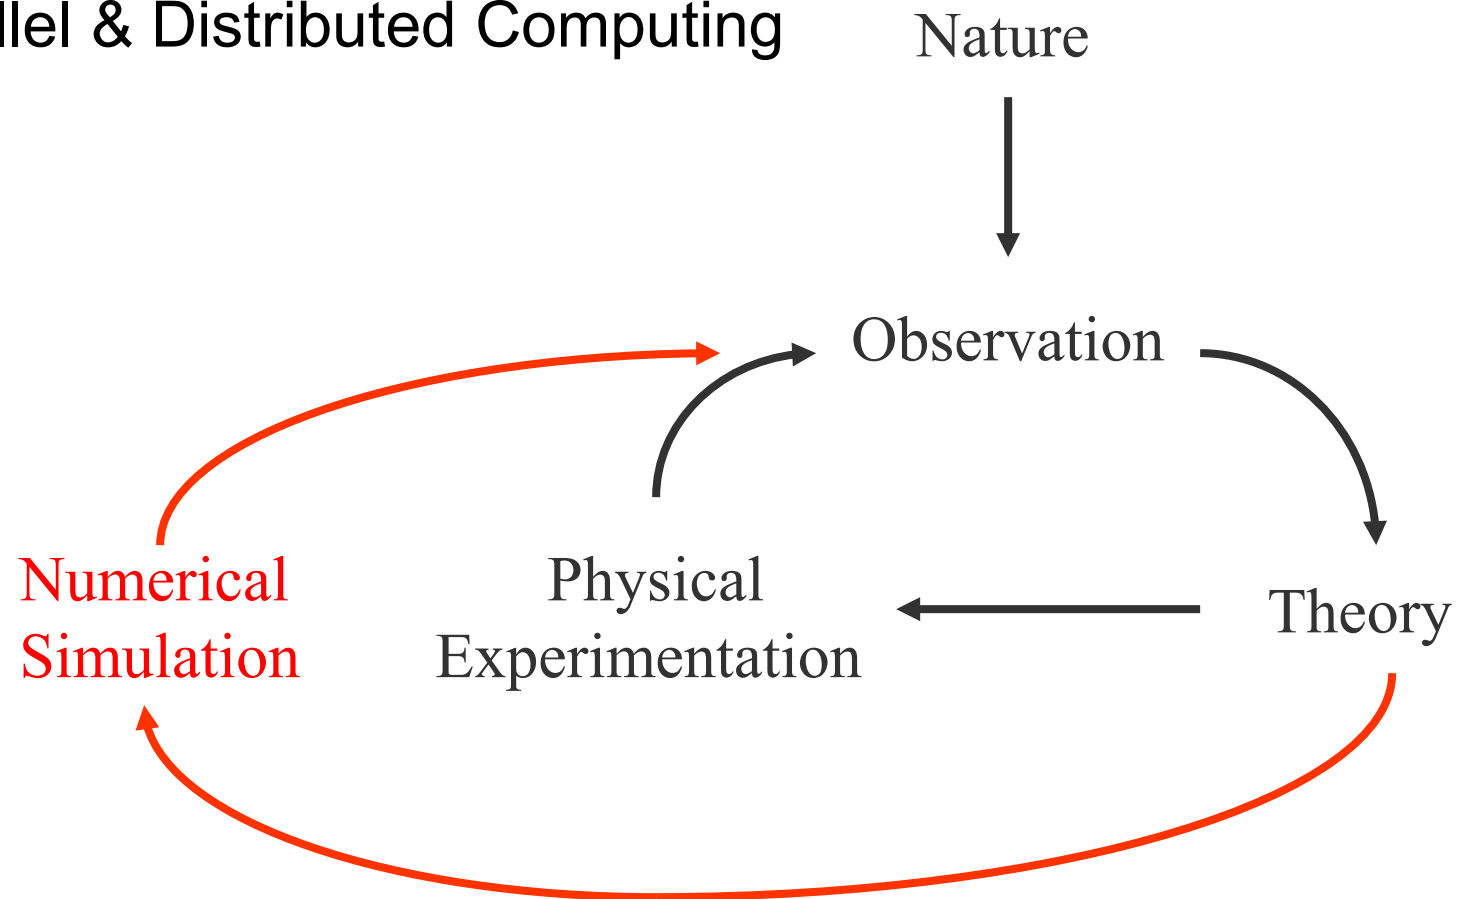
Uses of Parallel Computing

- Data mining, web search engines, web-based business services, pharmaceutical design, financial and economic modelling, etc.
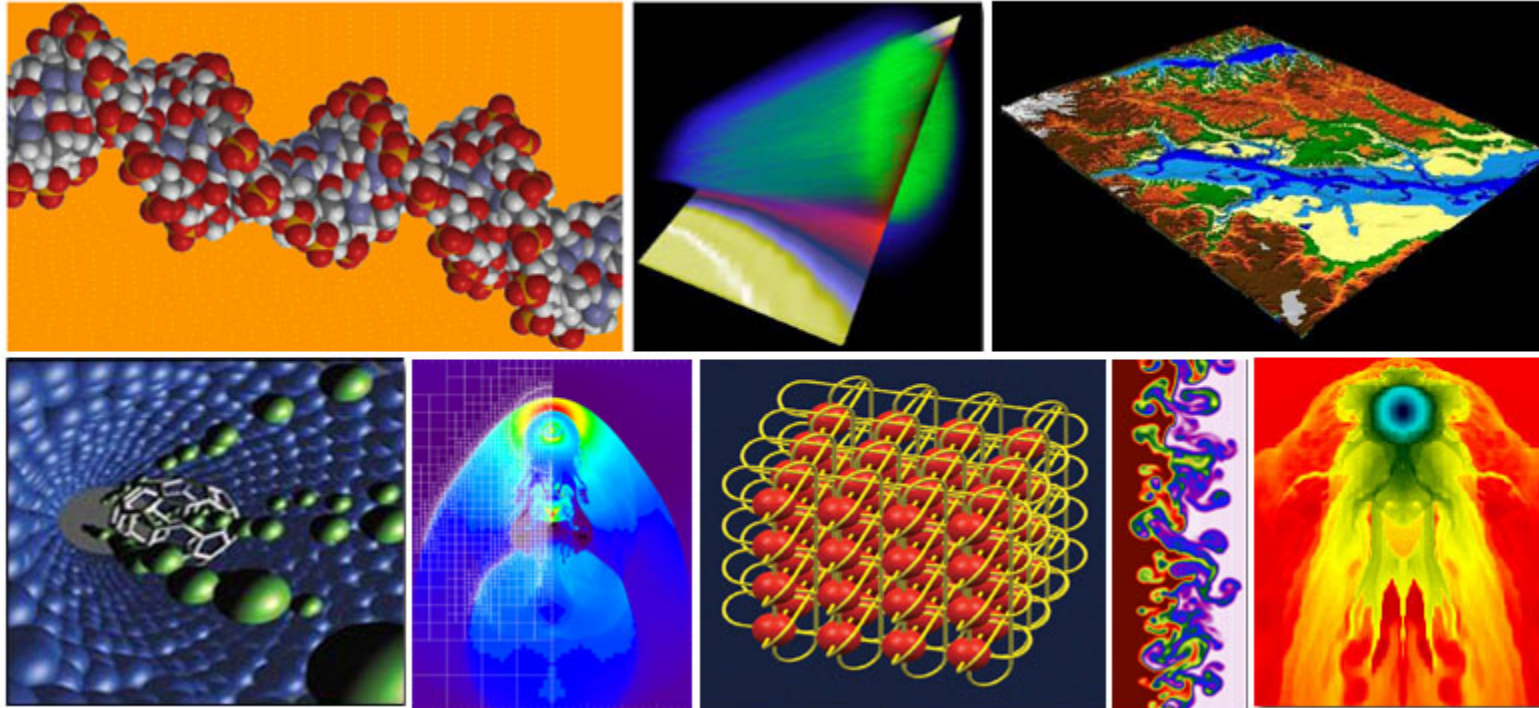
http://computing.llnl.gov/tutorials/parallel_comp

Types of Computer Systems

Computer Systems
- Single System
  - PC/Workstation
  - SMP/NUMA
  - Vector
  - Mainframe
- **Distributed Systems** (multiple systems)
  - Clusters
  - Client Server
  - Grids
  - Peer-to-Peer
  - Cloud

Control and Management

Centralised ←————————————————————→ Decentralized
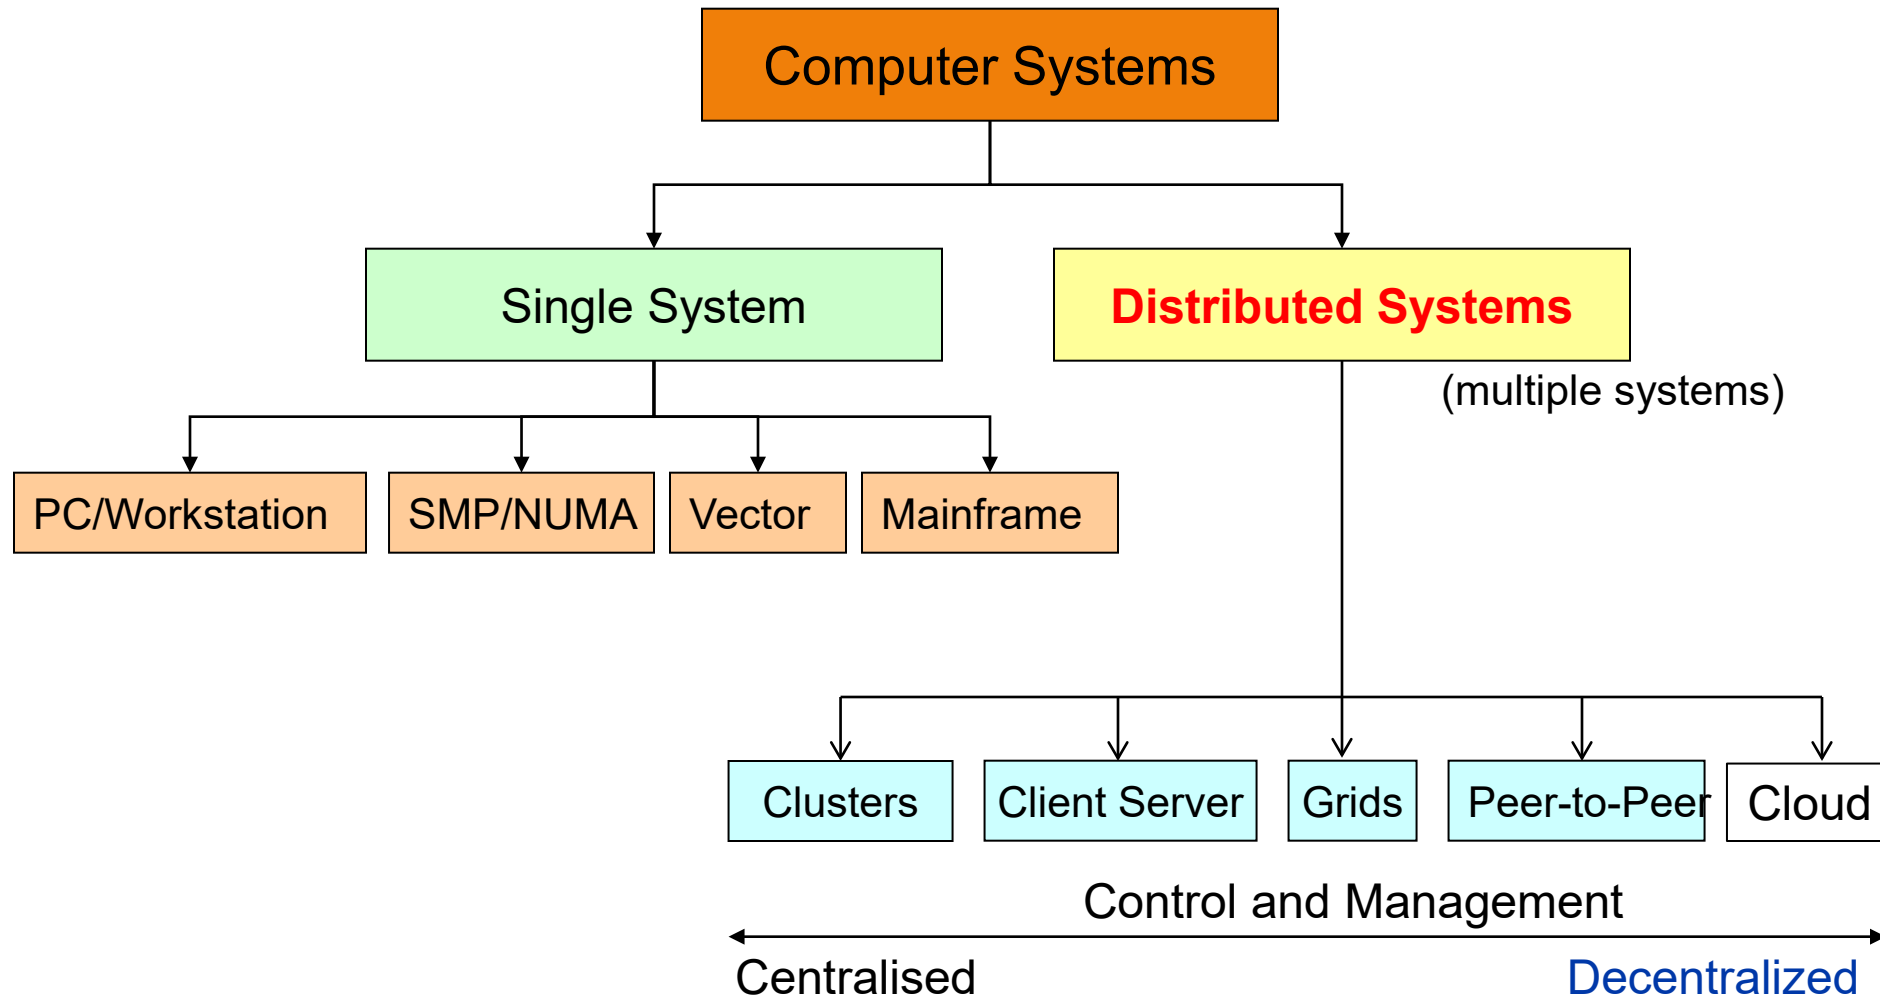
Parallel Computing "Food Chain"

Today
clusters, grids,
p2p, clouds, GPUs

mainframe

vector supercomputer

minicomputer

workstation

PC
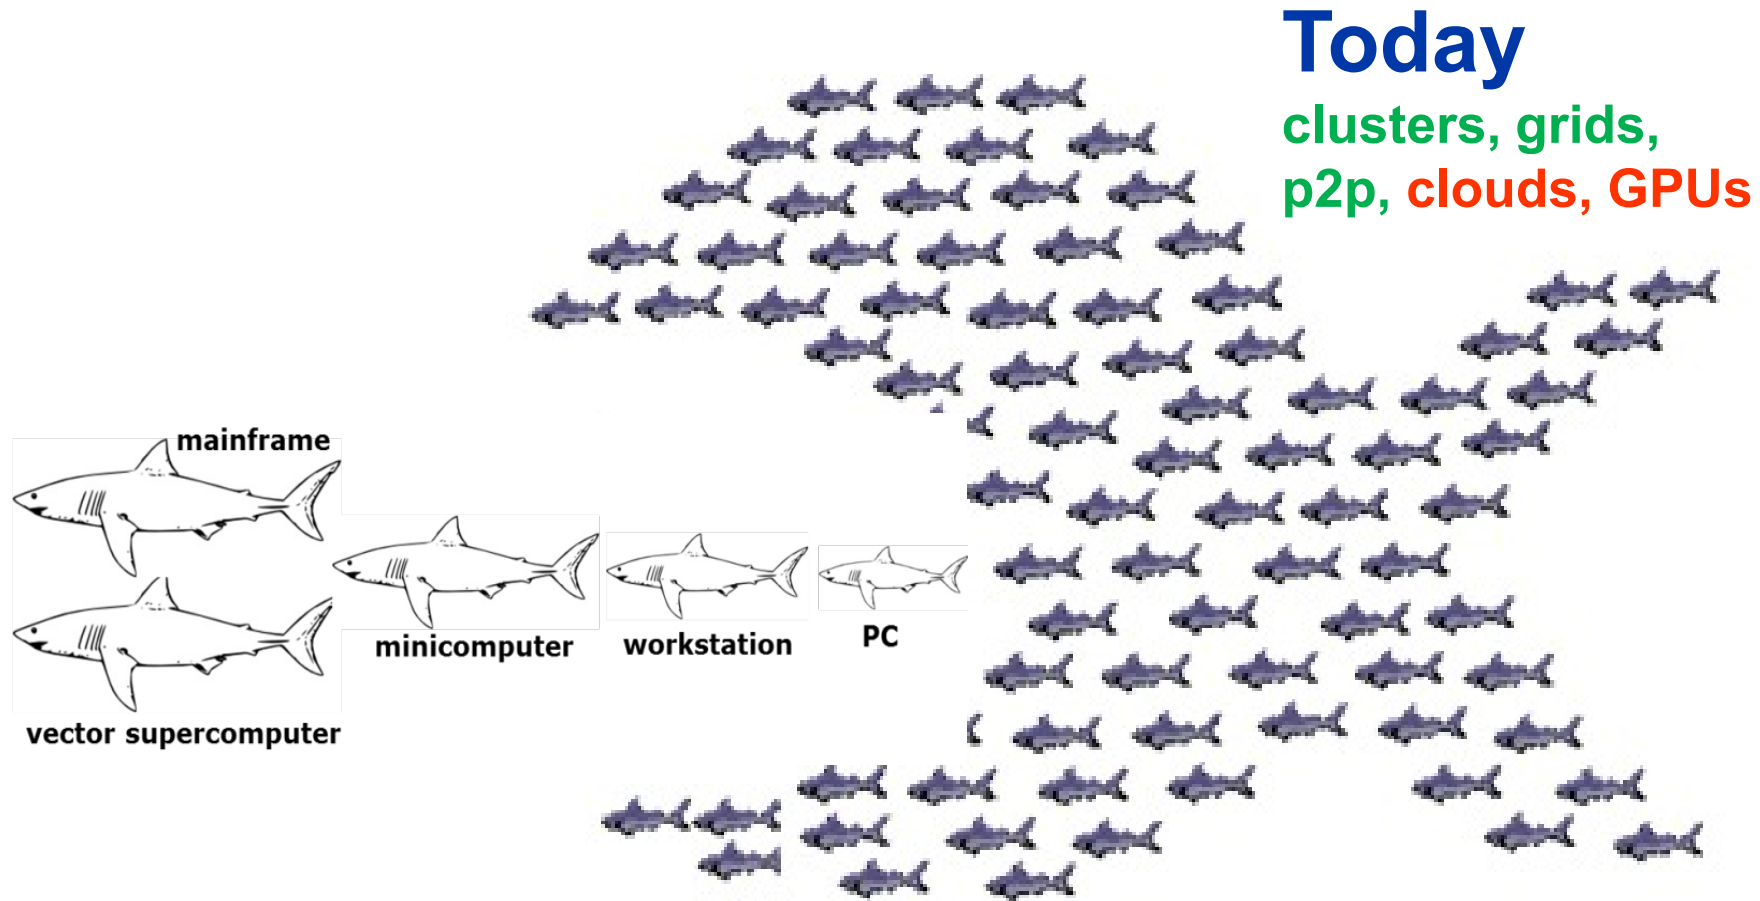
Parallel Computing and Parallel machines

In your pocket:
ARM big.LITTLE 64-bit
(Cortex-A53 + Cortex-A73)
4-8 cores, >2GHz, 128GB RAM

Intel Skylake architecture:
Quad-core CPU + multi-core
GPU integrated on one chip

2020 Fastest Supercomputer

Prefixes for representing orders of magnitude

Orders of magnitude (in base 10) are expressed using standard metric prefixes, which are abbreviated to single characters when prepended to other abbreviations, such as FLOPS and B (for byte):

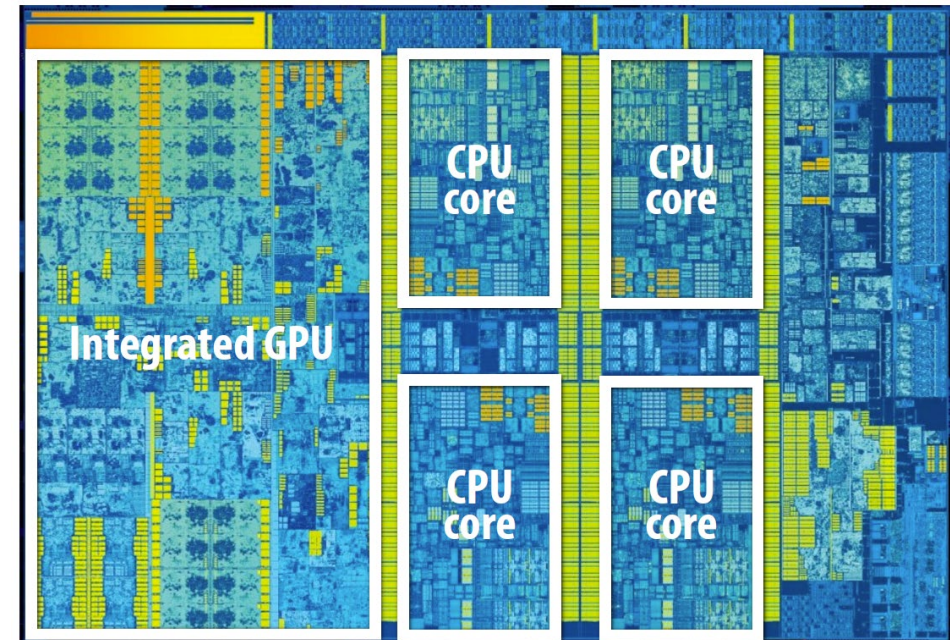| Prefix	Abbreviation	Order of magnitude (as a factor of 10)	Computer performance	Storage capacity
giga-	G	10^9	gigaFLOPS (GFLOPS)	gigabyte (GB)
tera-	T	10^{12}	teraFLOPS (TFLOPS)	terabyte (TB)
peta-	P	10^{15}	petaFLOPS (PFLOPS)	petabyte (PB)
exa-	E	10^{18}	exaFLOPS (EFLOPS)	exabyte (EB)
zetta-	Z	10^{21}	zettaFLOPS (ZFLOPS)	zettabyte (ZB)
yotta-	Y	10^{24}	yottaFLOPS (YFLOPS)	yottabyte (YB)

2020 Fastest Supercomputer

Supercomputer Fugaku

415 petaflops

~7300000 cores ARM A64FX

28 Mwatts

No. 9 in GREEN500

Video

2018-19 Fastest Supercomputer

Summit

148 petaflops
4356*22*2 cores IBM Power 9
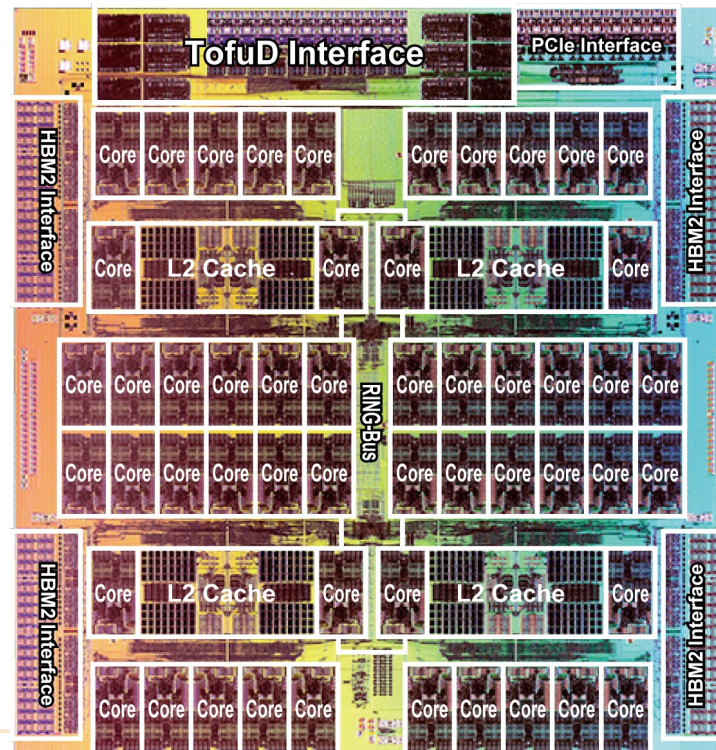**4356*6 NVIDIA Tesla V100 GPUs**
10.0 Mwatts
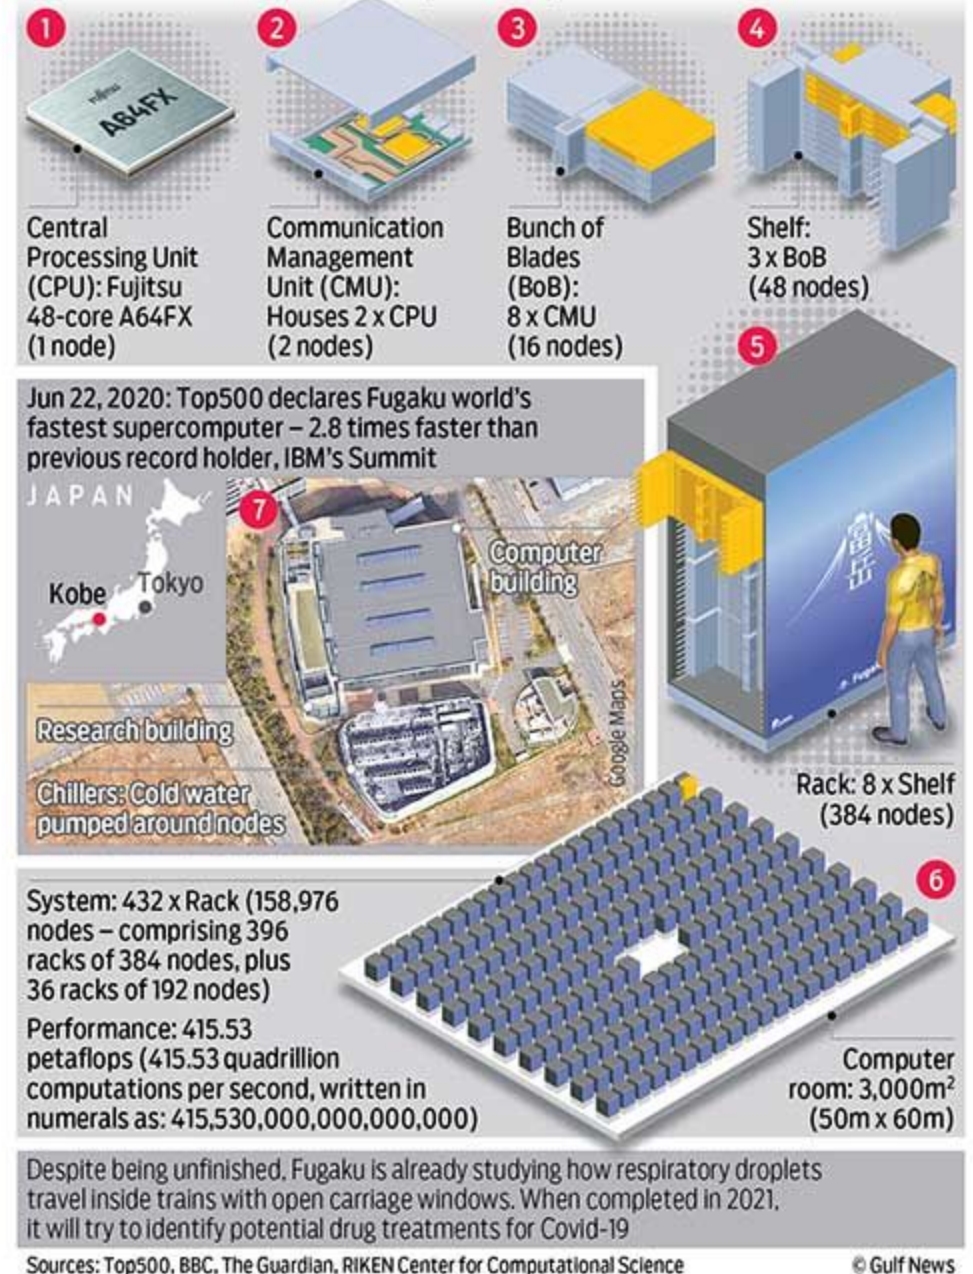**No. 8 in GREEN500**

Video

Deconstructing Fugaku

Processor (CPU):

- Fujitsu's 48-core A64FX SoC (optional assistant cores)

- 158,976 nodes (CPUs):
    396 racks of 384 nodes
    36 racks of 192 nodes
- Fujitsu Tofu (Torus Fusion) interconnect

Supercomputer to seek Covid-19 cure

The world's fastest supercomputer, Japan's $1.2 billion Fugaku, is to use its enormous power to try to identify treatments for Covid-19

① Central Processing Unit (CPU): Fujitsu 48-core A64FX (1 node)

② Communication Management Unit (CMU): Houses 2 x CPU (2 nodes)

③ Bunch of Blades (BoB): 8 x CMU (16 nodes)

④ Shelf: 3 x BoB (48 nodes)

Jun 22, 2020: Top500 declares Fugaku world's fastest supercomputer – 2.8 times faster than previous record holder, IBM's Summit

JAPAN — Kobe, Tokyo

Research building

Chillers: Cold water pumped around nodes

Computer building

⑤ Rack: 8 x Shelf (384 nodes)

System: 432 x Rack (158,976 nodes – comprising 396 racks of 384 nodes, plus 36 racks of 192 nodes)

Performance: 415.53 petaflops (415.53 quadrillion computations per second, written in numerals as: 415,530,000,000,000,000)

⑥ Computer room: 3,000m² (50m x 60m)

Despite being unfinished, Fugaku is already studying how respiratory droplets travel inside trains with open carriage windows. When completed in 2021, it will try to identify potential drug treatments for Covid-19

Sources: Top500, BBC, The Guardian, RIKEN Center for Computational Science

© Gulf News

Usage of Fugaku

- **Social and scientific priority issues**
  - Innovative drug discovery
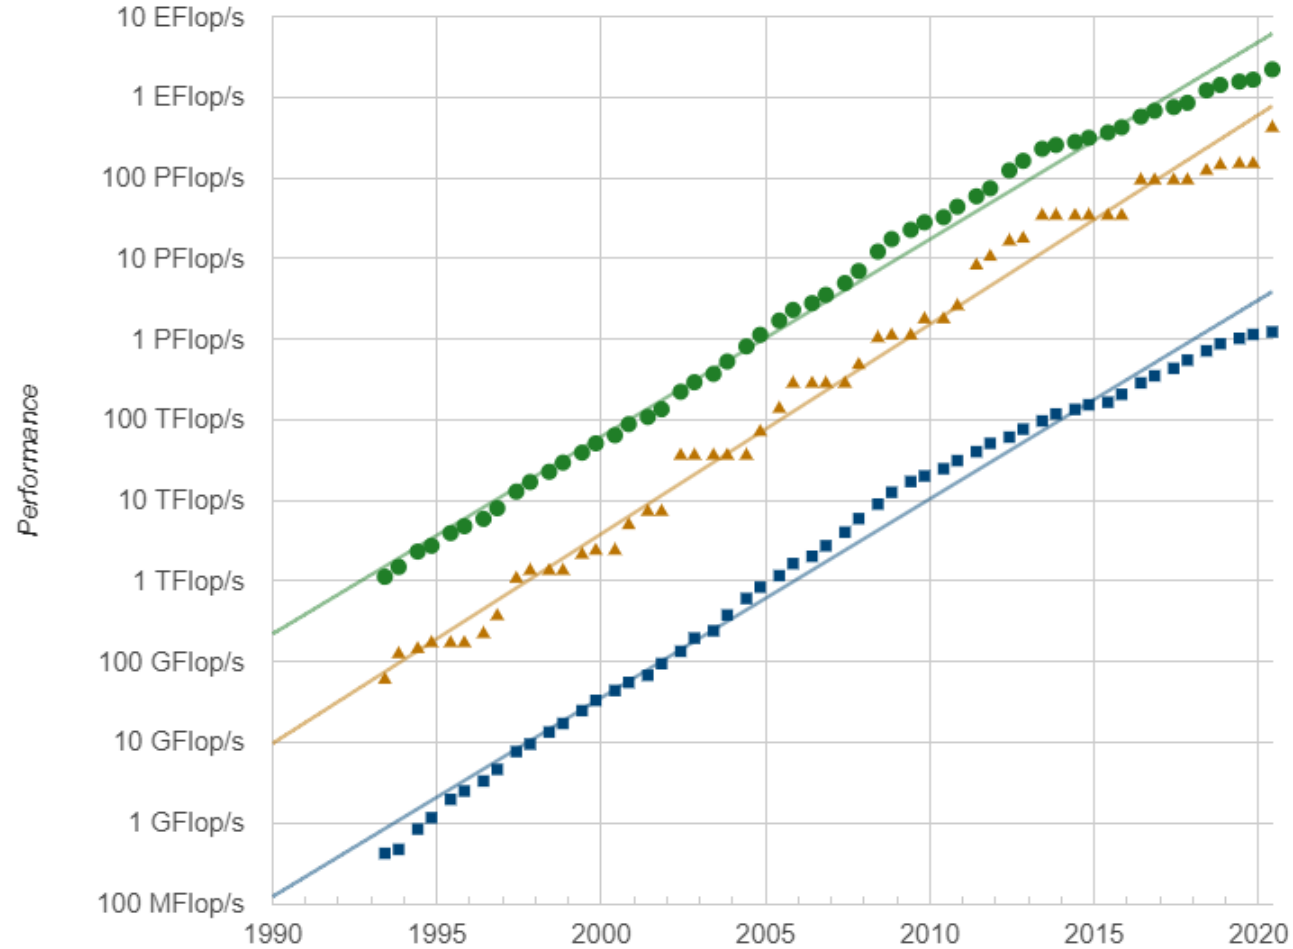 - Personalized and preventive medicine
- **Disaster prevention and global climate problems**
  - Meteorological and global environmental predictions
- **Innovative clean energy systems**
- **Evolution of the universe**

Supercomputer: Architecture (2018)
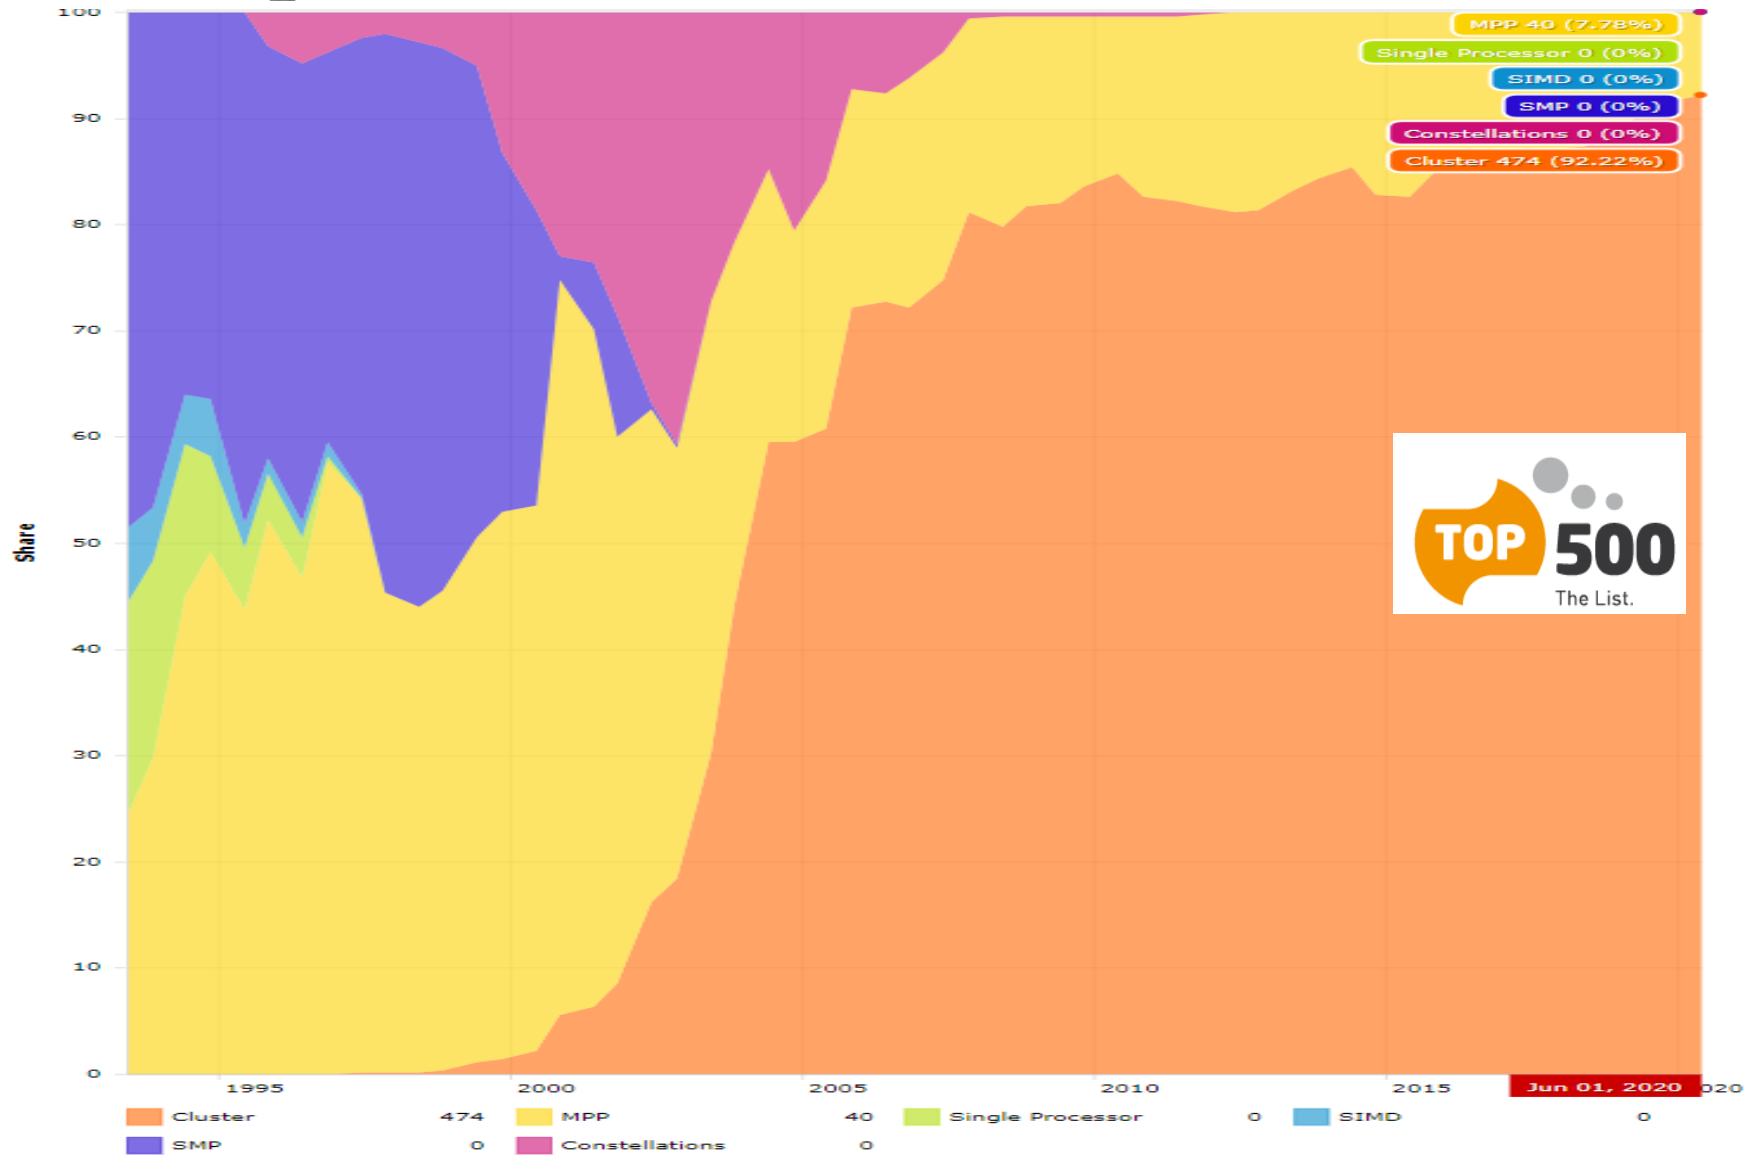
TOP500 Performance over Time

Projected Performance Development

source: top500.org

name	prefix	multiplier
exa	E	10^{18}
peta	P	10^{15}
tera	T	10^{12}
giga	G	10^{9}
mega	M	10^{6}
kilo	K	10^{3}

Summary

- Serial versus parallel computing

- Parallel computing - what, why and how

- Performance and challenges

- Landscape of parallel computing

Evolution of Prominent Computers

year	location	name	first
1834	Cambridge	Difference engine	Programmable computer
1943	Bletchley	Colossus	Electronic computer
1948	Manchester	SSEM (Baby)	Stored-program computer
1951	MIT	Whirlwind 1	Real-time I/O computer
1953	Manchester	Transistor computer	Transistorized computer
1971	California	Intel 4004	Mass-market CPU & IC
1979	Cambridge	Sinclair ZX-79	Mass-market home computer
1981	New York	IBM PC	Personal computer
1987	Cambridge	Acorn A400 series	High-street RISC PC sales
1990	New York	IBM RS6000	Superscalar RISC processor
1998	California	Sun picoIJAVA	Computer Based on a language

Progression of Computation Speed

year	Floating point operations per second (FLOPS)
1941	1
1945	100
1949	**1,000 (1 KiloFLOPS, kFLOPS)**
1951	10,000
1961	100,000
1964	**1,000,000 (1 MegaFLOPS, MFLOPS)**
1968	10,000,000
1975	100,000,000
1987	**1,000,000,000 (1 GigaFLOPS, GFLOPS)**
1992	10,000,000,000
1993	100,000,000,000
1997	**1,000,000,000,000 (1 TeraFLOPS, TFLOPS)**
2000	10,000,000,000,000
2007	478,000,000,000,000 (478 TFLOPS)
2009	**1,100,000,000,000,000 (1.1 PetaFLOPS)**

Computing Units

International System of Units (SI)

name	prefix	multiplier
exa	E	10^{18}
peta	P	10^{15}
tera	T	10^{12}
giga	G	10^{9}
mega	M	10^{6}
kilo	K	10^{3}
milli	m	10^{-3}
micro	μ	10^{-6}
nano	n	10^{-9}
pico	p	10^{-12}

International Electrotechnical Commission (IEC)

Name	prefix	multiplier
Exbi	Ei	2^{60}
pebi	Pi	2^{50}
tebi	Ti	2^{40}
gibi	Gi	2^{30}
mebi	Mi	2^{20}
kibi	Ki	2^{10}

Example of disk capacity:
20 Mibytes
20 MiB
20 Mebibytes
= 20 x 2^{20} = 20,971,520 bytes

Reference & Readings

- **Fugaku technical details:**
    - https://www.r-ccs.riken.jp/en/fugaku/project/outline

- **Summit technical details:**
    - https://www.tomshardware.com/news/summit-supercomputer-nvidia-ibm-power9-volta,35962.html

- Reading
    - Main reference book: chapter 1
    - A View of the Parallel Computing Landscape, CACM Vol. 52 No. 10, pp. 56-67, October 2009. [Technical Report: The Landscape of Parallel Computing Research: A View from Berkeley, Dec 2006]