


# Processes, Threads, and Synchronization

---

## Lecture 02

# Program Parallelization: Steps


- 3 main steps:
  1. **Decomposition** of the computations
  2. **Scheduling** (assignment of **tasks** to processes (or threads))
  3. **Mapping** of processes (or threads) to physical processors (or cores)



# Program Parallelization: Steps

## ■ 3 main steps:

1. **Decomposition** of the computations
2. **Scheduling** (assignment of **tasks** to processes (or threads))
3. **Mapping** of processes (or threads) to physical processors (or cores)



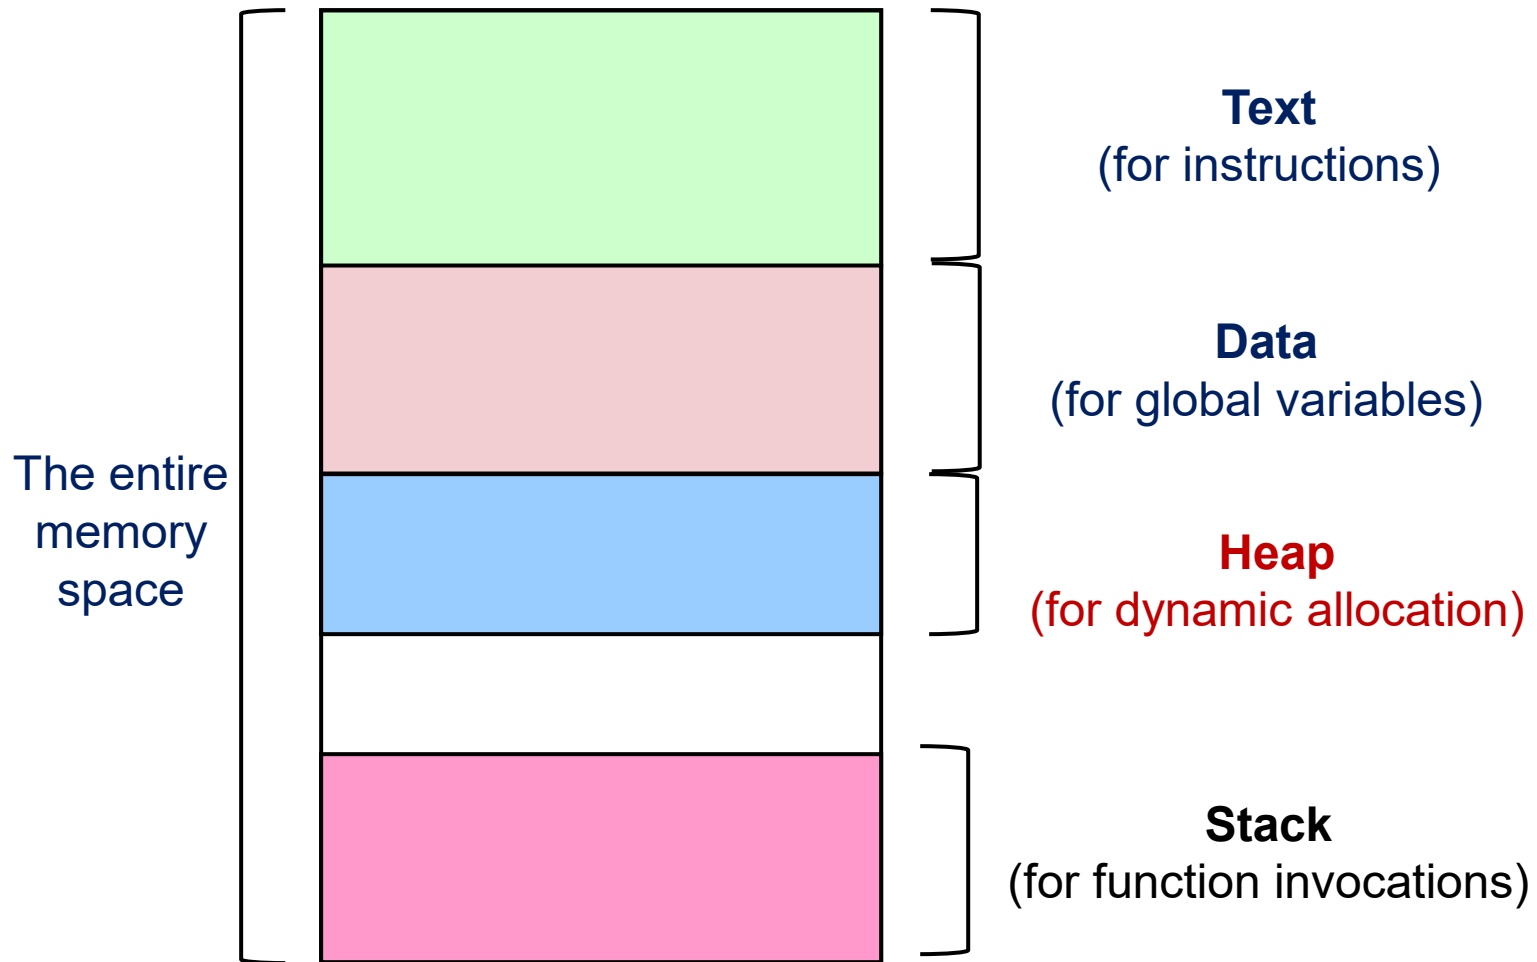
Abstractions of flow of control

# PROCESSES AND THREADS

# Processes

- A program in execution
  - Identified by PID (process ID)
  - Comprises:
    - executable program (PC),
    - global data
      - OS resources: open files, network connections
    - stack or heap
    - current values of the registers (GPRs and Special)
  - Own address space → exclusive access to its data
  - Two or more processes exchange data → need explicit communication

# Memory Illustration of a Process



# Multi-Programming (Multitasking)

- Several processes at different stages of execution
  - Need **context switch**, i.e., switching between processes
  - states of the suspended process must be saved → overhead
  - 2 types of execution:
    - Time slicing execution – pseudo-parallelism
    - Parallel execution of processes on different resources

# Create a New Process in Unix

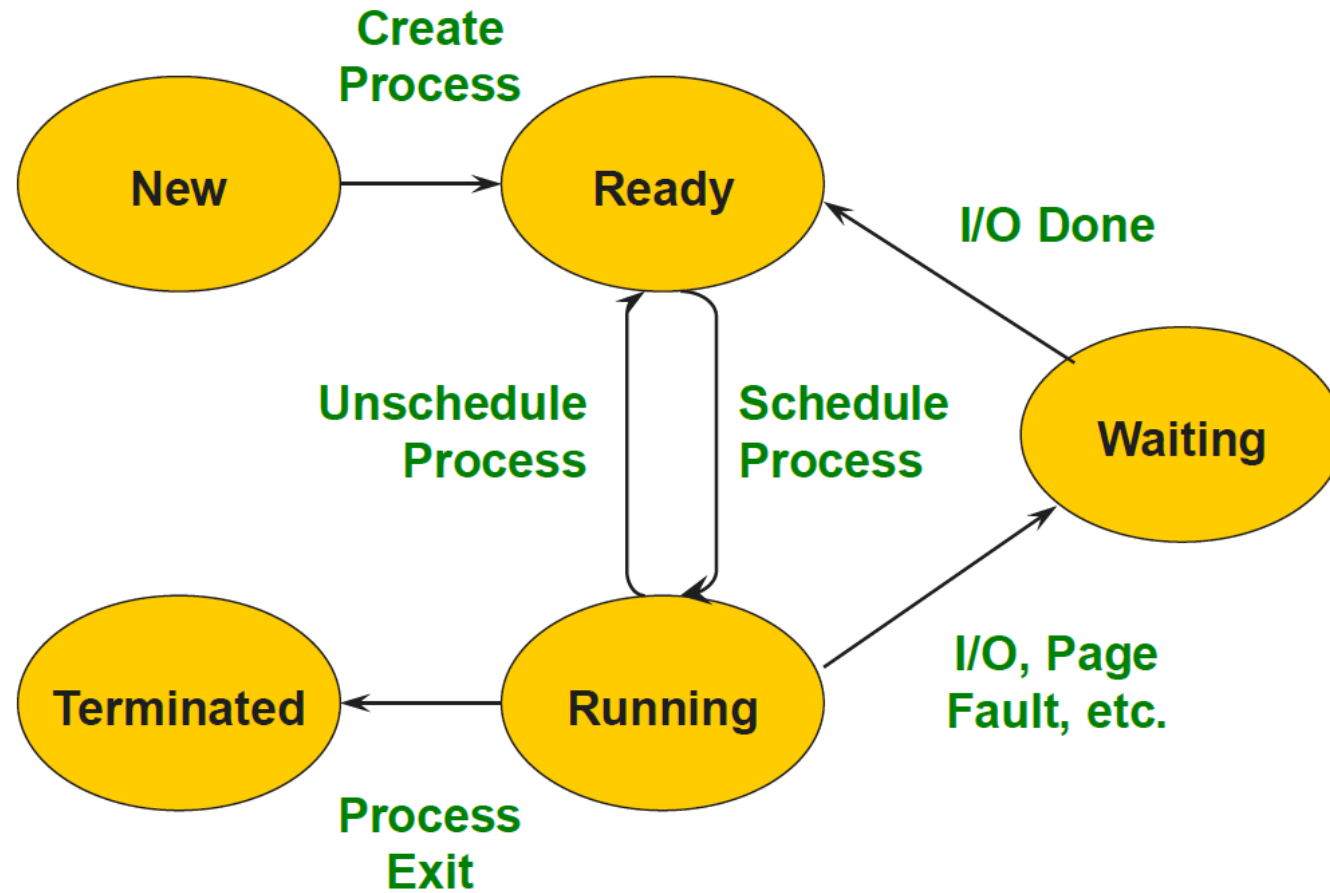
- Process  $P_1$  can create a new process  $P_2$ 
  - **fork** system call
  - **int exec(char \*prog, char \*argv[])**
- $P_2$  is an identical copy of  $P_1$  at the time of the fork call
  - $P_2$  works on a **copy** of the address space of  $P_1$
  - $P_2$  executes the same program as  $P_1$ , starting with the instruction following the fork call
- $P_2$  gets its own process number
  - Use **ps** or **top** in Unix console to see a list of processes
- $P_2$  can execute different statements as  $P_1$



# Fork()

```
int main(int argc, char *argv[])
{
    char *name = argv[0];
    int child_pid = fork();
    if (child_pid == 0) {
        printf("Child of %s is %d\n", name, getpid());
        return 0;
    } else {
        printf("My child is %d\n", child_pid);
        return 0;
    }
}
```

# Process State Graph



# Why fork()?

- Very useful when the child...
  - ❑ Is cooperating with the parent
  - ❑ Relies upon the parent's data to accomplish its task
- Example: web server

```
while(1) {  
    int sock = accept(); if ((child_pid = fork()) == 0) {  
        Handle client request  
    } else {  
        Close socket  
    }  
}
```

# Process Termination

- Use **exit(status)** in the child process
- Wait for a process in the parent process
  - ❑ **wait**
  - ❑ **waitpid (pid)**

# Inter-process Communication (IPC)

- Cooperating processes have to share information
  - Shared memory
    - Need to protect access when reading/writing with locks
  - Message passing
    - Blocking & non-blocking
    - Synchronous & asynchronous
  - Unix specific:
    - Pipes & Signal

# Process Interaction with OS

## Exceptions

- Executing a **machine level instruction** can cause exception
- For example: Overflow, Underflow, Division by Zero, Illegal memory address, Mis-aligned memory access
- **Synchronous**
  - Occur due to program execution
- Have to execute a **exception handler**

## Interrupts

- External events can interrupt the execution of a program
- Usually hardware related: Timer, Mouse Movement, Keyboard Pressed etc
- **Asynchronous**
  - Occur **independently** of program execution
- Have to execute an **interrupt handler**

# Disadvantages of Processes

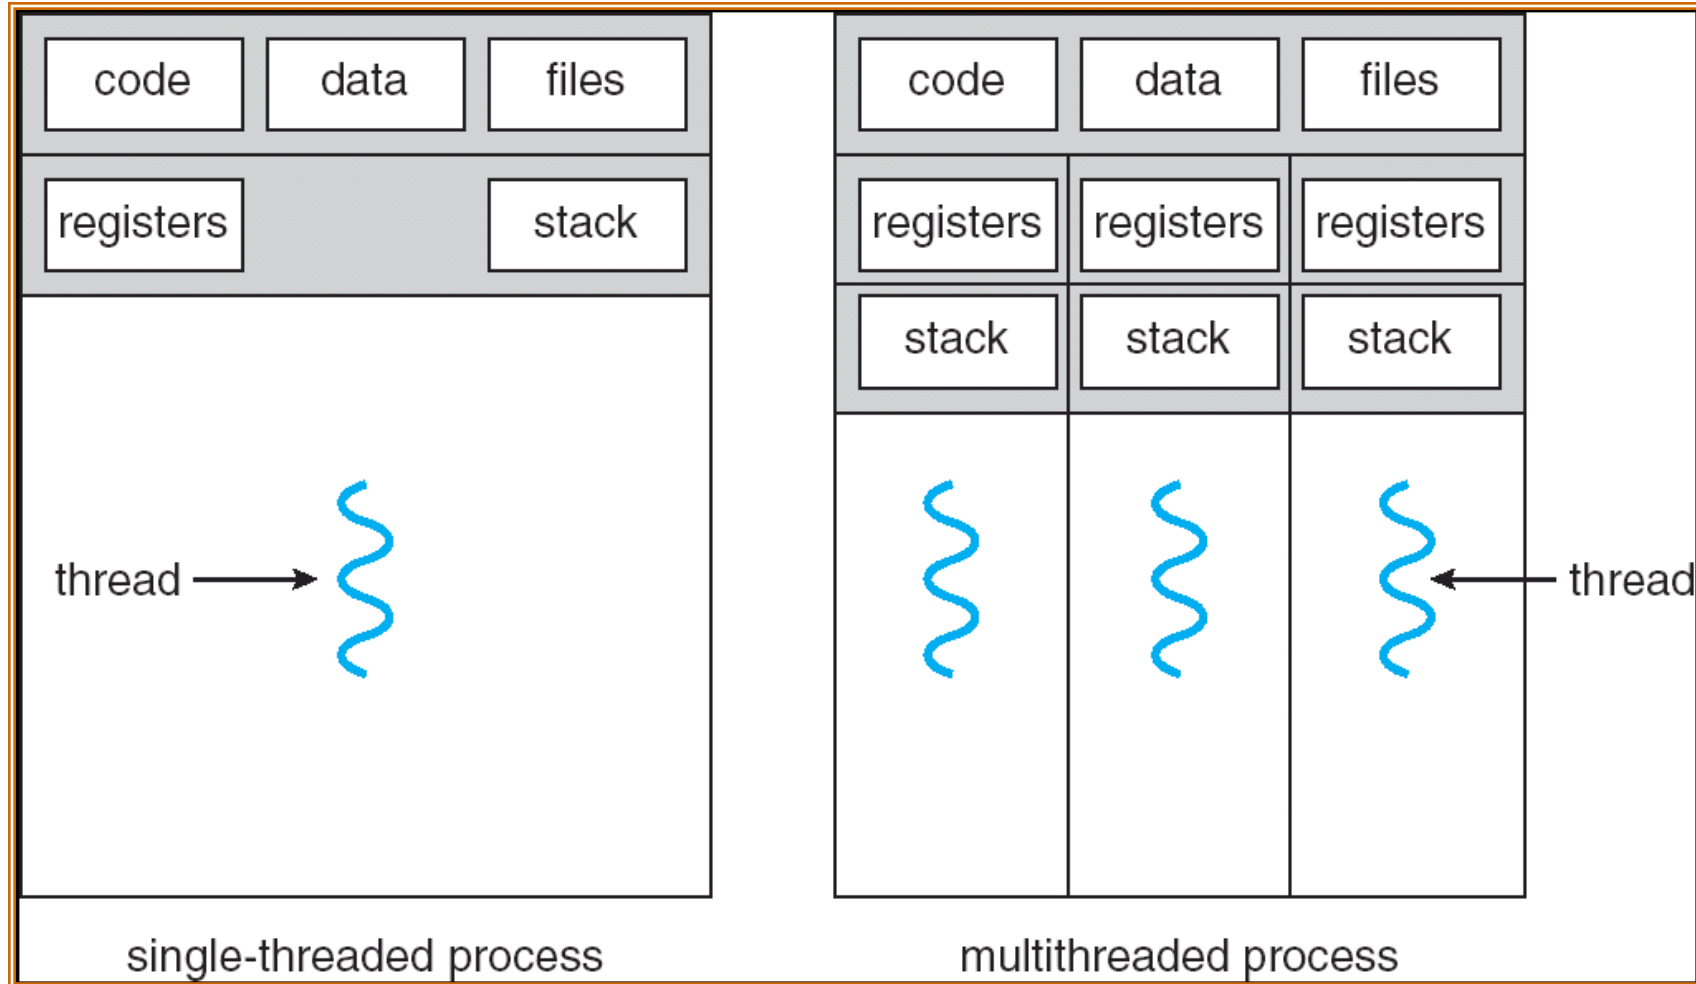
- Creating a new process is costly
  - ❑ Overhead of system calls
  - ❑ All data structures must be allocated, initialized and copied
- Communicating between processes costly
  - ❑ Communication goes through the OS

# Threads

- Extension of process model:
  - A process may consist of multiple independent control flows called **threads**
  - The thread defines a sequential execution stream within a process(PC, SP, registers)
- Threads share the address space of the process:
  - All threads belonging to the same process see the same value → **shared-memory architecture**



# Process and thread: Illustration



- Taken from Operating System Concepts (7<sup>th</sup> Edition) by Silberschatz, Galvin & Gagne, published by Wiley

# Threads (cont)

- Thread generation is faster than process generation
  - No copy of the address space is necessary
- Different threads of a process can be assigned run on different cores of a multicore processor
- 2 types of threads
  - User-level threads
  - Kernel threads

# User-Level Threads

- Managed by a thread library – OS unaware of user-level threads so no OS support
- Advantages – switching thread context is fast
- Disadvantages
  - ❑ OS cannot map different threads of the same process to different execution resources → no parallelism
  - ❑ OS cannot switch to another thread if one thread executes a blocking I/O operation

# Kernel Threads

- OS is aware of the existence of threads and can react correspondingly
- Avoid disadvantages of user-level threads
- Efficient use of the cores in a multicore system

# Mapping User-level to Kernel Threads

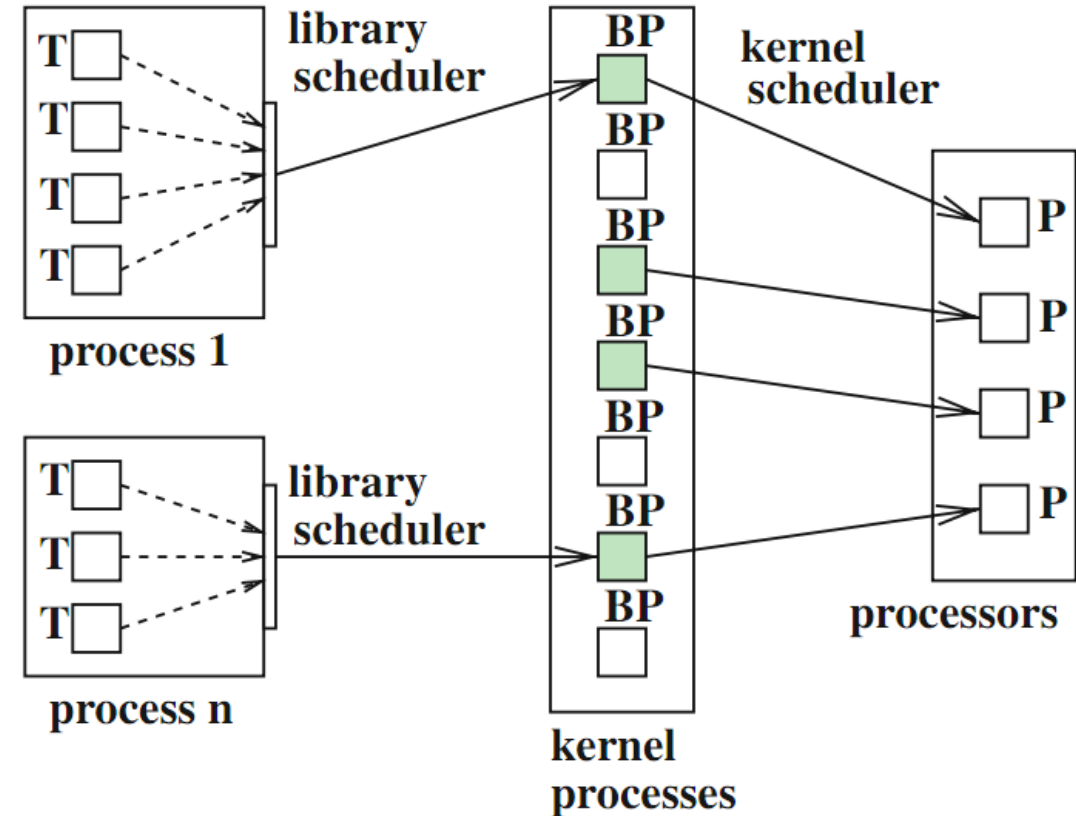
- User-level threads - Many-to-one mapping

- Kernel threads

- One-to-one mapping
- Many-to-many mapping

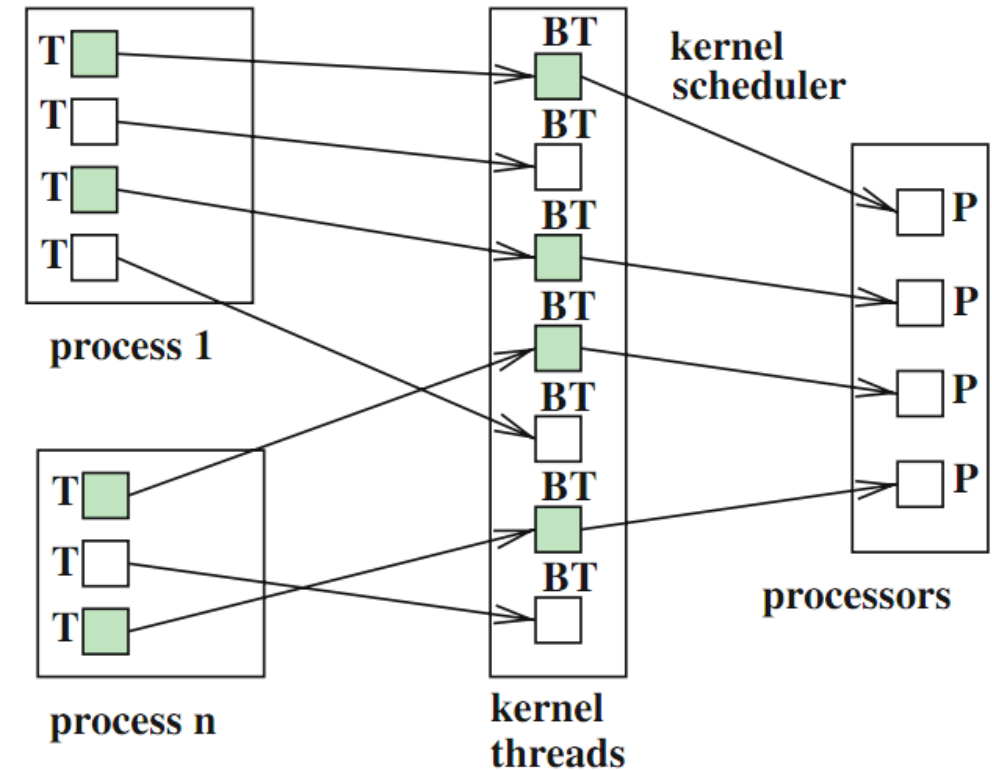
# Many-to-One Mapping

- All user-level threads are mapped to one process
- Thread library is responsible for the scheduling of user-level threads




# One-to-One Mapping

- Each user-level thread is assigned to exactly one kernel thread - no library scheduler needed
- OS is responsible for the scheduling and mapping of kernel threads



# Many-to-Many Mapping

- Library scheduler assigns the user-level threads to a given set of kernel threads
- Kernel scheduler maps the kernel threads to the available execution resources
- At different points in time, a user thread may be mapped to a different kernel thread



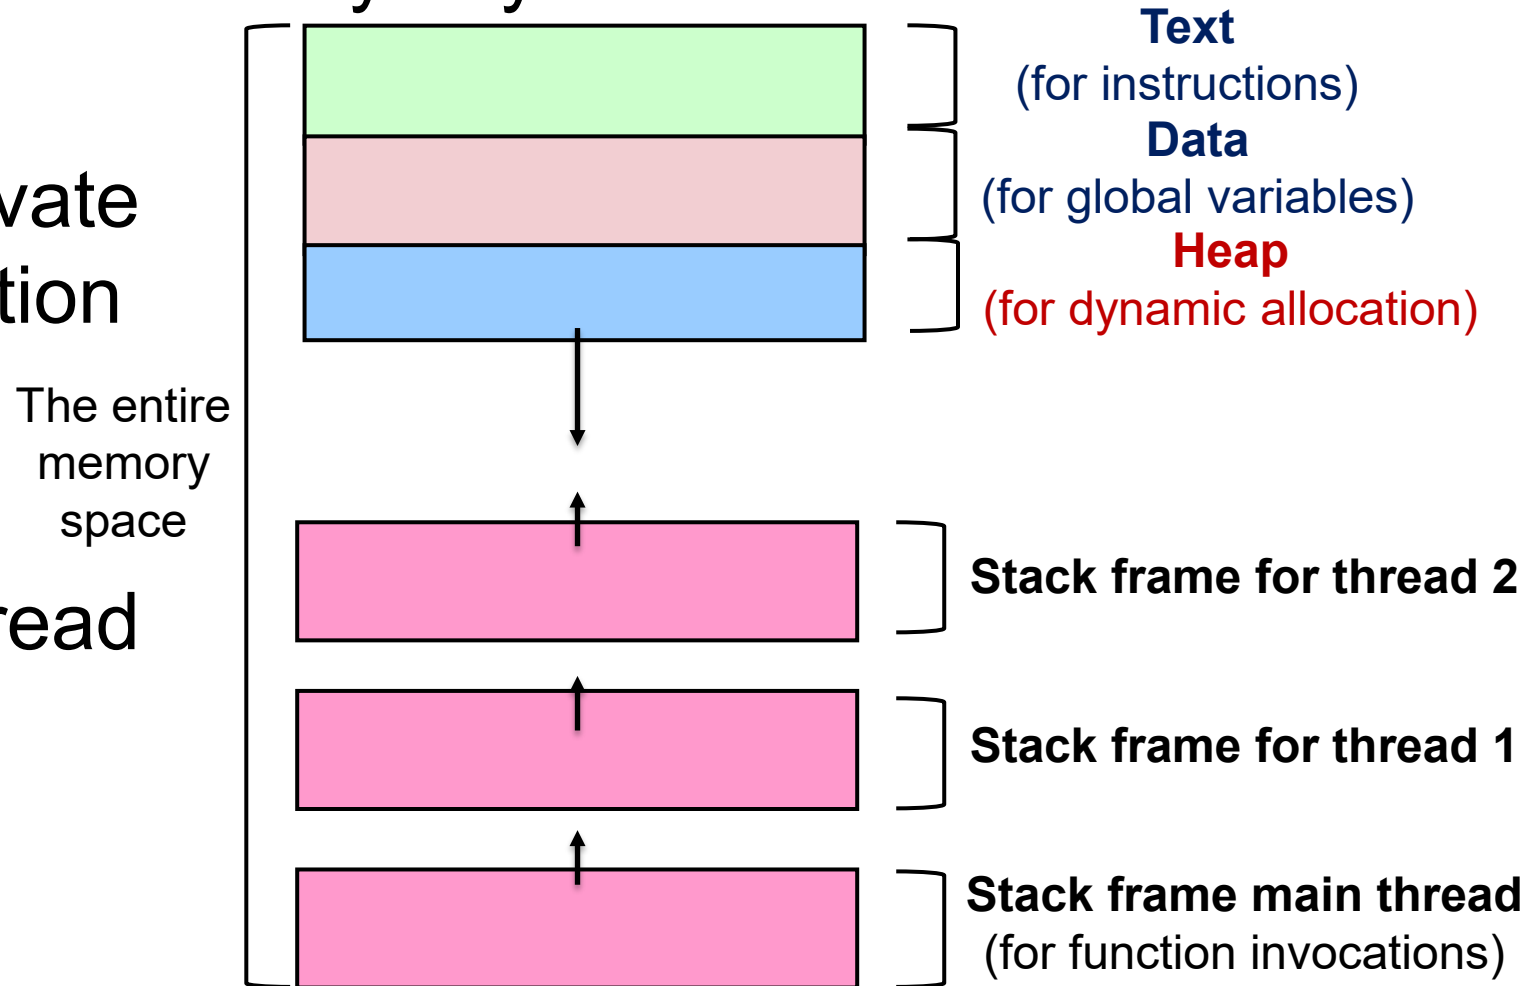


# Visibility of Data

- Global variables of a program and all dynamically allocated data objects can be accessed by any thread of this process

- Each thread has a private runtime stack for function stack frames

- Runtime stack of a thread exists iff the thread is active



# Number of Threads

- Number of threads should be
  - ❑ Suitable to parallelism degree of application
  - ❑ Suitable to available execution resources
  - ❑ Not be too large to keep the overhead for thread creation, management, and termination small

# POSIX Threads

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
void *print_message_function( void *ptr );
main()
{
    pthread_t thread1, thread2;
    char *message1 = "Thread 1";
    char *message2 = "Thread 2";
    int  iret1, iret2;
    /* Create independent threads each of which will execute function */
    iret1 = pthread_create( &thread1, NULL, print_message_function, (void*) message1);
    iret2 = pthread_create( &thread2, NULL, print_message_function, (void*) message2);

    /* Wait till threads are complete before main continues. Unless we */
    /* wait we run the risk of executing an exit which will terminate */
    /* the process and all threads before the threads have completed. */
    pthread_join( thread1, NULL);
    pthread_join( thread2, NULL);

    printf("Thread 1 returns: %d\n",iret1);
    printf("Thread 2 returns: %d\n",iret2);
    exit(0);
}
void *print_message_function( void *ptr )
{
    char *message;
    message = (char *) ptr;
    printf("%s \n", message);
}
```

[ CS3210 - AY2021S1 - L02 ]

# SYNCHRONIZATION

# Introduction

- Threads cooperate in multithreaded programs
    - Share resources, access shared data structures
    - Coordinate their execution
      - One thread executes relative to another
  - For correctness, control this cooperation
    - Threads interleave executions arbitrarily and at different rates
    - Scheduling is not under program control
  - Use synchronization
    - Restrict the possible interleaving of thread executions
- \*Discuss in terms of threads, also applies to processes

# Shared Resources

- Coordinating access to shared resources

- **Basic problem:**

- If two concurrent threads (processes) are accessing a shared variable, and that variable is read/ modified/ written by those threads, then access to the variable must be controlled to avoid erroneous behavior

- **Mechanisms to control access to shared resources**

- Locks, mutexes, semaphores, monitors, condition variables, etc.

- **Patterns for coordinating accesses to shared resources**

- Bounded buffer, producer-consumer, etc.

# Classic Example

- Implement a function to handle withdrawals from a bank account:

```
withdraw (account, amount) {  
    balance = get_balance(account);  
    balance = balance - amount;  
    put_balance(account, balance);  
    return balance;  
}
```

- 2 people share a bank account with a balance of \$1000
- Simultaneously withdraw \$100 from the account

# Classic Example - Threading

- Create a thread for each person to do the withdrawals
- These threads run on the same bank server:

```
withdraw (account, amount) {  
    balance = get_balance(account);  
    balance = balance - amount;  
    put_balance(account, balance);  
    return balance;  
}
```

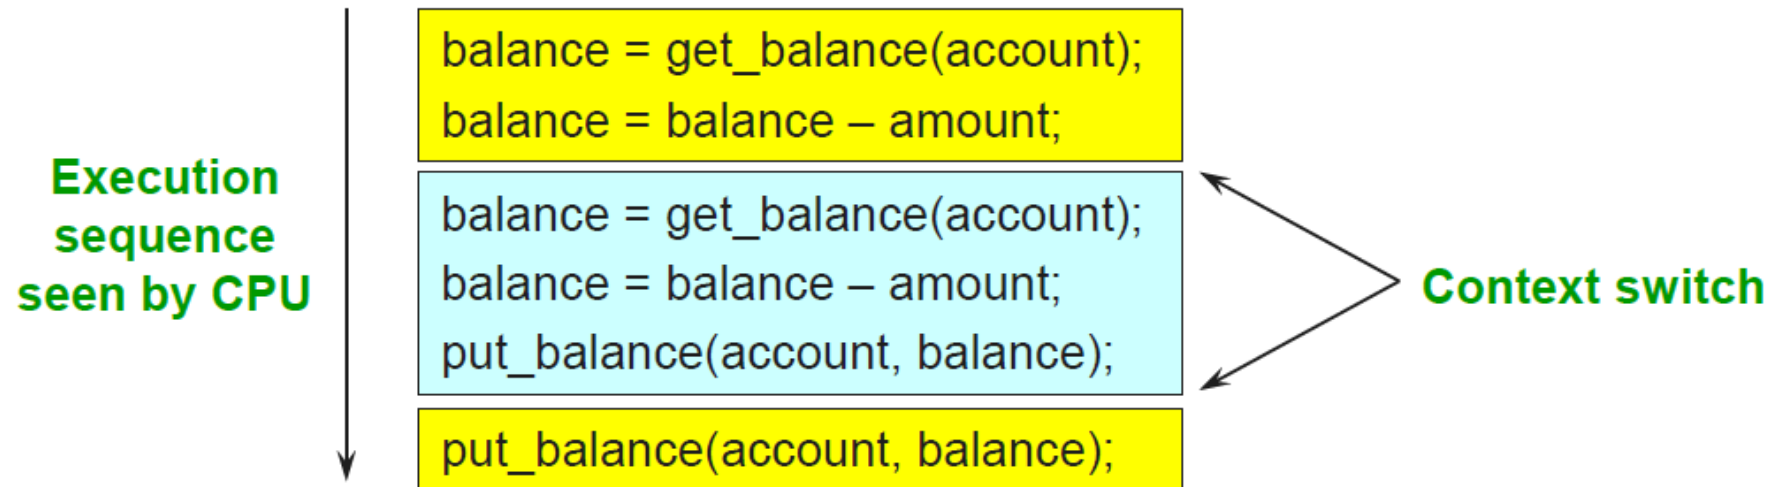
```
withdraw (account, amount) {  
    balance = get_balance(account);  
    balance = balance - amount;  
    put_balance(account, balance);  
    return balance;  
}
```

- Possible problems?



# Classic Example - Problem

- Execution of the two threads can be interleaved



# Race Condition

- Two concurrent threads (or processes) accessed a **shared resource** (account) without any **synchronization**
  - Known as a **race condition**
- Control access to these shared resources
- Necessary to synchronize access to **any shared data structure**
  - Buffers, queues, lists, hash tables, etc.

# Mutual Exclusion

- Use **mutual exclusion** to synchronize access to shared resources
  - This allows us to have large atomic blocks
- Code sequence that uses mutual exclusion is called **critical section**
  - Only one thread at a time can execute in the critical section
  - All other threads have to wait on entry
  - When a thread leaves a critical section, another can enter

# Critical Section Requirements

## 1) Mutual exclusion (mutex)

- ◆ If one thread is in the critical section, then no other is

## 2) Progress

- ◆ If some thread T is not in the critical section, then T cannot prevent some other thread S from entering the critical section
- ◆ A thread in the critical section will eventually leave it

## 3) Bounded waiting (no starvation)

- ◆ If some thread T is waiting on the critical section, then T will eventually enter the critical section

## 4) Performance

- ◆ The overhead of entering and exiting the critical section is small with respect to the work being done within it

# Critical Section Requirements - Details

- Requirements:
  - **Safety property**: nothing bad happens
    - Mutex
  - **Liveness property**: something good happens
    - Progress, Bounded Waiting
  - **Performance requirement**
- Properties hold for each run, while performance depends on all the runs
- Rule of thumb: When designing a concurrent algorithm, **worry about safety first (but don't forget liveness!)**

# Mechanisms

## ■ Locks

- ❑ Primitive, minimal semantics, used to build others

## ■ Semaphores

- ❑ Basic, easy to get the hang of, but hard to program with

## ■ Monitors

- ❑ High-level, requires language support, operations implicit

## ■ Messages

- ❑ Simple model of communication and synchronization based on atomic transfer of data across a channel
- ❑ Direct application to distributed systems
- ❑ Messages for synchronization are straightforward (once we see how the others work)

# Locks

- Two operations
  - `acquire()`: to enter a critical section
  - `release()`: to leave a critical section
- Pair calls to acquire and release
  - Between acquire/release, the thread holds the lock
  - Acquire does not return until any previous holder releases
  - What can happen if the calls are not paired?
- Locks can spin (a **spinlock**) or block (a **mutex**)

# Using Locks

```
withdraw (account, amount) {  
    acquire(lock);  
    balance = get_balance(account);  
    balance = balance – amount;  
    put_balance(account, balance);  
    release(lock);  
    return balance;  
}
```

**Critical  
Section**

```
acquire(lock);  
balance = get_balance(account);  
balance = balance – amount;
```

```
acquire(lock);
```

```
put_balance(account, balance);  
release(lock);
```

```
balance = get_balance(account);  
balance = balance – amount;  
put_balance(account, balance);  
release(lock);
```

- ◆ What happens when blue tries to acquire the lock?
- ◆ Why is the “return” outside the critical section? Is this ok?
- ◆ What happens when a third thread calls acquire?



# Semaphores

- Semaphores are an abstract data type that provide mutual exclusion through atomic counters
  - Described by Dijkstra in the “THE” system in 1968
- Semaphores are “integers” that support two operations:
  - Semaphore::Wait(): decrement, block until semaphore is open
    - Also P(), after the Dutch word for “try to reduce” (down)
  - Semaphore::Signal: increment, allow another thread to enter
    - Also V() after the Dutch word for increment (up)
  - Semaphore safety property: the semaphore value is always greater than or equal to 0

# Semaphore Types

- **Mutex** semaphore (or binary semaphore)
  - Represents single access to a resource
  - Guarantees mutual exclusion to a critical section
- Counting semaphore (or general **semaphore**)
  - Multiple threads can pass the semaphore
  - Number of threads determined by the semaphore “count”
    - mutex has count = 1, counting has count = N

# Example

```
struct Semaphore {  
    int value;  
    Queue q;  
} S;  
  
withdraw (account, amount) {  
    wait(S);  
    balance = get_balance(account);  
    balance = balance - amount;  
    put_balance(account, balance);  
    signal(S);  
    return balance;  
}
```

Threads  
block

critical  
section

It is undefined which  
thread runs after a signal

```
wait(S);  
balance = get_balance(account);  
balance = balance - amount;
```

```
wait(S);
```

```
wait(S);
```

```
put_balance(account, balance);  
signal(S);
```

```
...  
signal(S);
```

```
...  
signal(S);
```


# Semaphores Summary

- Semaphores can be used as a mutex
- However, they have some drawbacks
  - They are essentially shared global variables
    - Can potentially be accessed anywhere in program
  - No connection between the semaphore and the data being controlled by the semaphore
  - Used both for critical sections (mutual exclusion) and coordination (scheduling)
- Sometimes hard to use and prone to bugs

# Condition Variables

- Condition variables support three operations:
  - Wait—release monitor lock, wait for condition variable to be signaled
    - So condition variables have wait queues, too
  - Signal—wakeup one waiting thread
  - Broadcast—wakeup all waiting threads
- Condition variables are not boolean objects
  - “if (condition\_variable) then” ... does not make sense
  - “if (num\_resources == 0) then wait(resources\_available)” ... does

# Barrier



# Deadlock

- Definition:

- ❑ Deadlock exists among a set of processes if every process is waiting for an event that can be caused only by another process in the set.

- Deadlock is a problem that can arise:

- ❑ When processes compete for access to limited resources
  - ❑ When processes are incorrectly synchronized

# Condition for Deadlock

- Deadlock can exist if and only if the following four conditions hold simultaneously:
  1. Mutual exclusion – At least one resource must be held in a non-sharable mode
  2. Hold and wait – There must be one process holding one resource and waiting for another resource
  3. No preemption – Resources cannot be preempted (critical sections cannot be aborted externally)
  4. Circular wait – There must exist a set of processes  $[P_1, P_2, P_3, \dots, P_n]$  such that  $P_1$  is waiting for  $P_2$ ,  $P_2$  for  $P_3$ , etc.



# Dealing with Deadlock

- There are four approaches for dealing with deadlock:
  - ❑ Ignore it—how lucky do you feel?
  - ❑ Prevention—make it impossible for deadlock to happen
  - ❑ Avoidance—control allocation of resources
  - ❑ Detection and Recovery—look for a cycle in dependencies

# Starvation

- Starvation is a situation where a process is prevented from making progress because some other process has the resource it requires
- Starvation is a side effect of the scheduling algorithm
  - ❑ OS: A high priority process always prevents a low priority process from running on the CPU
  - ❑ One thread always beats another when acquiring a lock

# CLASSICAL SYNCHRONIZATION PROBLEMS

# Classic Synchronization Problems

- Producer-consumer
  - Infinite buffer
  - Finite buffer
- Readers-writers
- Dining philosophers
- Barbershop
- ...

# Producer-consumer

- Producers create items of some kind and add them to a data structure
- Consumers remove the items and process them
- Variables:
  - mutex = Semaphore (1)
  - items = Semaphore (0)

# Producer-consumer

## Producer

- `event = waitForEvent ()`
- `mutex.wait ()`
  - `buffer.add ( event )`
  - `items.signal ()`
- `mutex.signal ()`

## Consumer

- `items.wait ()`
- `mutex.wait ()`
  - `event = buffer.get ()`
- `mutex.signal ()`
- `event.process ()`

# Improved Producer-consumer

## Producer

- `event = waitForEvent ()`
- `mutex.wait ()`
  - `buffer.add ( event )`
- `mutex.signal ()`
- `items.signal ()`

## Consumer

- `items.wait ()`
- `mutex.wait ()`
  - `event = buffer.get ()`
- `mutex.signal ()`
- `event.process ()`

# Improved Producer-consumer

## Producer

- `event = waitForEvent ()`
- `mutex.wait ()`
  - `buffer.add ( event )`
- `mutex.signal ()`
- `items.signal ()`

## Consumer

- `items.wait ()`
- `mutex.wait ()`
  - `event = buffer.get ()`
- `mutex.signal ()`
- `event.process ()`

Kahoot quiz



# Broken Producer-consumer

## Producer

- `event = waitForEvent ()`
- `mutex.wait ()`
  - `buffer.add ( event )`
- `mutex.signal ()`
- `items.signal ()`

## Consumer

- `mutex.wait ()`
  - `items.wait ()`
  - `event = buffer.get ()`
- `mutex.signal ()`
- `event.process ()`

# Producer-consumer with Finite Buffer

## Producer

- `event = waitForEvent ()`
- `spaces.wait ()`
- `mutex.wait ()`
  - `buffer.add ( event )`
- `mutex.signal ()`
- `items.signal ()`

## Consumer

- `items.wait ()`
- `mutex.wait ()`
  - `event = buffer.get ()`
- `mutex.signal ()`
- `spaces.signal ()`
- `event.process ()`

# Readers-writers problem

- Any number of readers can be in the critical section simultaneously
- Writers must have exclusive access to the critical section
- Variables:
  - ❑ `int readers = 0`
  - ❑ `mutex = Semaphore (1)`
  - ❑ `roomEmpty = Semaphore (1)`

# Readers-writers

## Writers

- roomEmpty.wait ()
  - #critical section for writers
- roomEmpty.signal ()

## Readers

- mutex.wait ()
  - readers += 1
  - if readers == 1:
    - roomEmpty.wait () # first in locks
- mutex.signal ()
- # critical section for readers
- mutex.wait ()
  - readers -= 1
  - if readers == 0:
    - roomEmpty.signal () # last out unlocks
- mutex.signal ()

# Lightswitch Definition

class Lightswitch :

- def `__init__` ( self ):
  - `self.counter = 0`
  - `self.mutex = Semaphore (1)`
- def `lock` (self , semaphore ):
  - `self.mutex.wait ()`
    - `self.counter += 1`
    - if `self.counter == 1`:
      - `semaphore.wait ()`
    - `self.mutex.signal ()`
- def `unlock` (self , semaphore ):
  - `self.mutex.wait ()`
    - `self.counter -= 1`
    - if `self.counter == 0`:
      - `semaphore.signal ()`
  - `self.mutex.signal ()`

# Readers-writers with Lightswitch

## Writers

- roomEmpty.wait ()
    - #critical section for writers
  - roomEmpty.signal ()
- 
- #starving writers
  - Use a
    - turnstile = Semaphore (1)

## Readers

- readLightswitch.lock (roomEmpty )
  - # critical section
- readLightswitch.unlock (roomEmpty)

# No-starve Readers-writers

## Writers

- `turnstile.wait ()`
  - `roomEmpty.wait ()`
    - # critical section for writers
- `turnstile.signal ()`
- `roomEmpty.signal ()`

## Readers

- `turnstile.wait ()`
- `turnstile.signal ()`
- `readSwitch.lock ( roomEmpty )`
  - # critical section for readers
- `readSwitch.unlock ( roomEmpty )`

# Readers-writers with priorities

## Writers

- writeSwitch.lock (noReaders)
  - noWriters.wait ()
  - # critical section for writers
  - noWriters.signal()
- writeSwitch.unlock (noReaders)

## Readers

- noReaders.wait ()
  - readSwitch.lock (noWriters)
- noReaders.signal ()
- # critical section for readers
- readSwitch.unlock ( noWriters )



# Readings

- Main reference:
  - Chapter 3.8, 6.1
- CSE 120: Principles of Computer Operating Systems, UCSD, <http://cseweb.ucsd.edu/classes/fa16/cse120-a/>
- The Little Book of Semaphores by Allen Downey, <http://greenteapress.com/semaphores/LittleBookOfSemaphores.pdf>

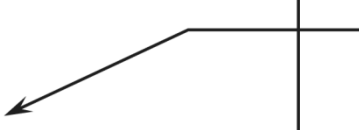
# LOCK IMPLEMENTATION

# Implementing Locks (1)

- An attempt:

```
struct lock {  
    int held = 0;  
}  
void acquire (lock) {  
    while (lock→held);  
    lock→held = 1;  
}  
void release (lock) {  
    lock→held = 0;  
}
```

busy-wait (spin-wait)  
for lock to be released

A black arrow originates from the right side of the code block, specifically pointing to the `while (lock→held);` line, and extends to the left side of the callout box.

- This is called a spinlock because a thread spins waiting for the lock to be released
- Does this work?

## Implementing Locks (2)

- No. Two independent threads may both notice that a lock has been released and thereby acquire it.

```
struct lock {  
    int held = 0;  
}  
void acquire (lock) {  
    while (lock→held);  
    lock→held = 1;  
}  
void release (lock) {  
    lock→held = 0;  
}
```

A context switch can occur here, causing a race condition

# Implementing Locks (3)

- The problem: implementation of locks has critical sections, too
  - How do we stop the recursion?
- The implementation of acquire/release must be atomic
  - An atomic operation is one which executes as though it could not be interrupted
  - Code that executes “all or nothing”
- Need help from hardware
  - Atomic instructions (e.g., test-and-set)
  - Disable/enable interrupts (prevents context switches)

# Atomic Instructions: Test-and-set

- The semantics of test-and-set are:
  - ❑ Record the old value
  - ❑ Set the value to indicate available
  - ❑ Return the old value
- Hardware executes it atomically!

```
bool test_and_set (bool *flag) {  
    bool old = *flag;  
    *flag = True;  
    return old;  
}
```

# Lock with Test-and-set

```
struct lock {  
    int held = 0;  
}  
void acquire (lock) {  
    while (test-and-set(&lock→held));  
}  
void release (lock) {  
    lock→held = 0;  
}
```

# Problems with Spinlocks

- Spinlocks are wasteful
  - If a thread is spinning on a lock, then the thread holding the lock cannot make progress (on a uniprocessor)
- How did the lock holder give up the CPU in the first place?
  - Lock holder calls yield or sleep
  - Involuntary context switch



# Higher-level Synchronization

- All synchronization requires atomicity
  - Use “atomic”
- Look at two common high-level mechanisms
  - **Semaphores**: binary (mutex) and counting
  - **Monitors**: mutexes and condition variables locks as primitives