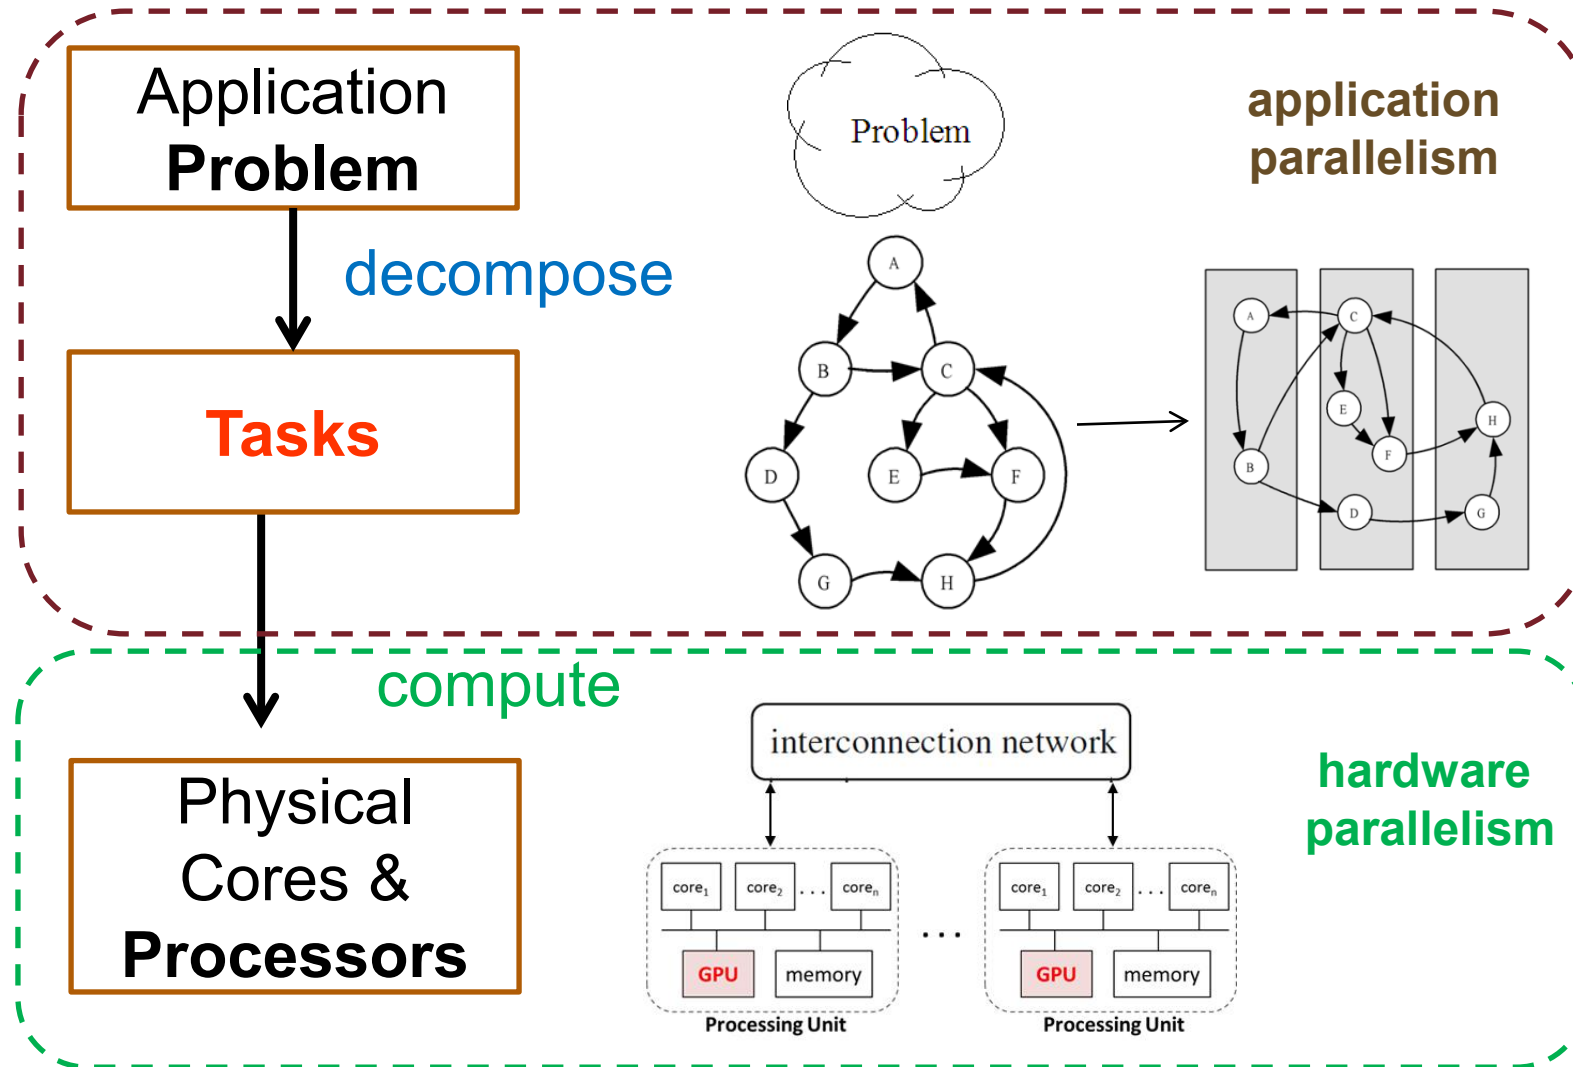


# Parallel Computing Platforms

## Lecture 03

# Parallel Computing



# Computer architecture

- Key concepts about how modern computers work
  - Concerns on parallel execution
  - Challenges of accessing memory
- Understanding these architecture basics help you
  - **Understand and optimize the performance** of your parallel programs
  - **Gain intuition** about what workloads might benefit from fast parallel machines

# Outline


- Processor Architecture and Technology Trends
  - Various forms of parallelism
- Flynn's Parallel Architecture Taxonomy
- Architecture of Multicore Processors
- Memory Organization
  - Distributed-memory Systems
  - Shared-memory Systems
  - Hybrid (Distributed-Shared Memory) Systems

# Source of Processor Performance Gain

- Parallelism of various forms are the main source of performance gain

- Let us understand parallelism at the:

- ❑ Bit Level
- ❑ Instruction Level
- ❑ Thread Level
- ❑ Process Level
- ❑ Processor Level:
  - Shared Memory
  - Distributed Memory



# Recap: Execution Time

$$\text{CPU Time} = \frac{\text{Seconds}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Cycle}}$$

Component	Affected By
Instruction / Program	<ul style="list-style-type: none"><li>• Compiler</li><li>• ISA</li></ul>
Cycles / Instruction	<ul style="list-style-type: none"><li>• Processor Implementation</li></ul>
Seconds / Cycle	<ul style="list-style-type: none"><li>• Clock Speed (Clock Frequency)</li></ul>

# Bit Level Parallelism

- Word size may mean:
  - ❑ Unit of transfer between processor  $\leftrightarrow$  memory
  - ❑ Memory address space capacity
  - ❑ Integer size
  - ❑ Single precision floating point number size
- Word size trend:
  - ❑ Varied in the 50s to 70s
  - ❑ Following x86 processors:
    - 16 bits (8086 – 1978)
    - 32 bits (80386 – 1985)
    - 64 bits (Pentium 4 / Opteron – 2003)


# Instruction Level Parallelism

- Execute instructions in parallel:
  - Pipelining (parallelism across time)
  - Superscalar (parallelism across space)
- Pipelining:
  - Split instruction execution in multiple stages, e.g.
    - Fetch (IF), Decode (ID), Execute (EX), Write-Back (WB)
  - Allow multiple instructions to occupy different stages in the same clock cycle
    - Provided there is no data / control dependencies
  - Number of pipeline stages == Maximum achievable speedup




# Pipelined Execution: Illustration

## Non-pipelined




## Pipelined



- Disadvantages
  - ❑ Independence
  - ❑ Bubbles
  - ❑ Hazards: data and control flow
- Speculation
- Out-of-order execution
  - ❑ Read-after-write

# The end of ILP



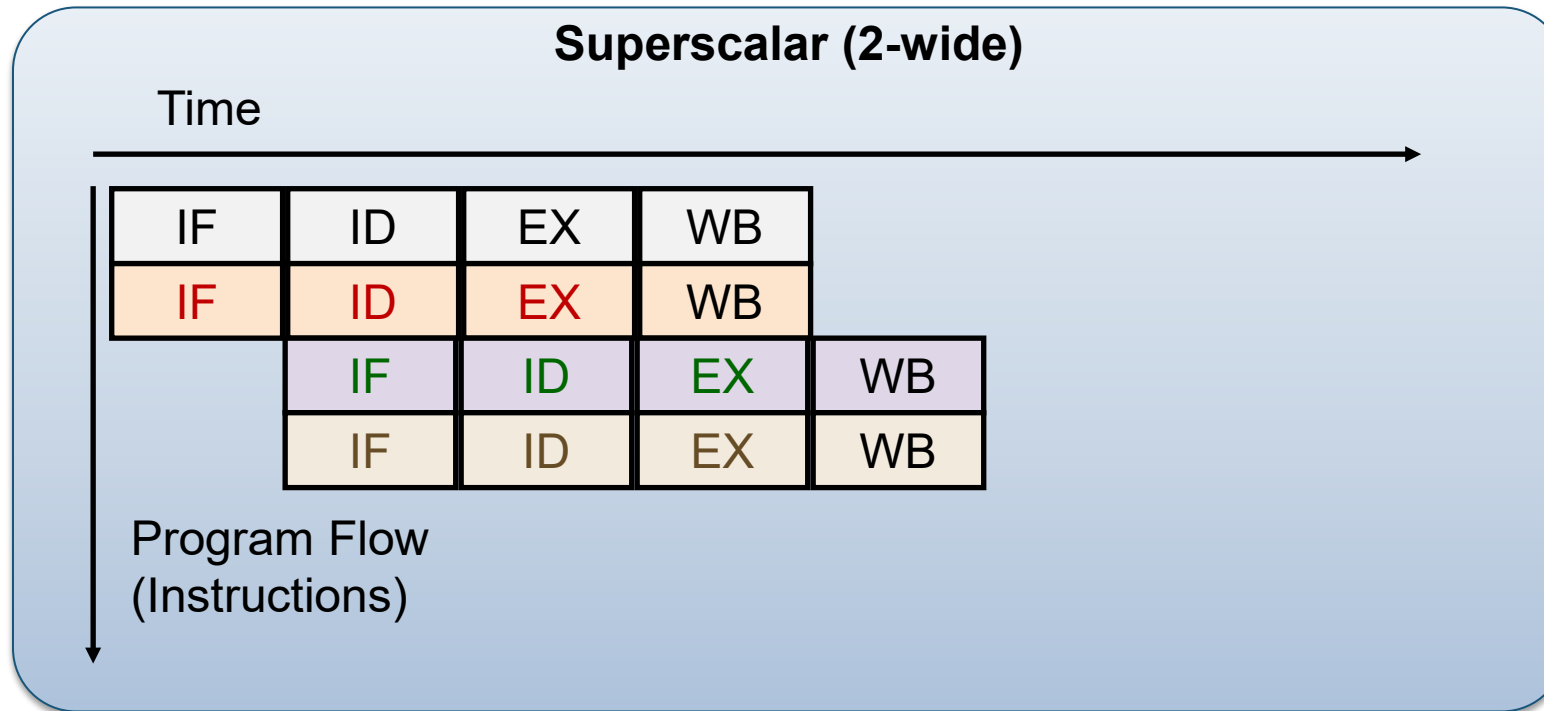
Processor clock rate stops increasing

No further benefit from ILP

# Instruction Level Parallelism: **Superscalar**

- Duplicate the pipelines:
  - Allow multiple instructions to pass through the same stage
  - Scheduling is challenging (decide which instructions can be executed together):
    - Dynamic (Hardware decision)
    - Static (Compiler decision)
- Most modern processors are superscalar
  - e.g. each intel i7 core has 14 pipeline stages and can execute 6 micro-ops in the same cycle

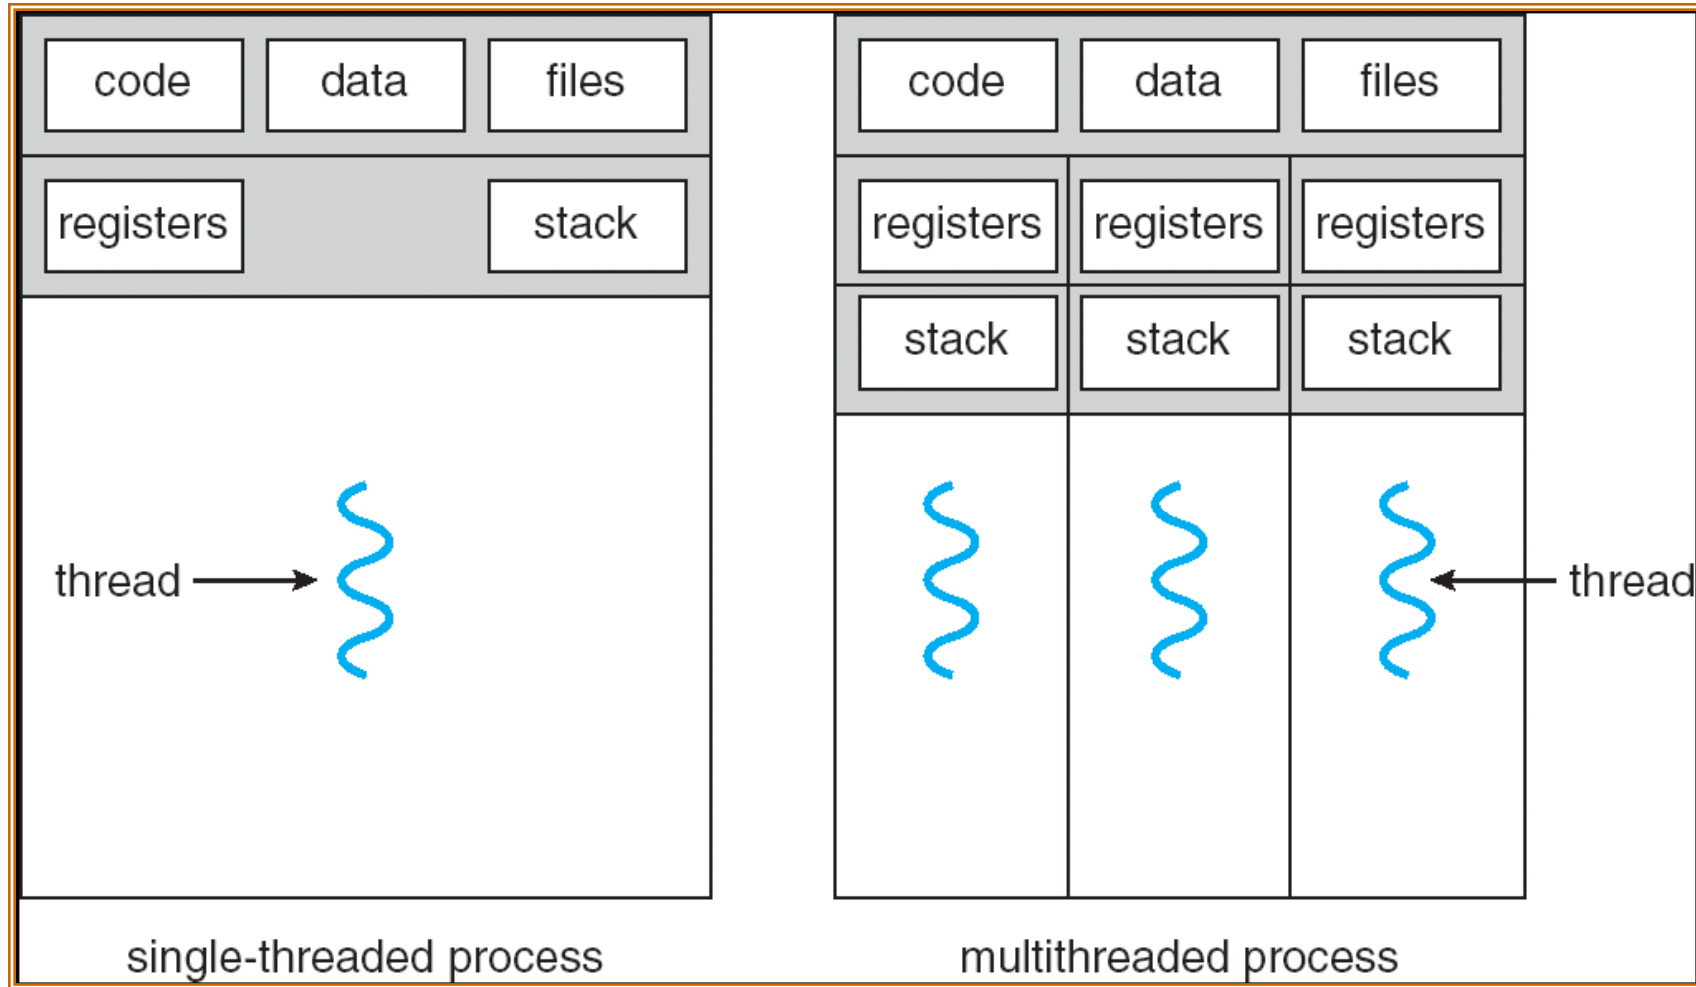
# Superscalar Execution: Illustration



Disadvantages:  
structural hazard

- Cycles-per-instruction (CPI) → Instructions-per-cycle (IPC) (Why?)
- How does this change the execution time calculation?

# Recap: Process vs Thread



- Taken from Operating System Concepts (7<sup>th</sup> Edition) by Silberschatz, Galvin & Gagne, published by Wiley

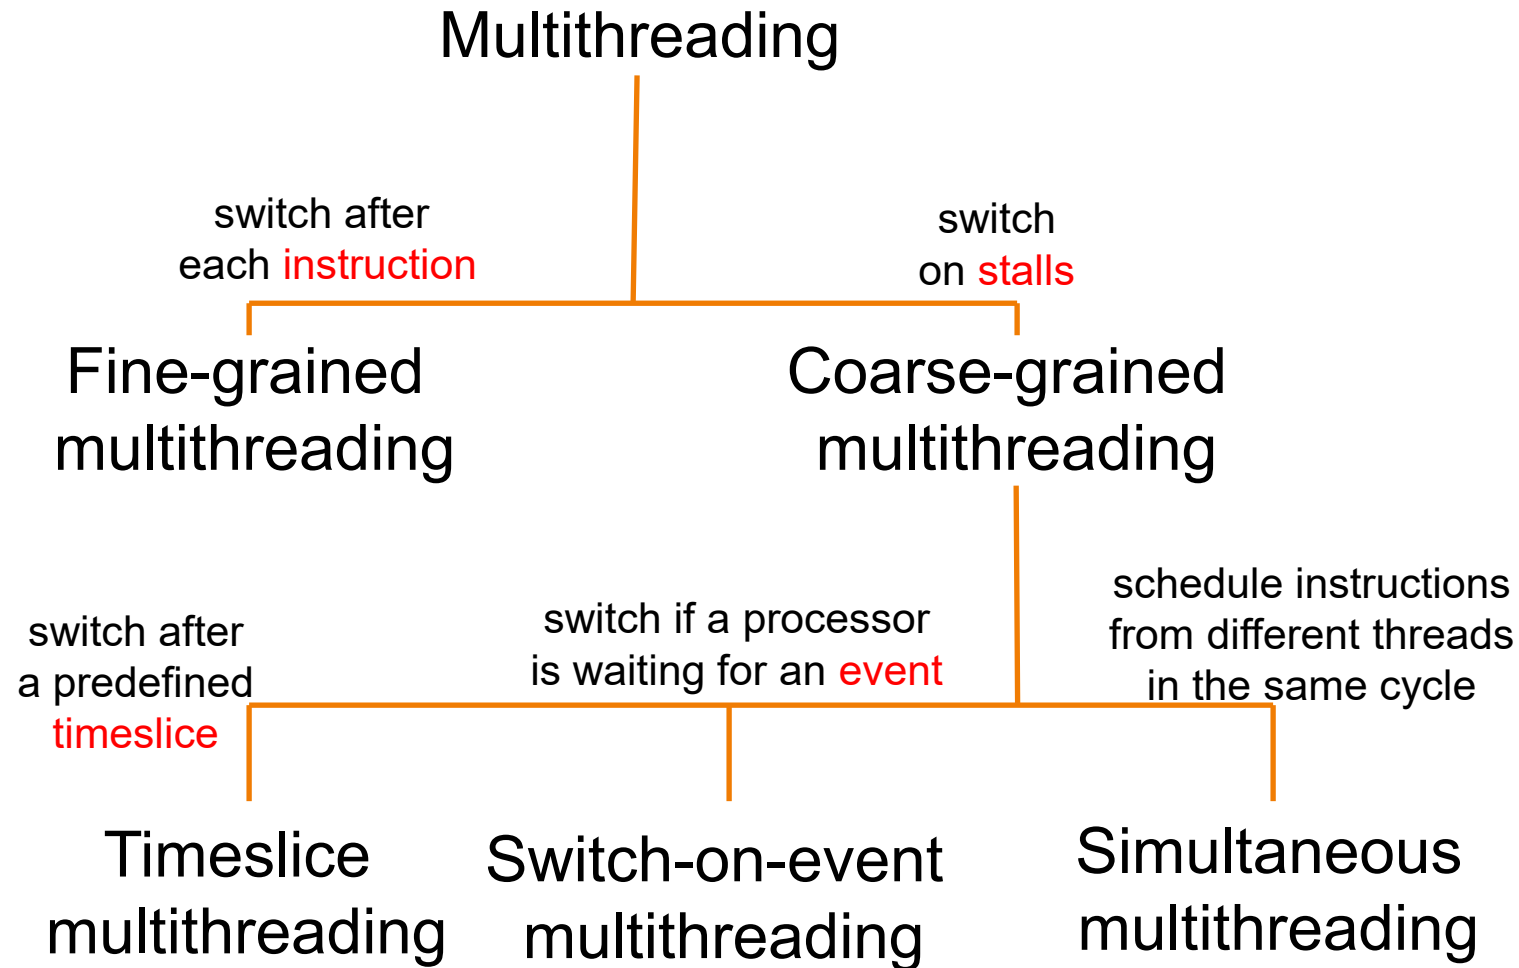
# Thread Level Parallelism: Motivation

- Instruction level parallelism is **limited**
  - ❑ For typical programs only 2-3 instructions can be executed in parallel (either pipelined / superscalar )
  - ❑ Due to data / control dependencies
- Multithreading was originally a software mechanism
  - ❑ Allow multiple parts of the same program to execute concurrently
- Key idea:
  - ❑ What if the processor can execute the threads in parallel?

# Thread Level Parallelism: Processor


- Processor can provide hardware support for the "thread context"
  - ❑ Information specific to each thread, e.g. Program Counter, Registers, etc
  - ❑ Software threads can then execute in parallel
  - ❑ Many implementation approaches
- Example:
  - ❑ [Simultaneous Multi-Threading] intel processors with **hyper-threading** technology, e.g. each i7 core can execute 2 threads at the same time

# Multithreading Implementations



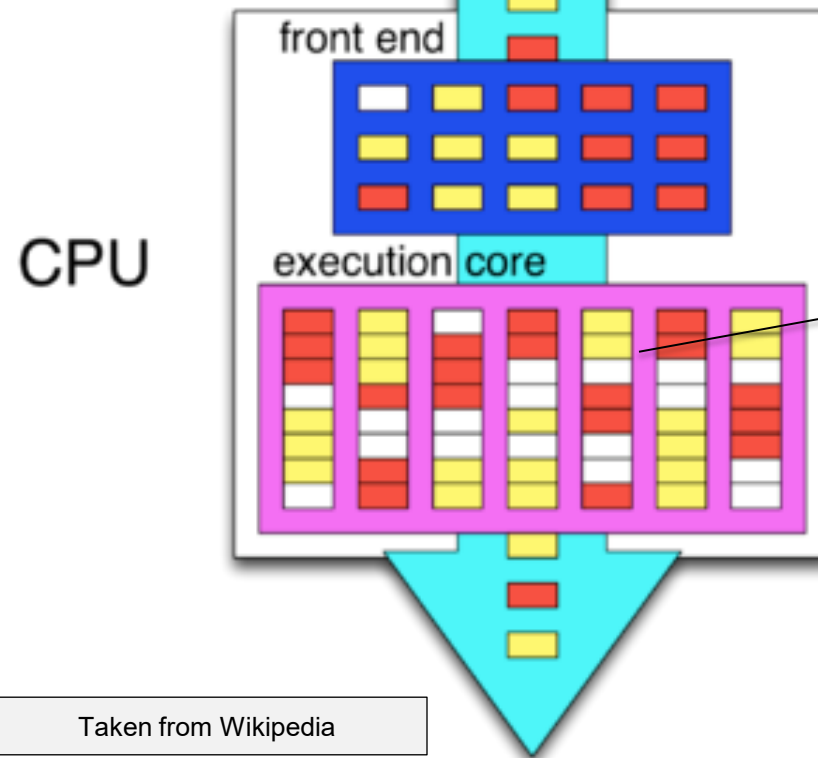


# Thread Level Parallelism: Illustration



Each box represents an instruction

Color represents different threads,  
e.g. there are 4 threads here



instructions from different threads  
are executed in parallel

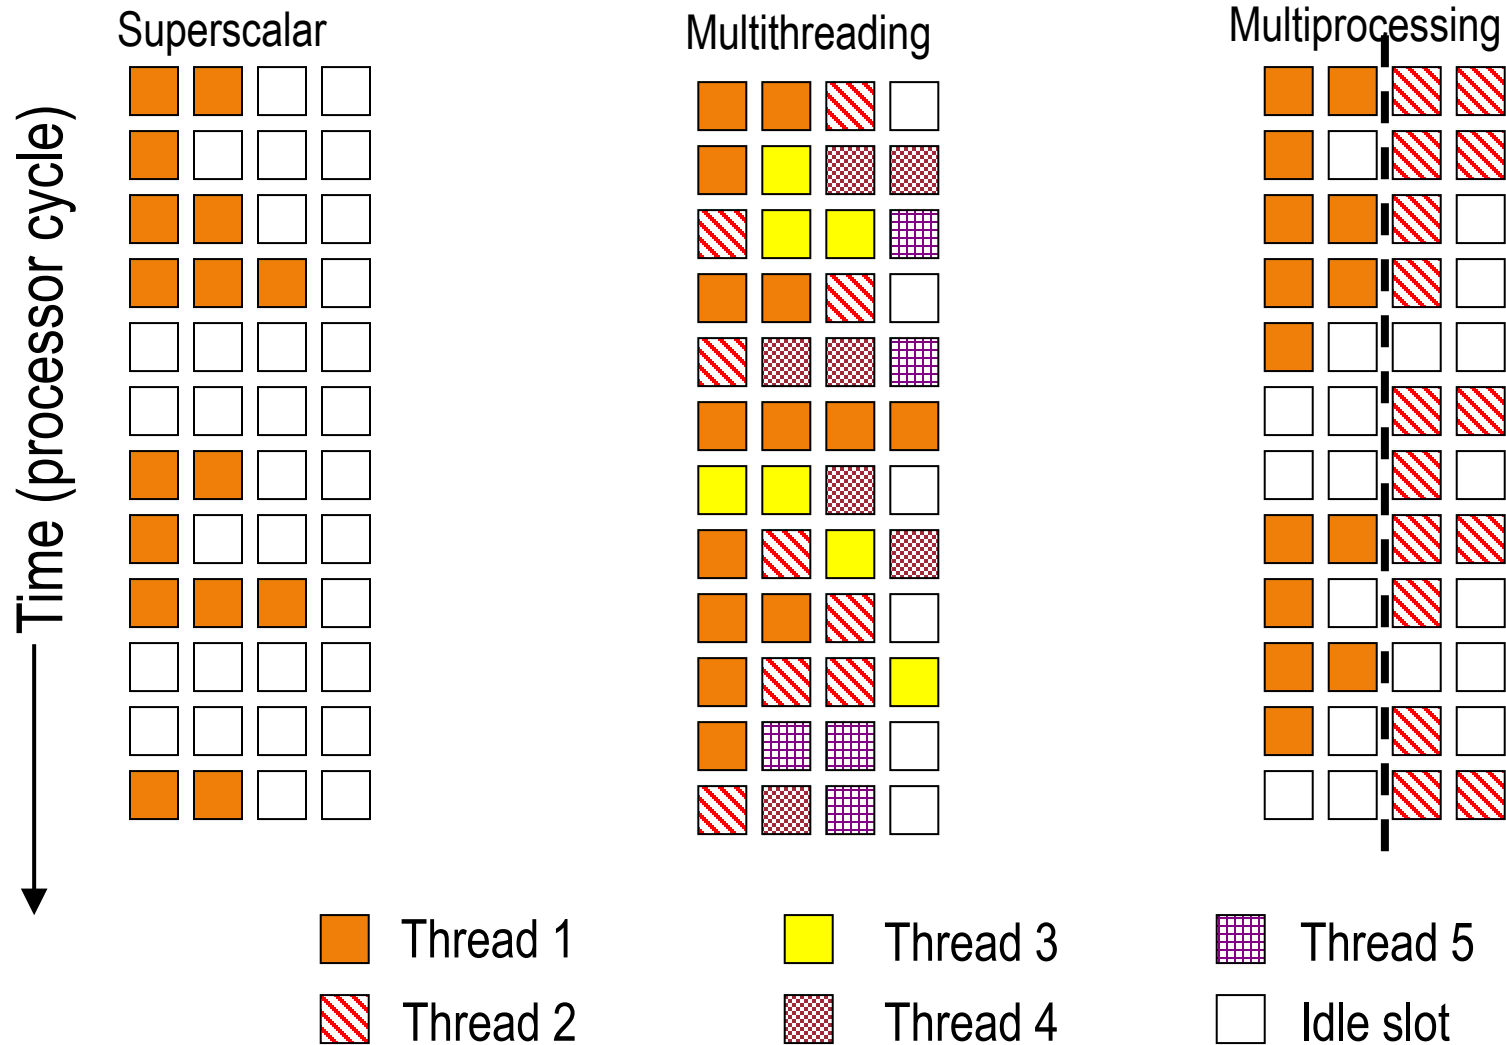
This example shows instructions  
from 2 threads are able to be  
executed together

Taken from Wikipedia

# Process Level Parallelism

- Instead of multiple threads, we can use multiple **processes** to work in parallel
  - Process has independent memory space → need special mechanism for communication
  - Operating system commonly provides IPC (Inter-Process communication) mechanisms
- Each process needs an independent set of processor context  
→ can be mapped to multiple processor cores

# Single Processor Parallelism: Summary




# Flynn's Parallel Architecture Taxonomy

- One commonly used taxonomy of parallel architecture:
  - ❑ Based on the parallelism of instructions and data streams in the most constrained component of the processor
  - ❑ Proposed by M.Flynn in 1972(!)
- Instruction stream:
  - ❑ A single execution flow
  - ❑ i.e. a single Program Counter (PC)
- Data stream:
  - ❑ Data being manipulated by the instruction stream

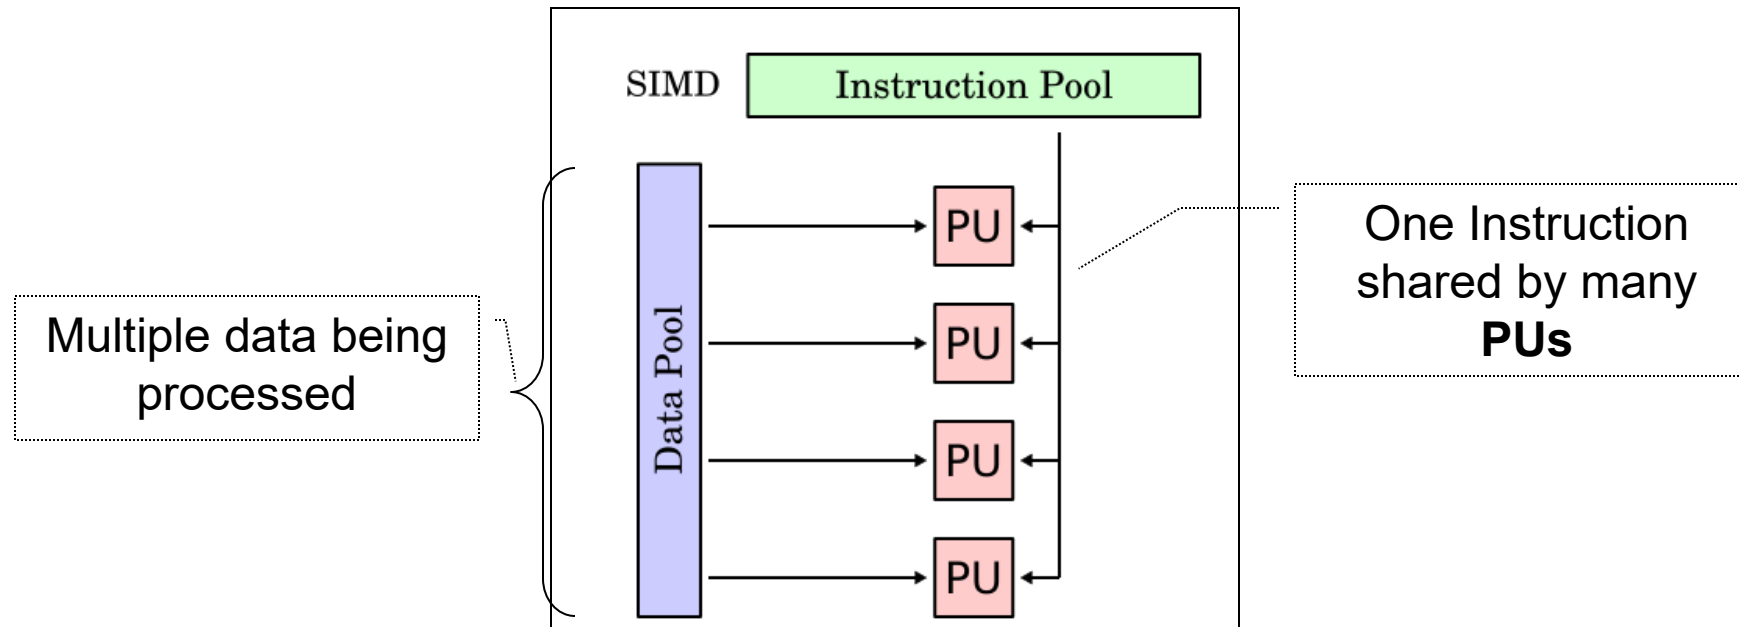
# Single Instruction Single Data (SISD)

- A single instruction stream is executed
- Each instruction work on single data
- Most of the uniprocessors fall into this category



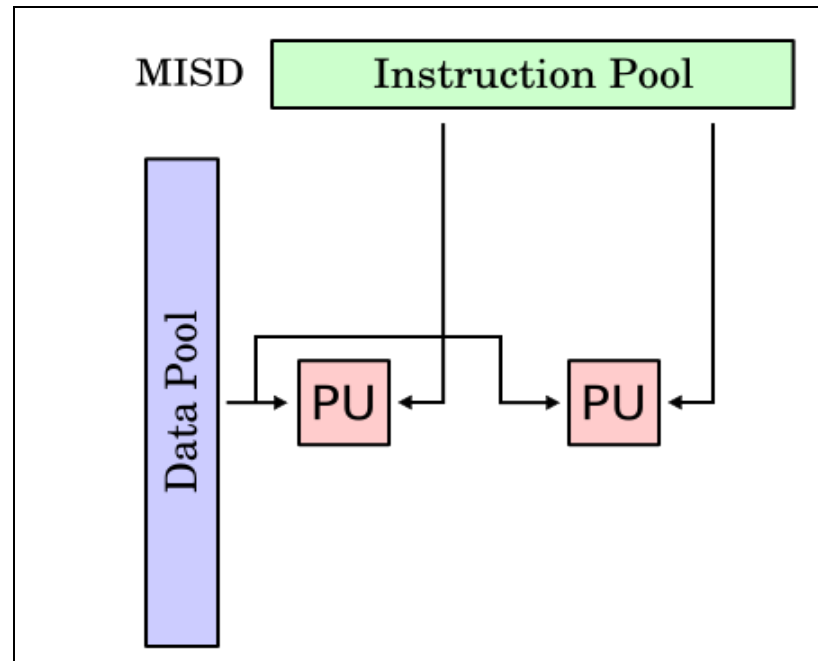
# Single Instruction Multiple Data (SIMD)

- A single stream of instructions
- Each instruction works on multiple data
- Popular model for supercomputer during 1980s:
  - ❑ Exploit **data parallelism**, commonly known as vector processor
- Modern processor has some forms of SIMD:
  - ❑ E.g. the SSE, AVX instructions in intel x86 processors




# Multiple Instruction Single Data (MISD)

- Multiple instruction streams
- All instruction work on the same data at any time
- No actual implementation // systolic array



# Multiple Instruction Multiple Data (MIMD)

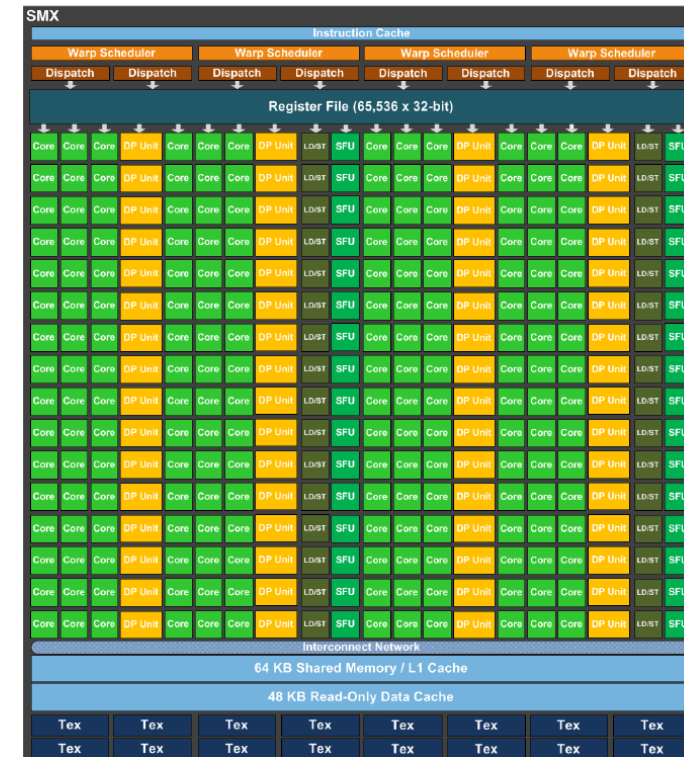
- Each PU fetch its own instruction
- Each PU operates on its data
- Currently the most popular model for multiprocessor





# Variant – SIMD + MIMD

- Stream processor (nVidia GPUs)
  - ❑ A set of threads executing the same code (effectively SIMD)
  - ❑ Multiple set of threads executing in parallel (effectively MIMD at this level)




# MULTICORE ARCHITECTURE

# Architecture of Multicore Processors

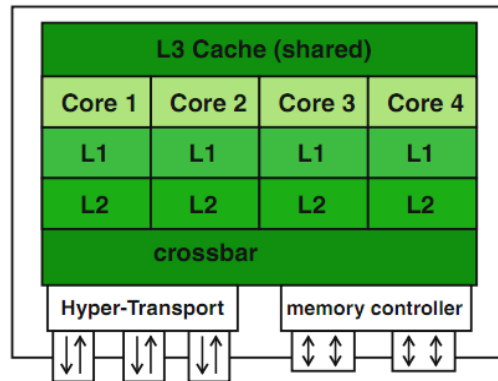
- Hierarchical design
- Pipelined design
- Network-based design

# Hierarchical Design

- Multiple cores share multiple caches
- Cache size increases from the leaves to the root
- Each core can have a separate L1 cache and shares the L2 cache with other cores
- All cores share the common external memory
- Usages
  - ❑ Standard desktop
  - ❑ Server processors
  - ❑ Graphics processing units

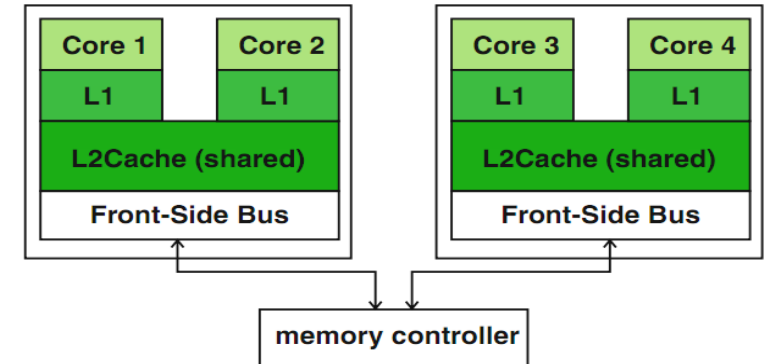


# Hierarchical Design - Examples

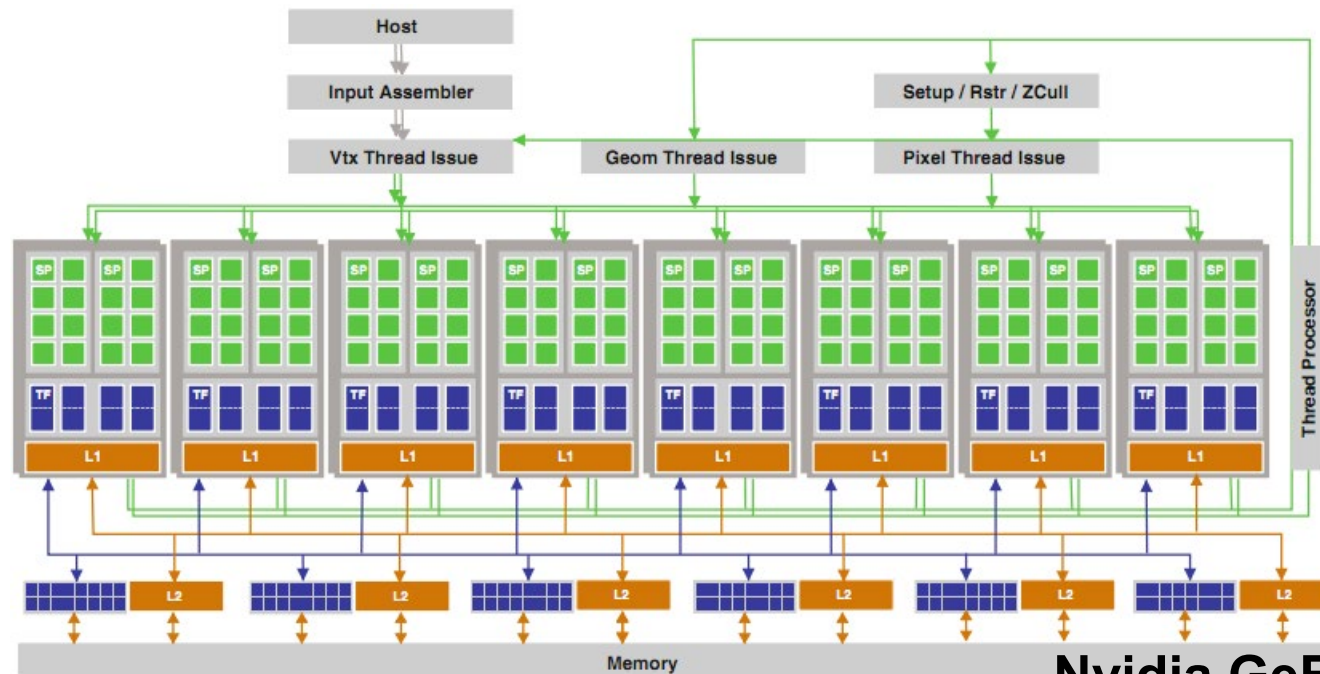


**Quad-Core AMD Opteron**

Each core is sophisticated, out-of-order processor to maximize ILP



**Intel Quad-Core Xeon**




128 stream processors (SP),  
16 texture / process clusters  
each with 8 SPs

Each core processors vectors of  
data

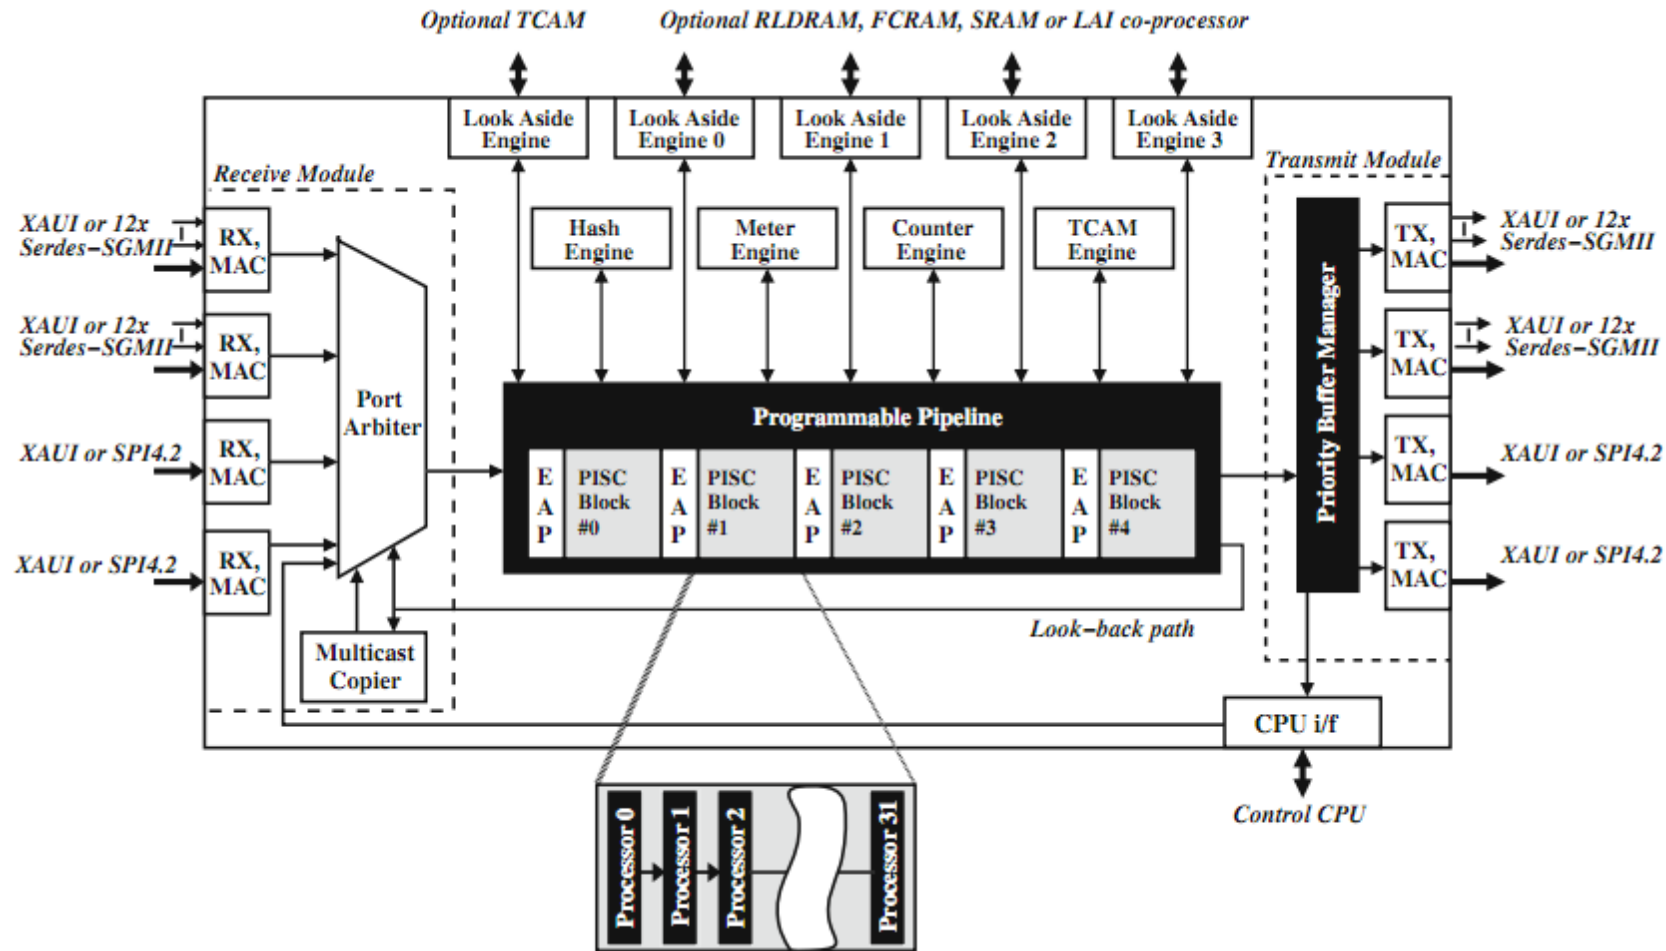
**Nvidia GeForce 8800**

# Pipelined Design

- Data elements are processed by multiple execution cores in a **pipelined way**
- Useful if same computation steps have to be applied to a long sequence of data elements
  - ❑ E.g. processors used in routers and graphics processors



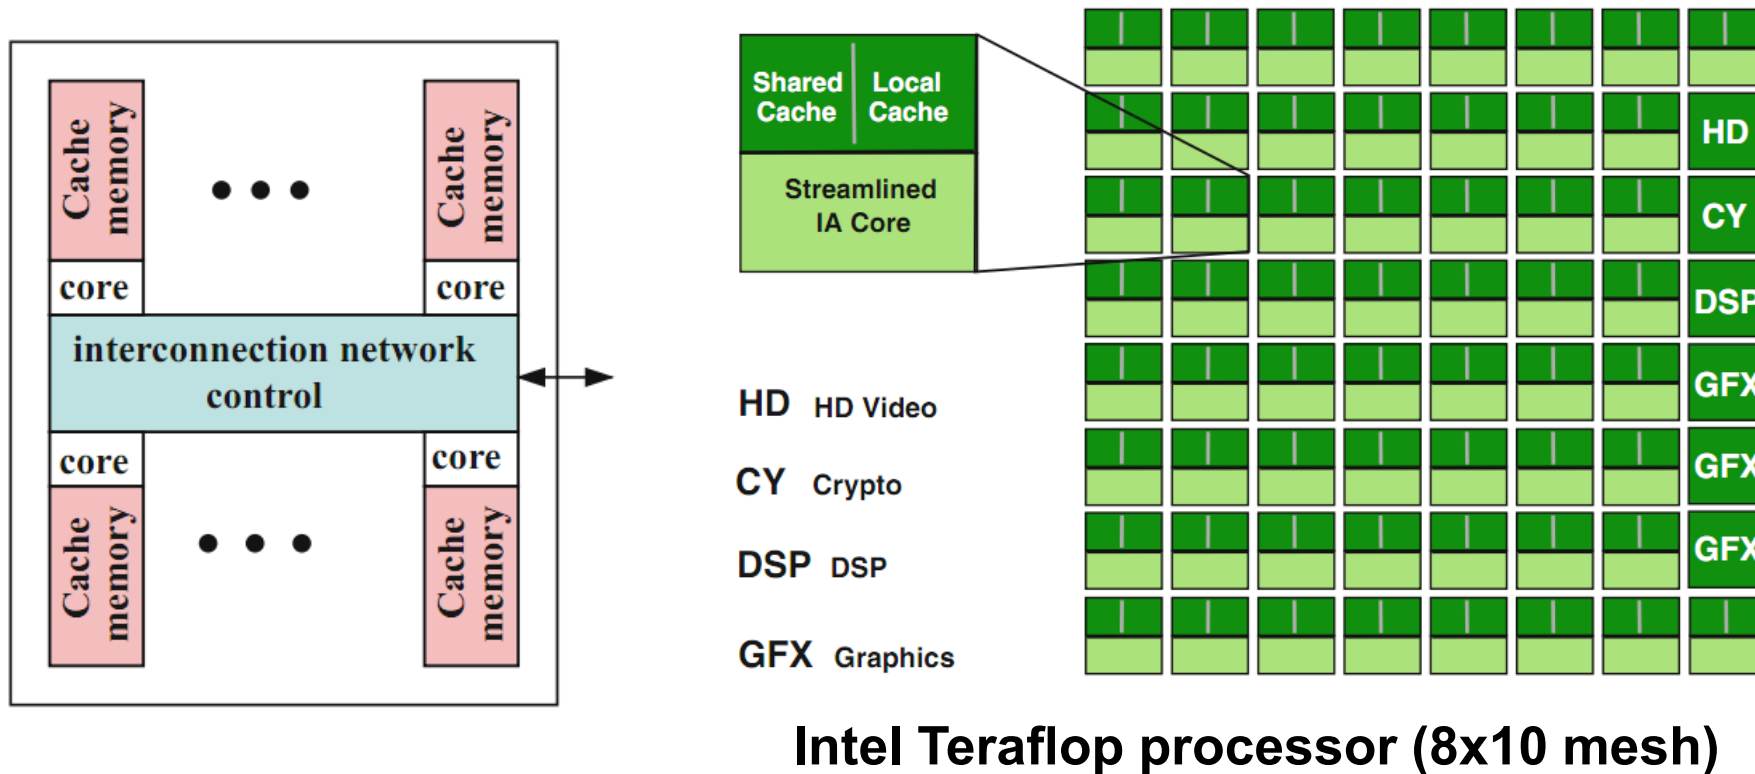
# Example: Pipelined Design



**Xelerator X11 network processor**

# Network-Based Design

- Cores and their local caches and memories are connected via an interconnection network






# Future Trends

- Efficient on-chip interconnection
  - ❑ Enough bandwidth for data transfers between the cores
  - ❑ Scalable
  - ❑ Robust to tolerate failures
  - ❑ Efficient energy management
  - ❑ Reduce memory access time
- Key word: Network on Chip (NoC)

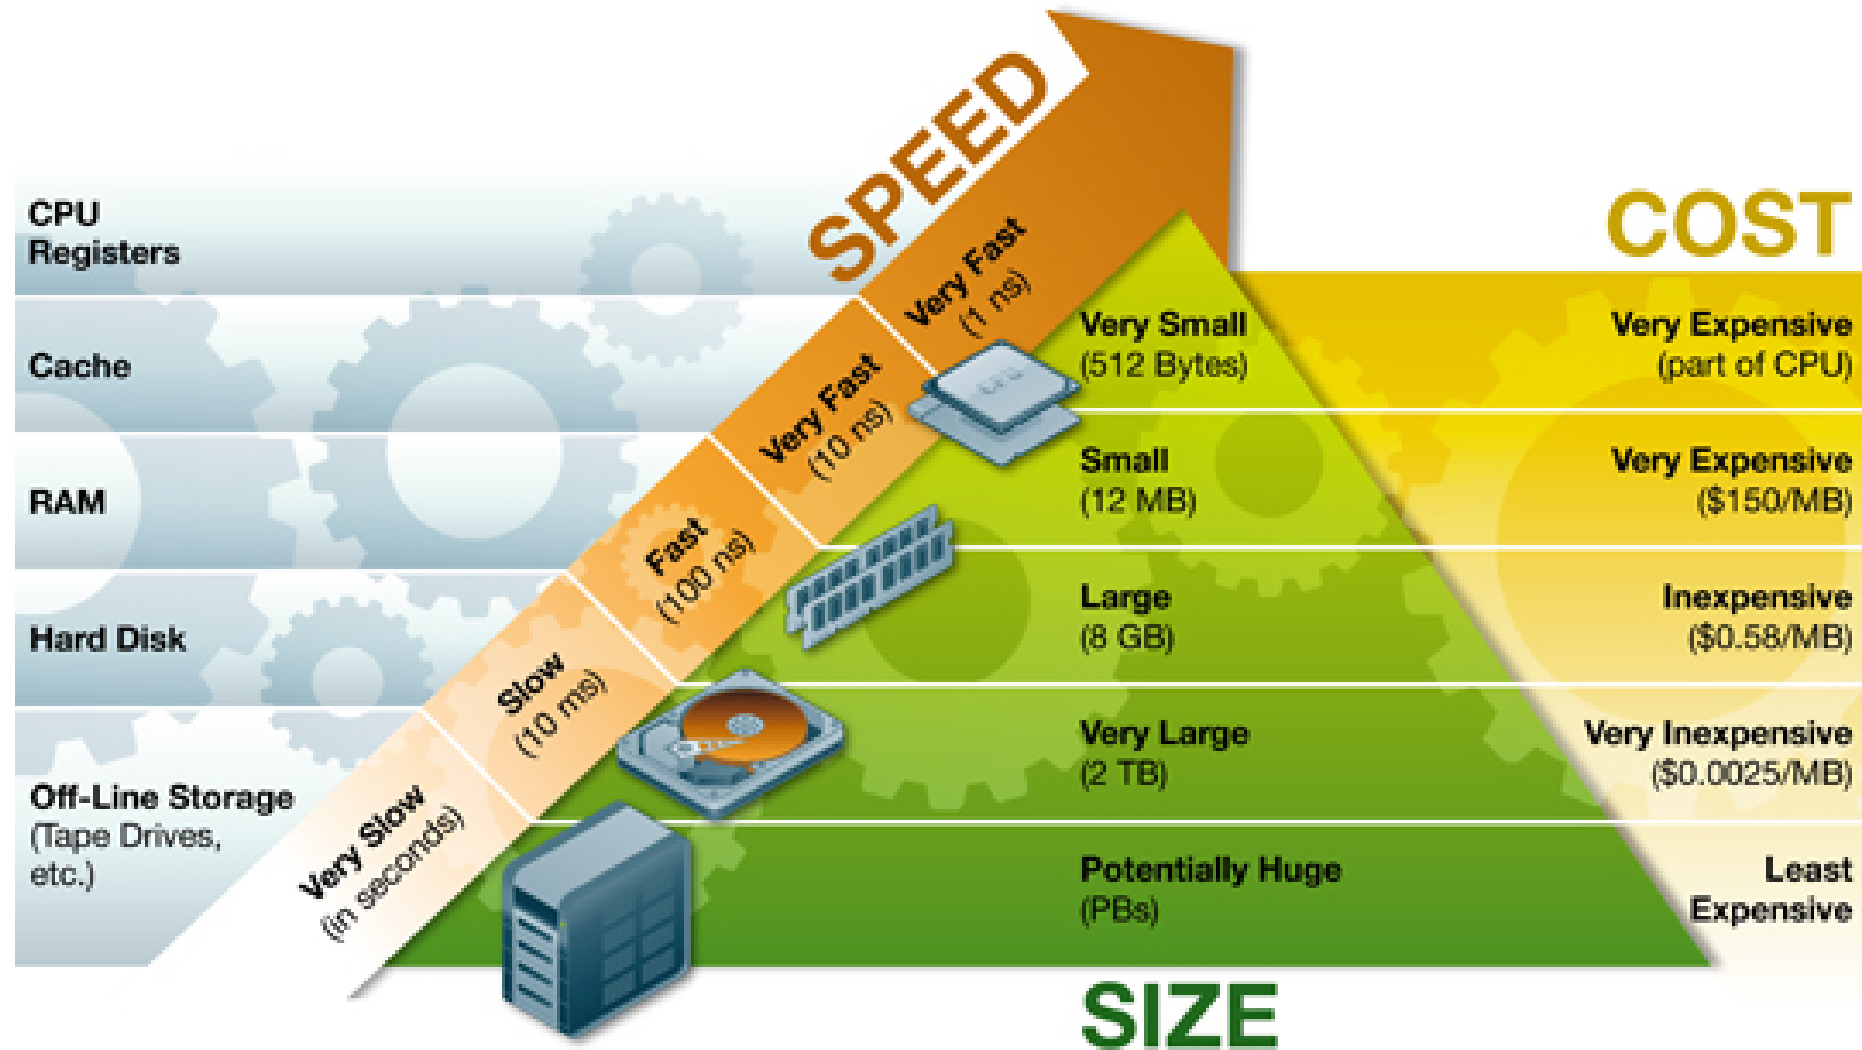
# MEMORY ORGANIZATION

# Parallel Computer Component



- Typical uniprocessor components:
  - ❑ Processor
  - ❑ One or more level of caches
  - ❑ Memory module
  - ❑ Other (e.g. I/O)
- These components similarly present in a parallel computer setup
- Processors in a parallel computer systems is also commonly known as **processing element**

# The Memory Hierarchy

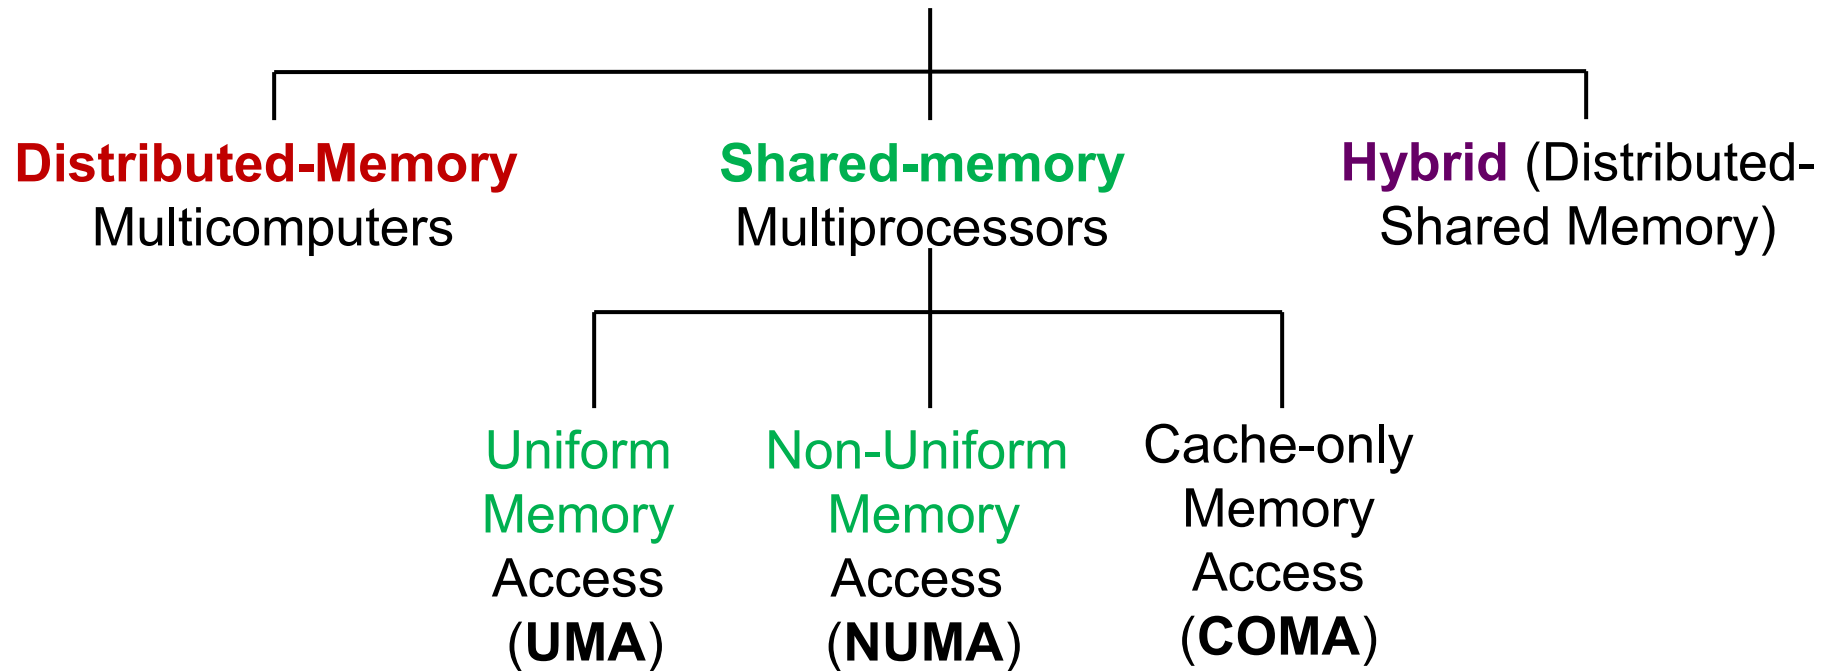


# Recap: memory latency and bandwidth


- Memory latency: the amount of time for a memory request (e.g., load, store) from a processor to be serviced by the memory system
  - Example: 100 cycles, 100 nsec
- Memory bandwidth: the rate at which the memory system can provide data to a processor
  - Example: 20 GB/s
- Processor “stalls” when it cannot run the next instruction in an instruction stream because of a dependency on a previous instruction

# Memory Organization of Parallel Computers

## Parallel Computers

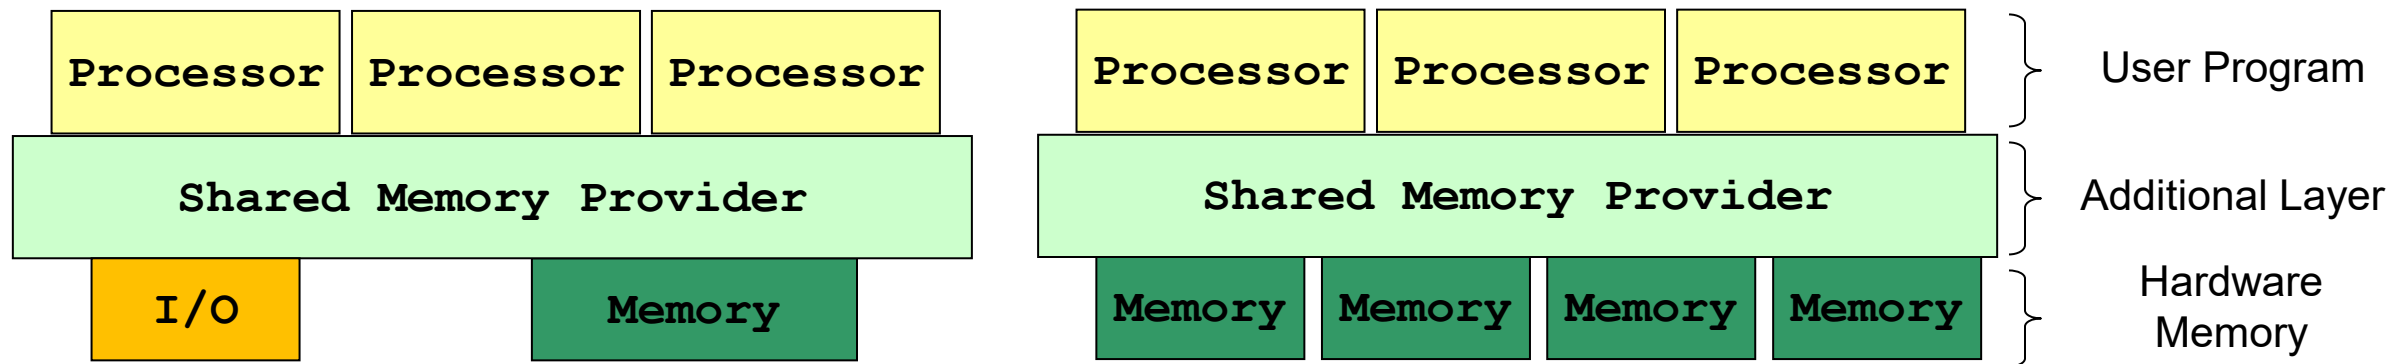


# Distributed-Memory Systems



- Each node is an independent unit:
  - With processor, memory and, sometimes, peripheral elements
- Physically distributed memory module:
  - ➔ Memory in a node is private
- Data exchanges between nodes
  - ➔ message-passing (more in later lectures)

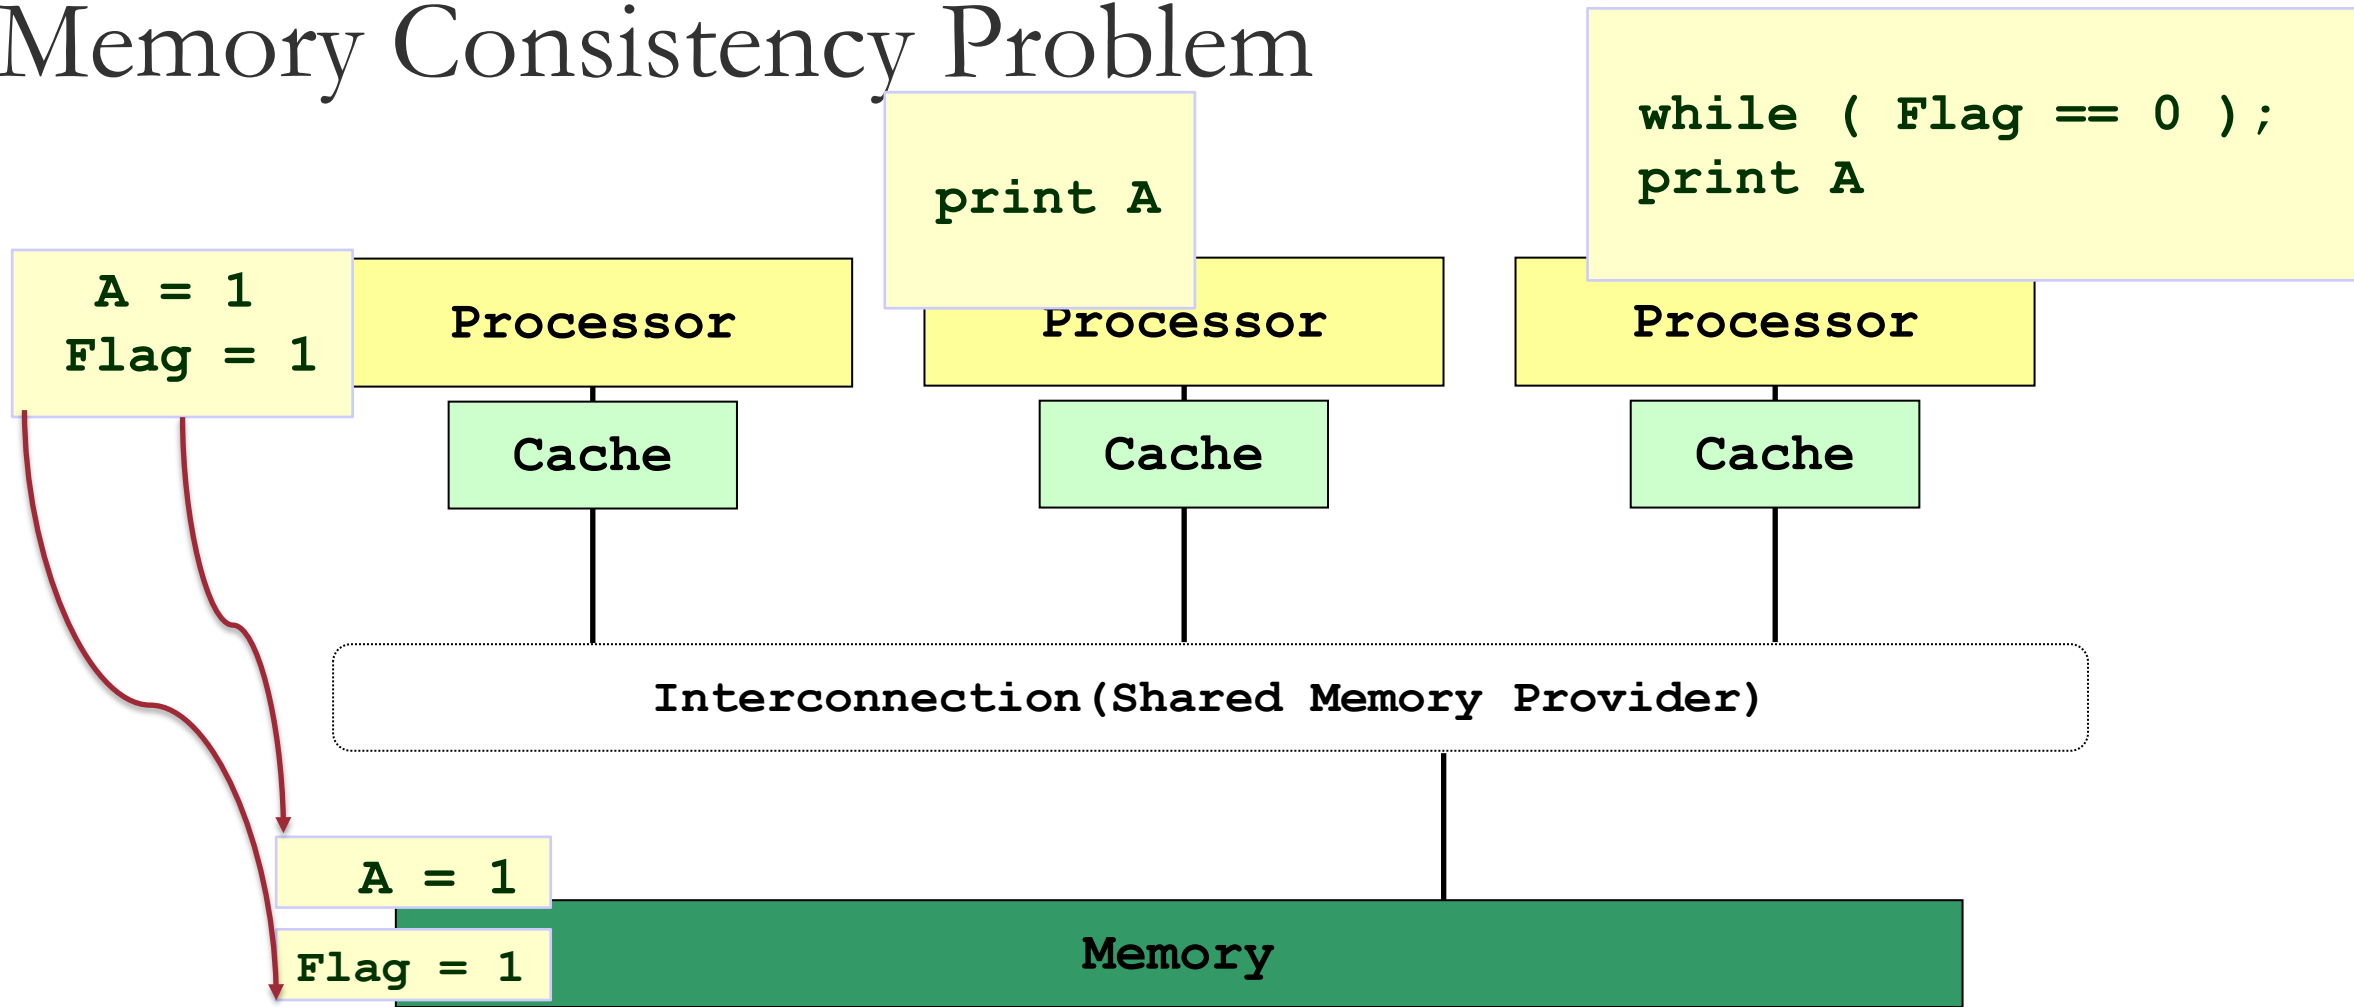
# Shared Memory System



- Parallel programs / threads access memory through the shared memory provider:
  - which maintain the illusion of shared memory
- Program is unaware of the actual hardware memory architecture
- Data exchanges between nodes
  - ➔ shared variables




# Memory Consistency Problem



# Recap: why do modern processors have cache?


- Processors run efficiently when data is resident in caches
  - Caches reduce memory access latency \*

\* Caches provide high bandwidth data transfer to CPU



# Cache Coherence Problem

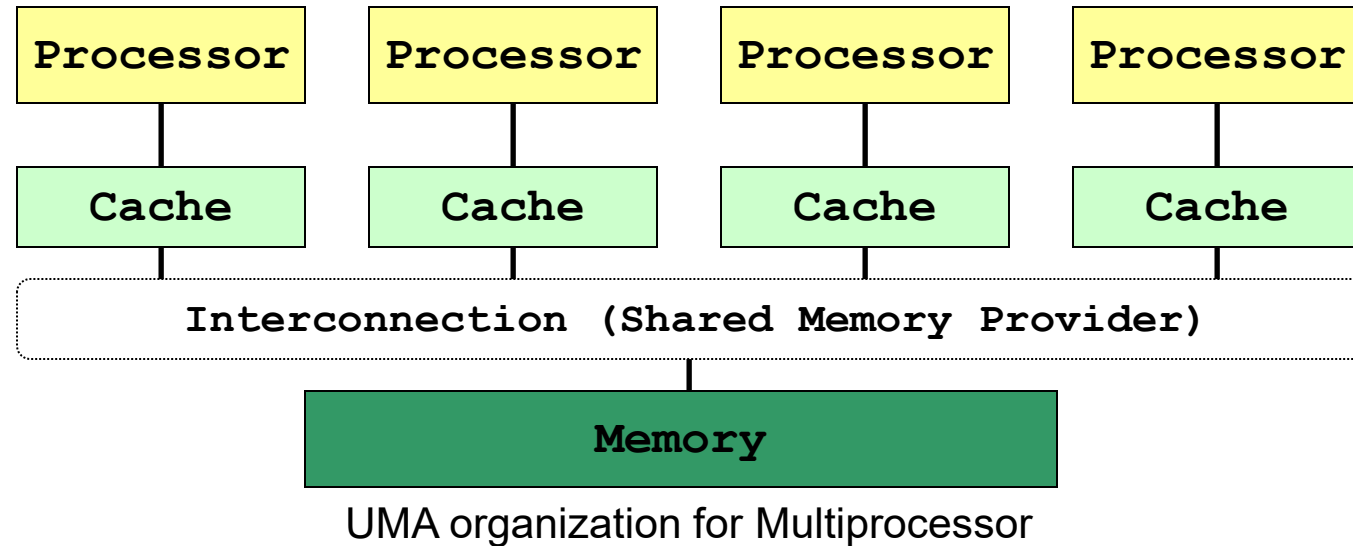
- Multiple copies of the same data exists on different caches
- Local update by processor → Other processors may still see the unchanged data



# Further Classification – Shared Memory



- Two factors can further differentiate **shared memory systems**:
  - Processor to Memory Delay (**UMA** / **NUMA**)
    - Whether delay to memory is uniform
  - Presence of a local cache with cache coherence protocol (**CC/NCC**):
    - Same shared variable may exist in multiple caches
    - Hardware ensures correctness via cache coherence protocol

# Uniform Memory Access (Time) (UMA)



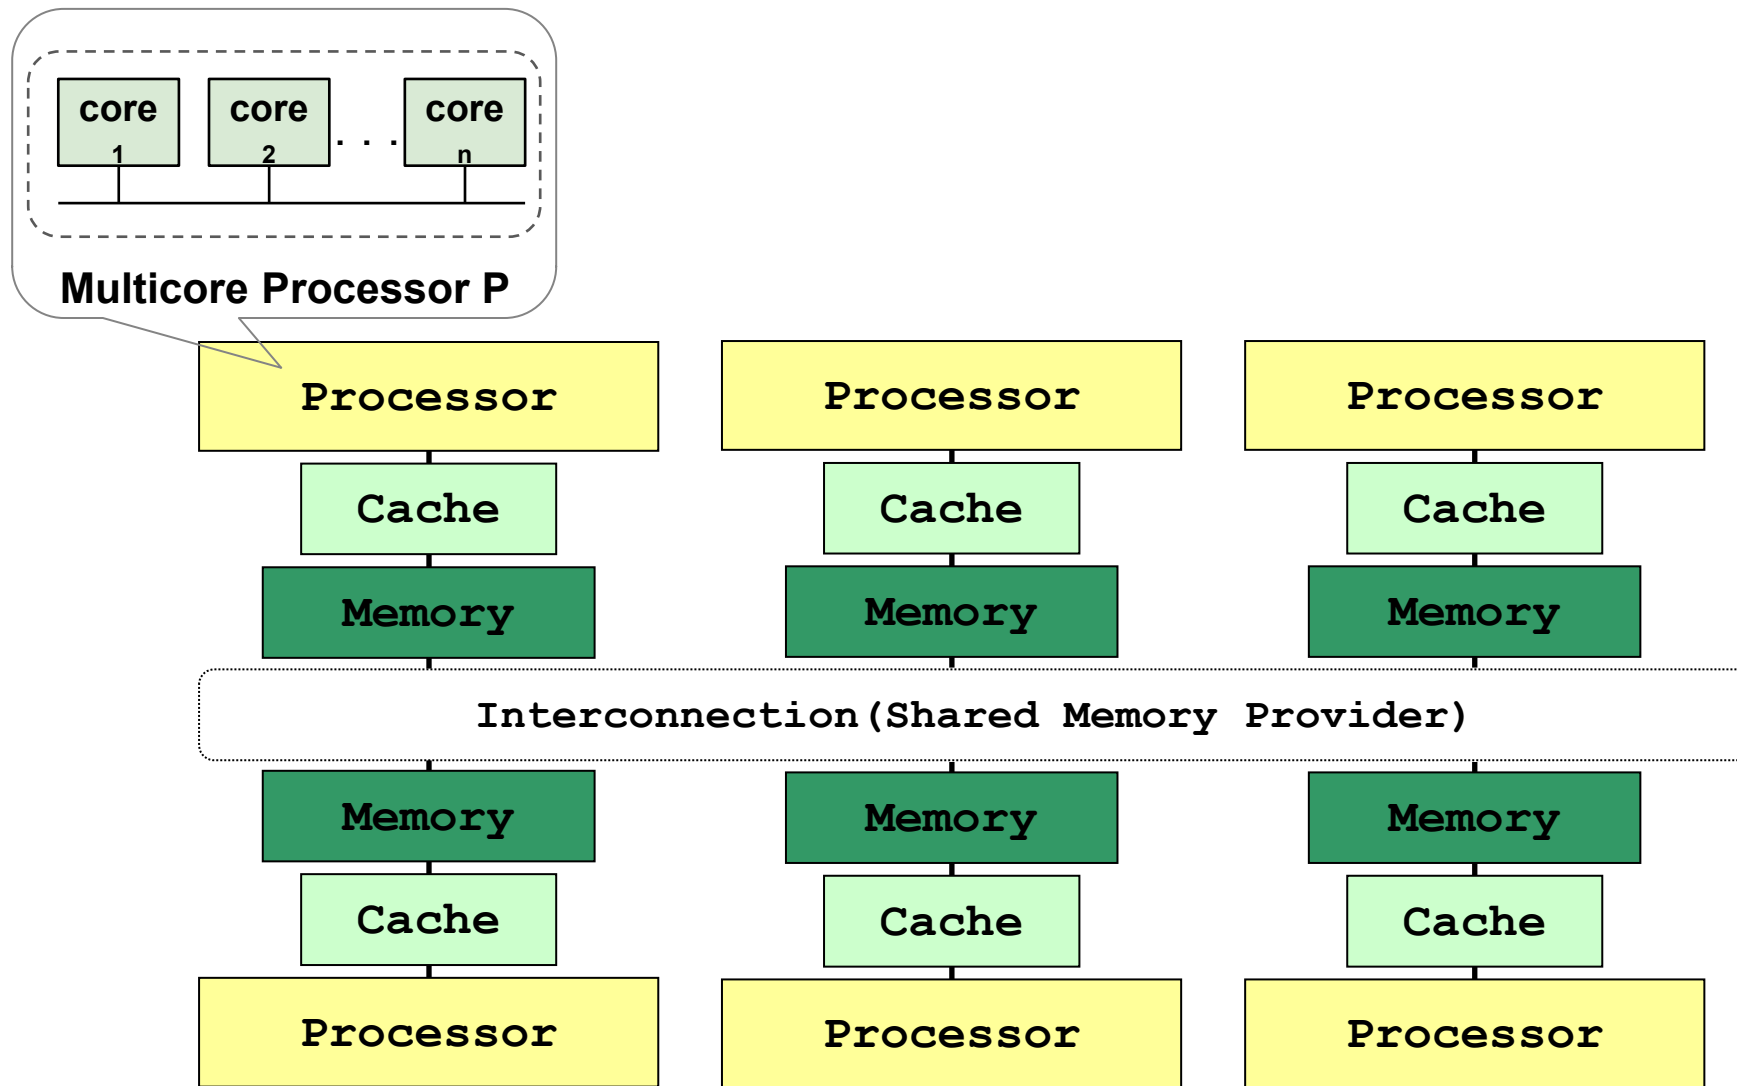
- Latency of accessing the main memory is the same for every processor:
  - Uniform access time, hence the name
- Suitable for small number of processors – due to **contention**
  - Related: Symmetric Multiprocessor (SMP)

# Non-Uniform Memory Access (NUMA)




- Physically distributed memory of all processing elements are combined to form a global shared-memory address space:
  - also called **distributed shared-memory**
- Accessing local memory is faster than remote memory for a processor
  - Non-uniform access time

# Example: Multicore NUMA



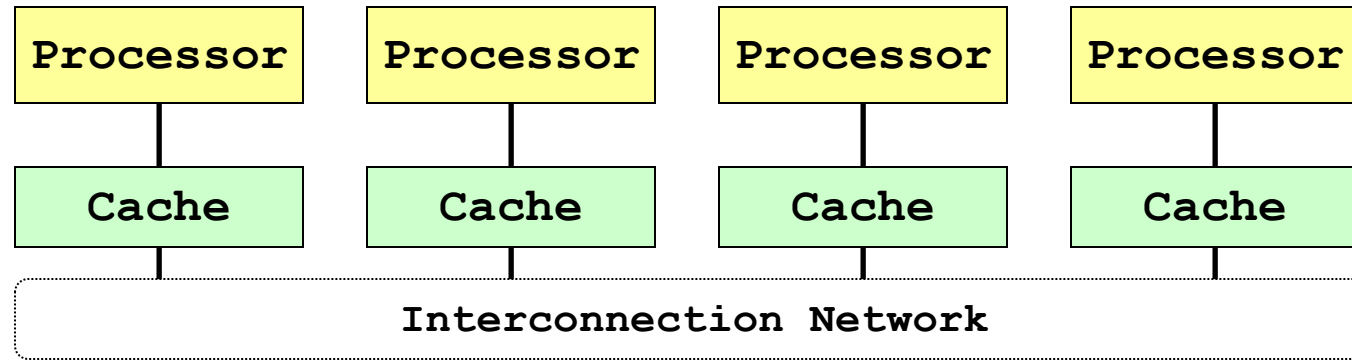
# ccNUMA



- Cache Coherent Non Uniform Memory Access
  - Each node has cache memory to reduce contention



# COMA



## ■ Cache Only Memory Architecture

- ❑ Each memory block works as cache memory
- ❑ Data migrates dynamically and continuously according to the cache coherence scheme

# Summary: Shared Memory Systems

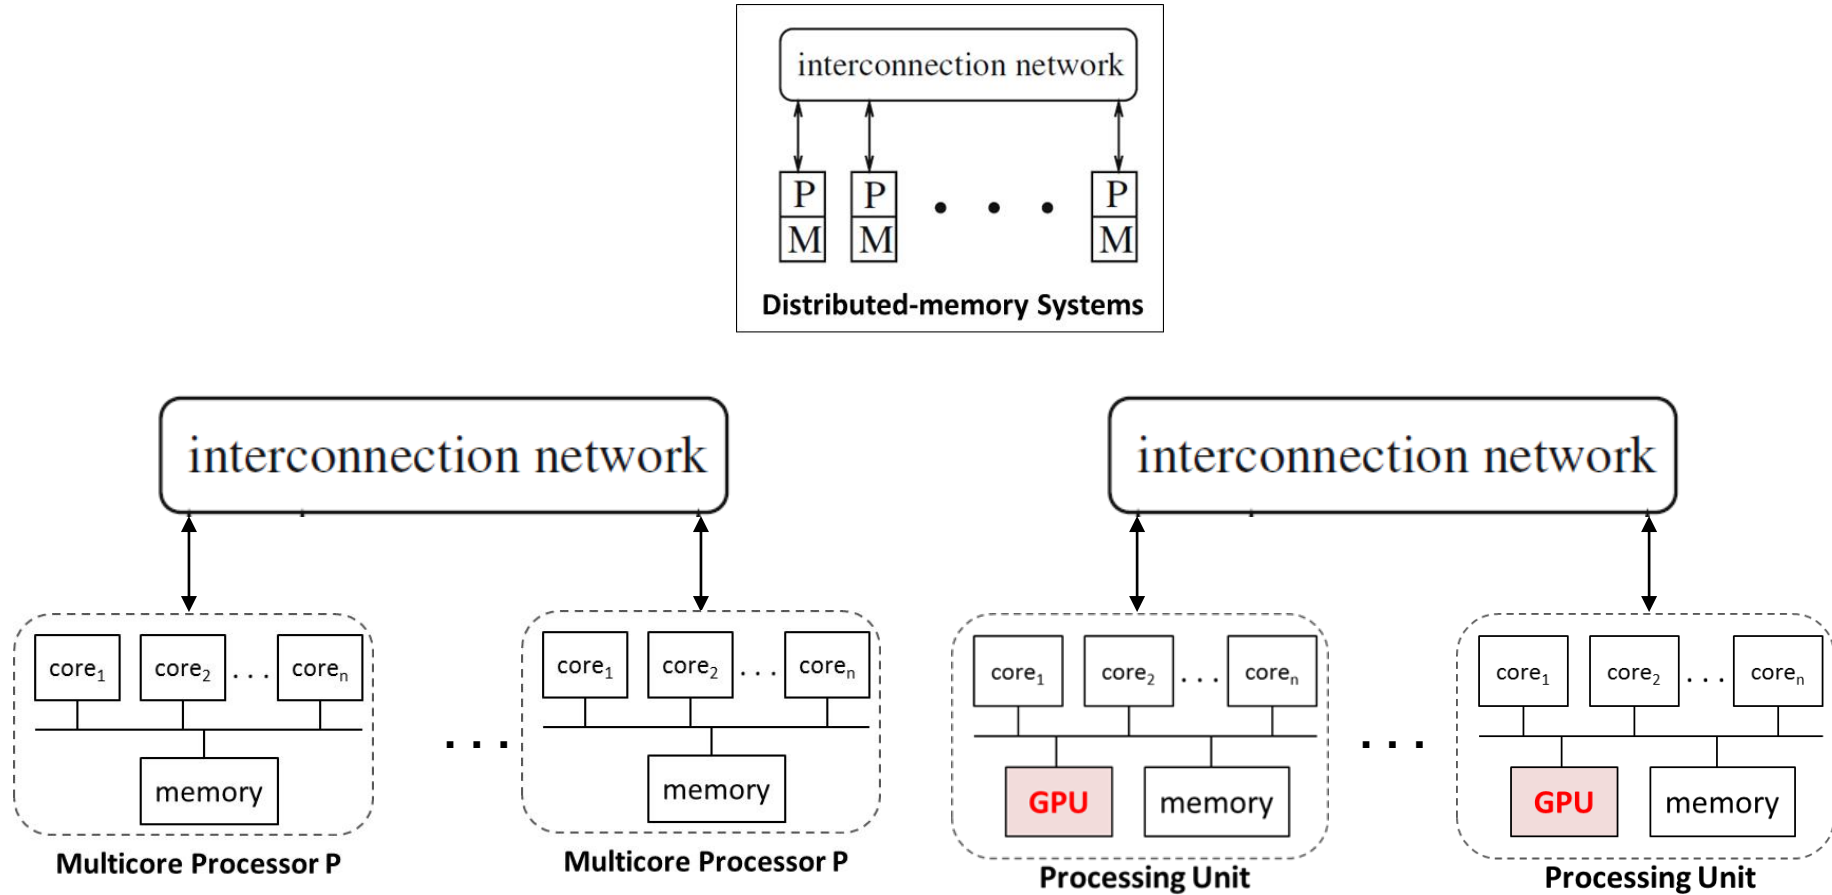
## ■ Advantages:

- ❑ No need to partition code or data
- ❑ No need to physically move data among processors → communication is efficient

## ■ Disadvantages:

- ❑ Special synchronization constructs are required
- ❑ Lack of scalability due to contention

# Hybrid (Distributed-Shared Memory)



**Hybrid with Shared-memory  
Multicore Processors**

**Hybrid with Shared-memory Multicore  
Processor and Graphics Processing Unit**

# Summary

- Goal of parallel architecture is to reduce the average time to execute an instruction
- Various forms of parallelism
- Different types of multicore processors
- Different type of parallel systems and different communication models

# Reading

- Main reference:
  - Chapter 2
- Platform 2015: Intel Processor and Platform Evolution for the Next Decade
  - Intel white paper, 2005