
Parallel Programming Models - II

Lecture 08

Outline

- Data (work) Distribution
 - 1D array
 - 2D array
- Information Exchange
 - Shared variables
 - Communication operations
- Summary


Data Distribution

- Parallel computing problems are commonly based on array of various dimensions
- Useful to study how to decompose the arrays for distribution on multiple processors
 - known as **data distribution / work distribution / decomposition / partitioning**
- For problems exhibiting data parallelism, data distribution can be used as a simple parallelization strategy

Data Distribution for 1D Arrays

- Assumptions for discussion:
 - ❑ p identical processors, P_1, P_2, \dots, P_p , and with processor rank i in $\{1, 2, \dots, p\}$
 - ❑ Array elements numbered from 1 to n
- Given a one dimensional array, common distribution patterns:
 - ❑ Blockwise data distribution
 - ❑ Cyclic data distribution

Blockwise and Cyclic Data Distribution




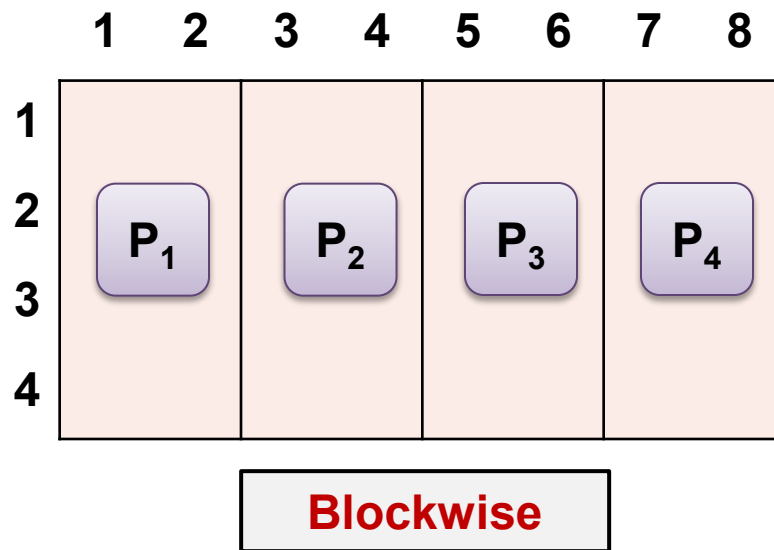
- Block size, $B = \left\lceil \frac{n}{p} \right\rceil$

- P_j takes elements $[(j-1) \times B + 1 \dots j \times B]$

- P_j takes elements $[j, j+p, \dots, j + (\left\lceil \frac{n}{p} \right\rceil - 1) \times p]$ if $j \leq n \bmod p$
 $[j, j+p, \dots, j + (\left\lceil \frac{n}{p} \right\rceil - 2) \times p]$ otherwise

Data distribution for 2D Arrays


- Combination of blockwise / cyclic distribution in one or both dimensions can be used
- One dimension distributions
 - Use the **column dimension** for illustration:



Data distribution for 2D Arrays


- One dimension distributions

- **Block-Cyclic** is a new distribution pattern
- Form blocks of size **b**, then perform cyclic (round robin) allocation



Two Dimensional Distributions

- Processors are **virtually organized into 2D mesh of $R \times C$** :
 - i.e. each Processor now has a row and column number:



- **Checkerboard** distribution can then be applied:
 - **Blockwise**: elements split into blocks along both dimensions depending on R and C
 - **Cyclic**: cyclic assignment of elements according to processor mesh
 - **Block-Cyclic**: elements split into $b_1 \times b_2$ size blocks, then cyclical assignment to processors

Checkerboard Distribution

	1	2	3	4	5	6	7	8
1	P_1				P_2			
2								
3	P_3				P_4			
4								

Blockwise

	1	2	3	4	5	6	7	8
1	P_1	P_2	P_1	P_2	P_1	P_2	P_1	P_2
2	P_3	P_4	P_3	P_4	P_3	P_4	P_3	P_4
3	P_1	P_2	P_1	P_2	P_1	P_2	P_1	P_2
4	P_3	P_4	P_3	P_4	P_3	P_4	P_3	P_4

Cyclic

	1	2	3	4	5	6	7	8	9	10	11	12
1	P_1	P_2	P_1	P_2	P_1	P_2	P_1	P_2	P_1	P_2	P_1	P_2
2	P_1	P_2	P_1	P_2	P_1	P_2	P_1	P_2	P_1	P_2	P_1	P_2
3	P_3	P_4	P_3	P_4	P_3	P_4	P_3	P_4	P_3	P_4	P_3	P_4
4	P_3	P_4	P_3	P_4	P_3	P_4	P_3	P_4	P_3	P_4	P_3	P_4

Block-Cyclic with $b_1 = 2, b_2 = 2$

Exercise: Matrix Multiplication

- To illustrate the effect of data distribution on the computation
- Assume:
 - ❑ $A \times B = C$, all matrices of $N \times N$
 - ❑ There are p processors, p will be specified
 - ❑ For each value of p , suggest a data distribution pattern for the matrices A and B

Exercise: Matrix Multiplication

1. $1 < p \leq N$, you can use $p = N$ as a start


- ❑ A distributed as _____
- ❑ B distributed as _____
- ❑ Each processor calculate _____

2. $p = N^2$

- ❑ A distributed as _____
- ❑ B distributed as _____
- ❑ Each processor calculate _____

Exercise: Heat Transfer Simulation

- A simplistic simulation of heat transfer on a metal plate
- The metal plate is modeled as:
 - 2D integer array
 - Each integer represent the temperature of a "point" on the plate
- The temperature is calculated iteratively:
 - Temperature of a point = Average of the top, left, right and down points



Exercise: Heat Transfer Simulation

- If we have a $N \times N$ metal plate and p processor, where $p < N$:
 - Suggest at least two data distribution patterns and discuss their pro / cons

I'll trade my B for your A

INFORMATION EXCHANGE

Information Exchange

■ Purpose

- Information exchange between the executing processors is necessary for controlling the coordination of different parts of a parallel program execution

■ Shared address space

- use **Shared variables**

■ Distributed address space

- use **Communication operations**

Shared Variables

- Shared memory programming models assume a global memory accessible by all processors
 - ➔ Information exchange through shared variables
 - ➔ Need synchronization operations for safe concurrent access
- Flow of control abstractions
 - ➔ processes or threads
- Each thread:
 - Executed by one processor or one core in multicore processors
 - Have shared variables and may have private variables

Synchronized Access

- **Race condition:** multiple threads accessing (read and write) the same shared variable
 - Computation result depends on the execution order of threads
 - May lead to **non-deterministic** behavior
 - Can be avoided using **critical section** mechanism
- **Critical section:**
 - A program part in which concurrent access should be avoided
 - i.e. only one thread can execute at any point in time
 - Use **mutual exclusion (mutex)** to provide critical section

Example: OpenMP

■ Race condition:

```
void main () {  
    int count = 0;  
    #pragma omp parallel  
    {  
        count = count + 1; // race  
    }  
    printf("count = %d\n", count);  
}
```

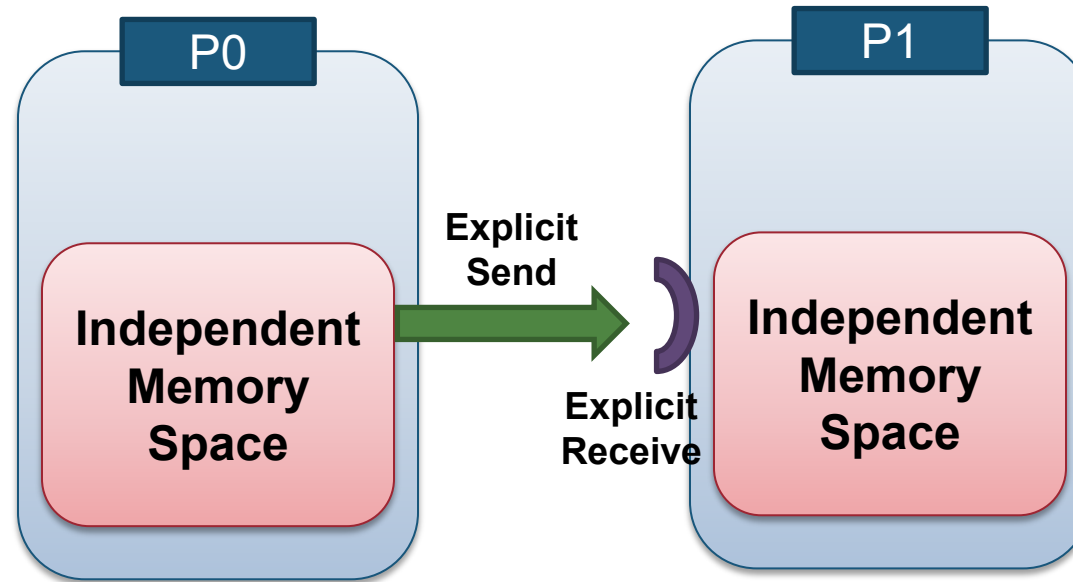
■ Mutual exclusion:

```
void main () {  
    int count = 0;  
    omp_lock_t lock;  
    omp_init_lock(&lock);  
  
    #pragma omp parallel  
    {  
        omp_set_lock(&lock);  
        count = count + 1;  
        omp_unset_lock(&lock);  
    }  
    printf("count = %d\n", count);  
}
```

Communication Operations

- Distributed memory programming models assume disjoint memory space:
 - Exchange of data between processors through **dedicated communication operations**
- One common communication model send / receive messages between participating processors:
 - known as **message-passing programming model**
- Two main types of data exchange:
 - **point-to-point** and global communication

Principles of Message Passing Model



- Data explicitly partitioned for each process
- All interaction requires both parties to participate
- ➔ Programmer has to explicitly express parallelism

Principles of Message Passing Model

- **Loosely synchronous** paradigm:
 - Tasks or subsets of tasks synchronize to perform interactions
 - Between these interactions, tasks execute completely **asynchronously**

So, you talk, I talk?

COMMUNICATION PROTOCOLS

Send and Receive Operations

```
a = 100;  
send(&a, P1);  
a = 0;
```

Process P0

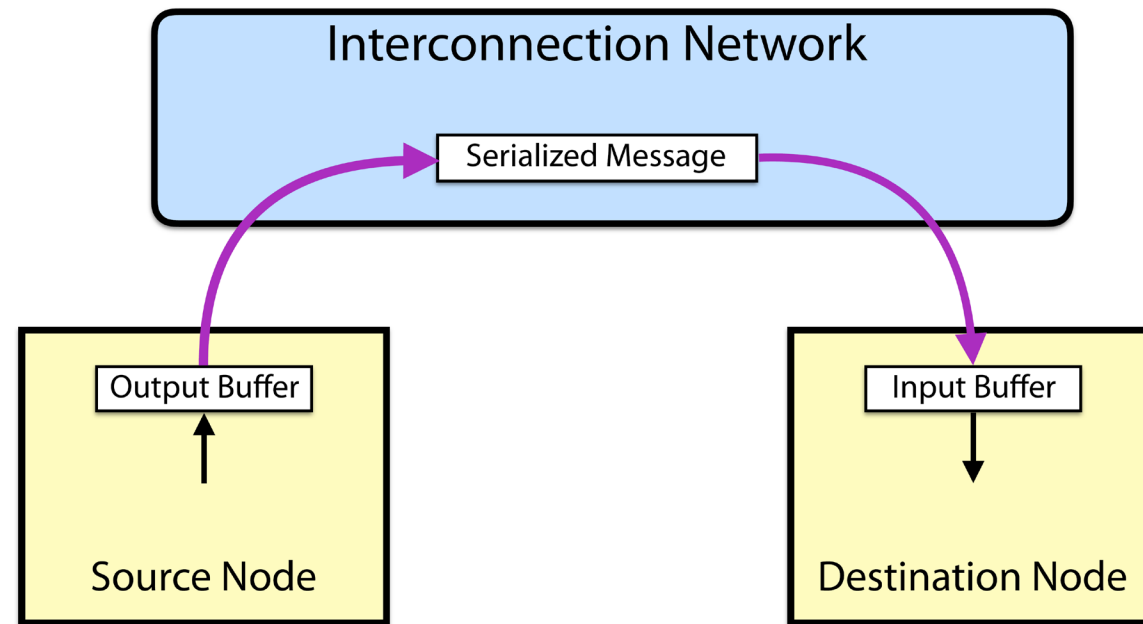
```
receive(&a, P0);  
printf("%d\n", a);
```

Process P1

- Semantic of the send():
 - The value received by P1 should be 100
- Motivates the design of the underlying communication protocols

Point-to-point Communication (Buffered)

- In a distributed memory system, over a network
 - One-way transfer



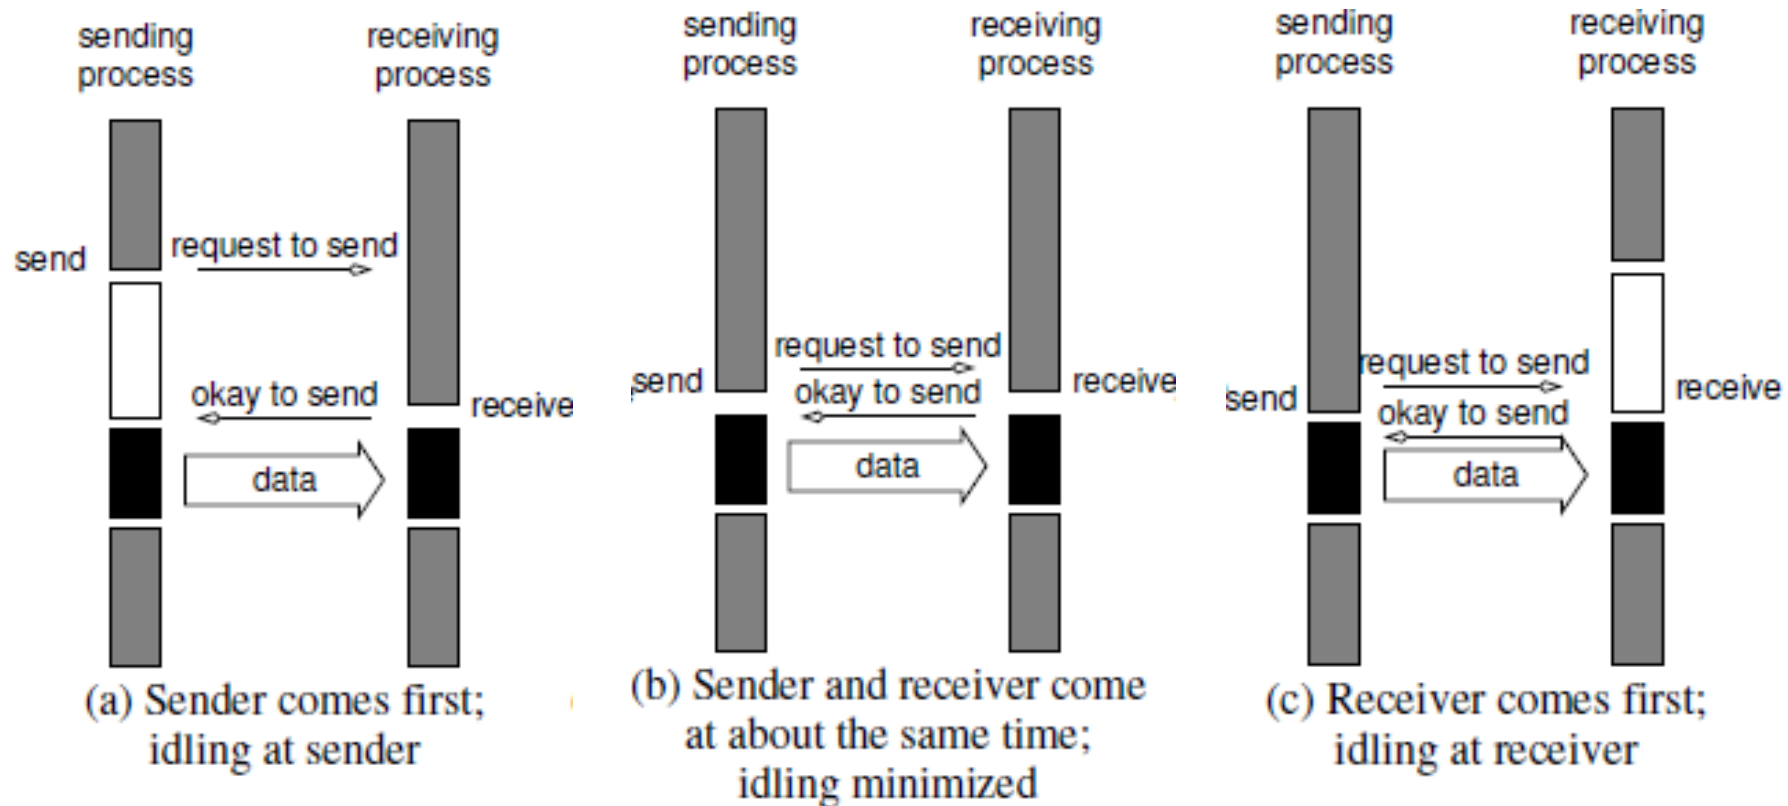
Send and Receive **Protocols Possibilities**

	Blocking Operations	Non-Blocking Operations
Buffered	Sending process returns after data has been copied into communication buffer	Sending process returns after initiating the transfer to buffer. This operation might not be completed on return.
Non-buffered	Sending process blocks until matching receive operation has been encountered.	
	Send and receives semantics assured by corresponding operation.	Programmer must explicitly ensure completion of the operation by polling.

Non-Buffered + Blocking Operations

- Send operation blocks until it is **safe** to do so
 - "Safe" refers to the integrity of the data to be sent
- Non-buffered blocking send:
 - The operation blocks until the matching receive has been performed by the receiving process
 - Idling and deadlocks are major issues with non-buffered blocking sends

Non-Buffered + Blocking Operations

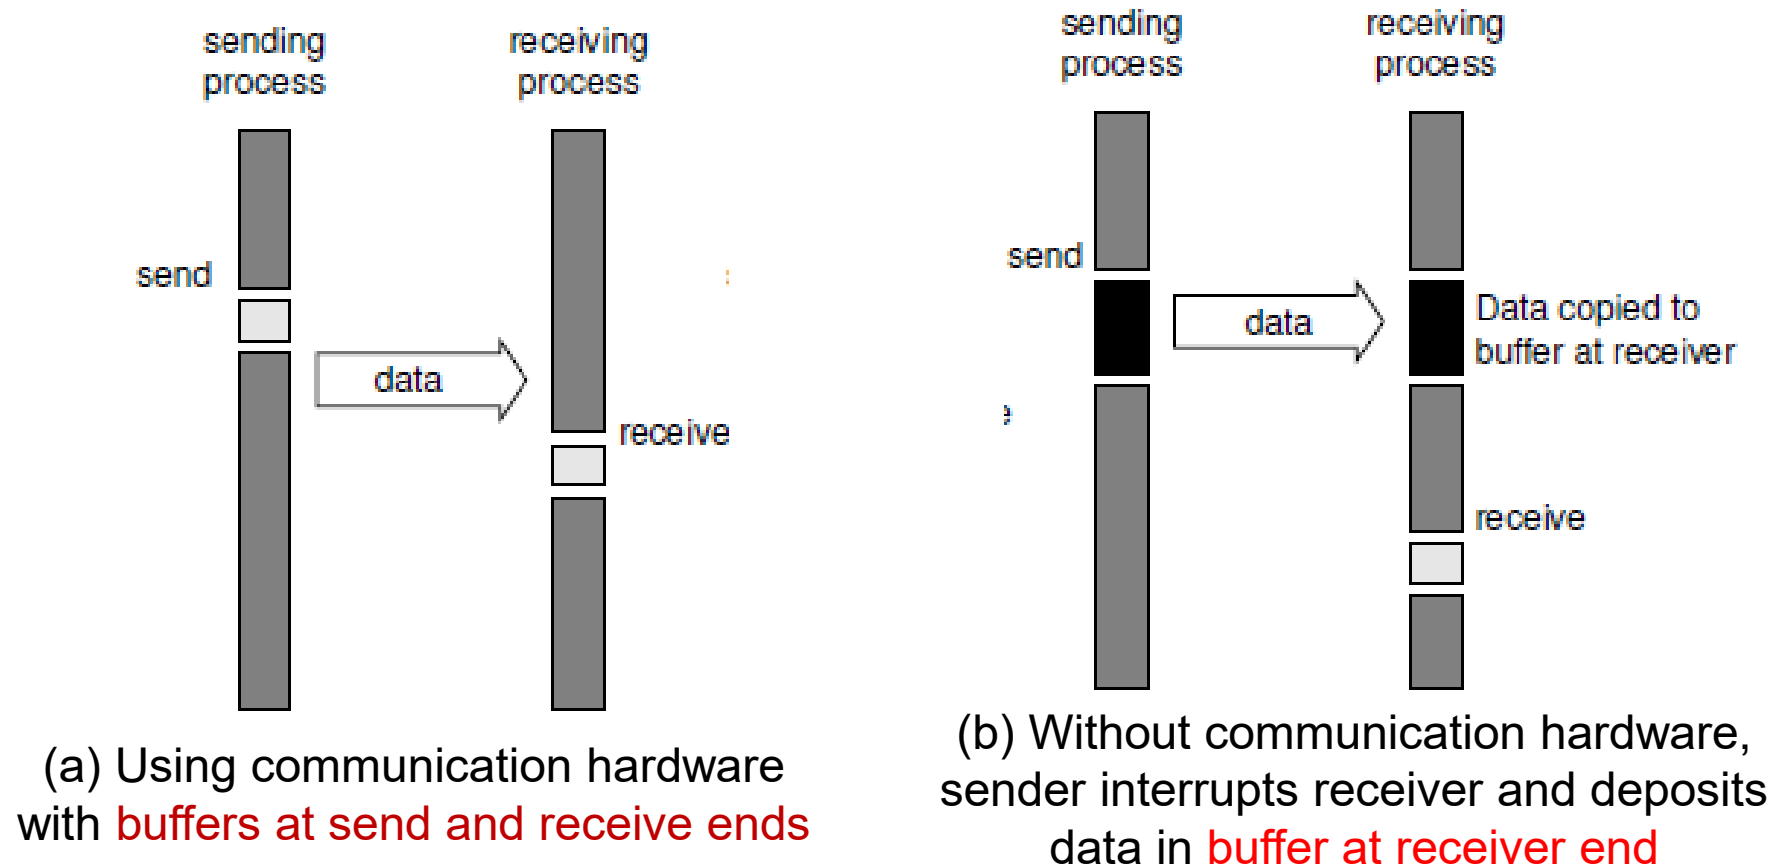


- Considerable idling overheads
 - Due to the mismatch in timing between sender and receiver

Buffered + Blocking Operations

- To reduce idling overhead:
 - Utilize **buffers** at both ends
- **Sender** simply copies the data into the designated buffer and **returns** after the copy operation has been completed
- **Receiver** similarly buffered the incoming data
- Buffering trades off **idling overhead** for **buffer copying overhead**

Buffered + Blocking Operations



Blocking buffered transfer protocols

Bounded Buffer Size: Impact

```
for (i = 0; i < 1000; i++) {  
    produce(&a);  
    send(&a, P1);  
}
```

Process P0

```
for (i = 0; i < 1000; i++) {  
    receive(&a, P0);  
    consume(&a);  
}
```

Process P1

- What if consumer was much slower than producer?
 - Think "behind the scene"....

Deadlock

- Deadlocks are still possible with buffering since receive operations block:

```
receive (&a, P1) ;  
send (&b, P1) ;
```

Process P0

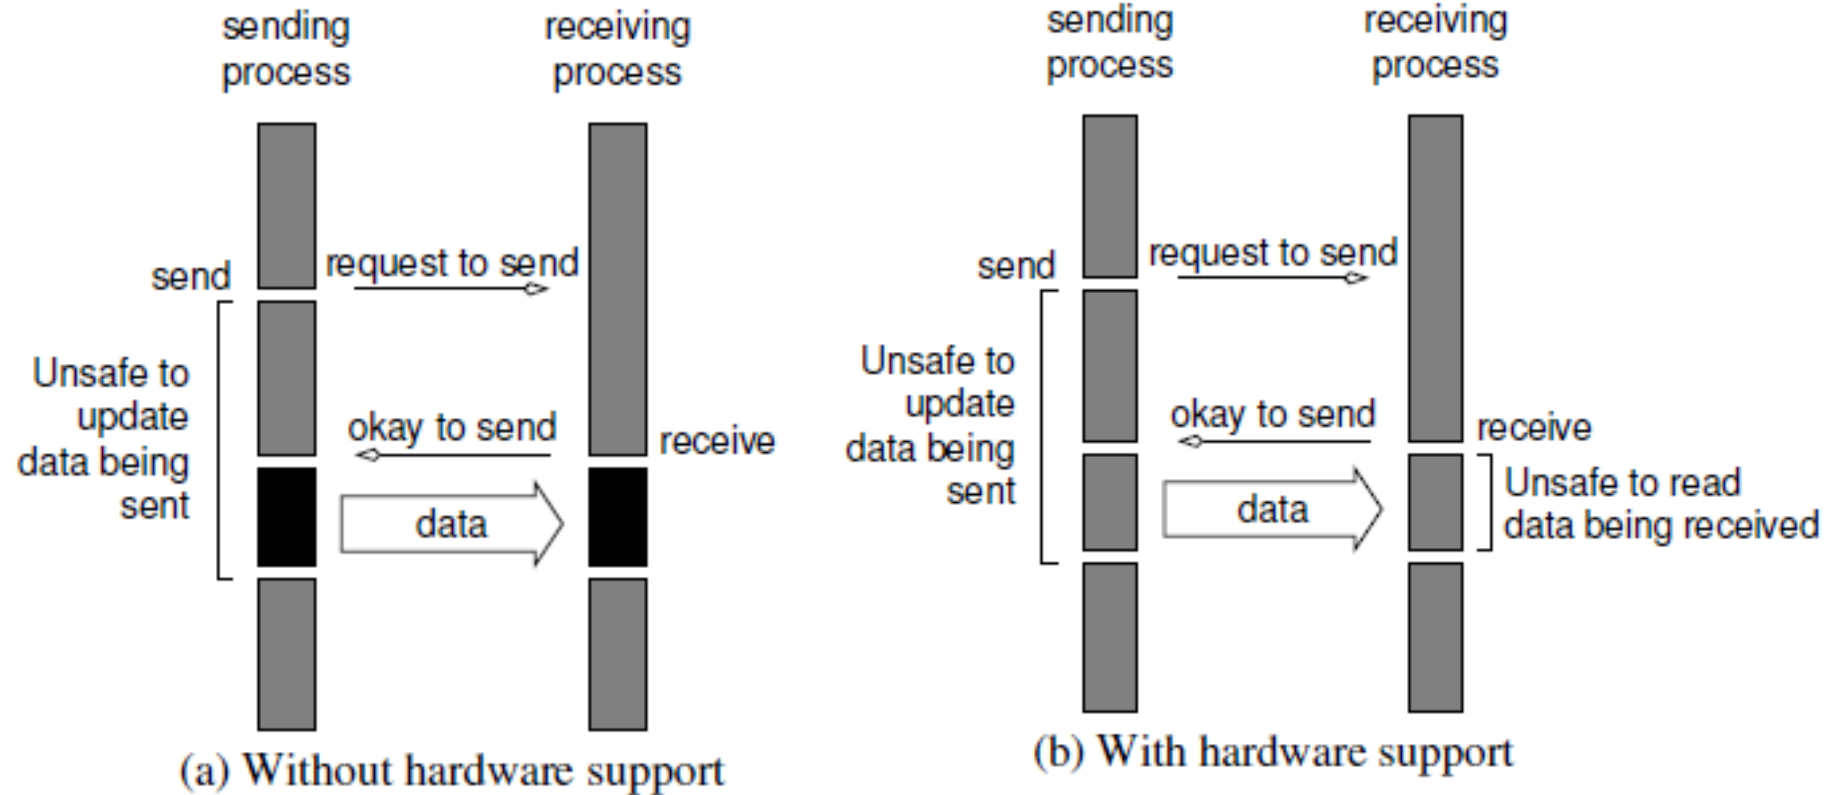
```
receive (&a, P0) ;  
send (&b, P0) ;
```

Process P1

Non-Blocking Operations

- Send / Receive returns before it is semantically safe
 - Non-blocking operations are generally accompanied by a ***check-status*** operation
 - The programmer must ensure semantics of the operations
- When used correctly, these primitives are capable of overlapping communication overheads with useful computations
- Message passing libraries typically provide both blocking and non-blocking primitives

Non-Blocking + Non-Buffered Operations



Non-blocking non-buffered send and receive operations

Semantic of Send/Receive Operations

Local view

Blocking

Return from a library call indicates the user is allowed to reuse resources specified in the call

Non-blocking

A procedure may return before the operation completes, and before the user is allowed to reuse resources specified in the call

Global view

Synchronous

Communication operation does not complete before both processes have started their communication operation

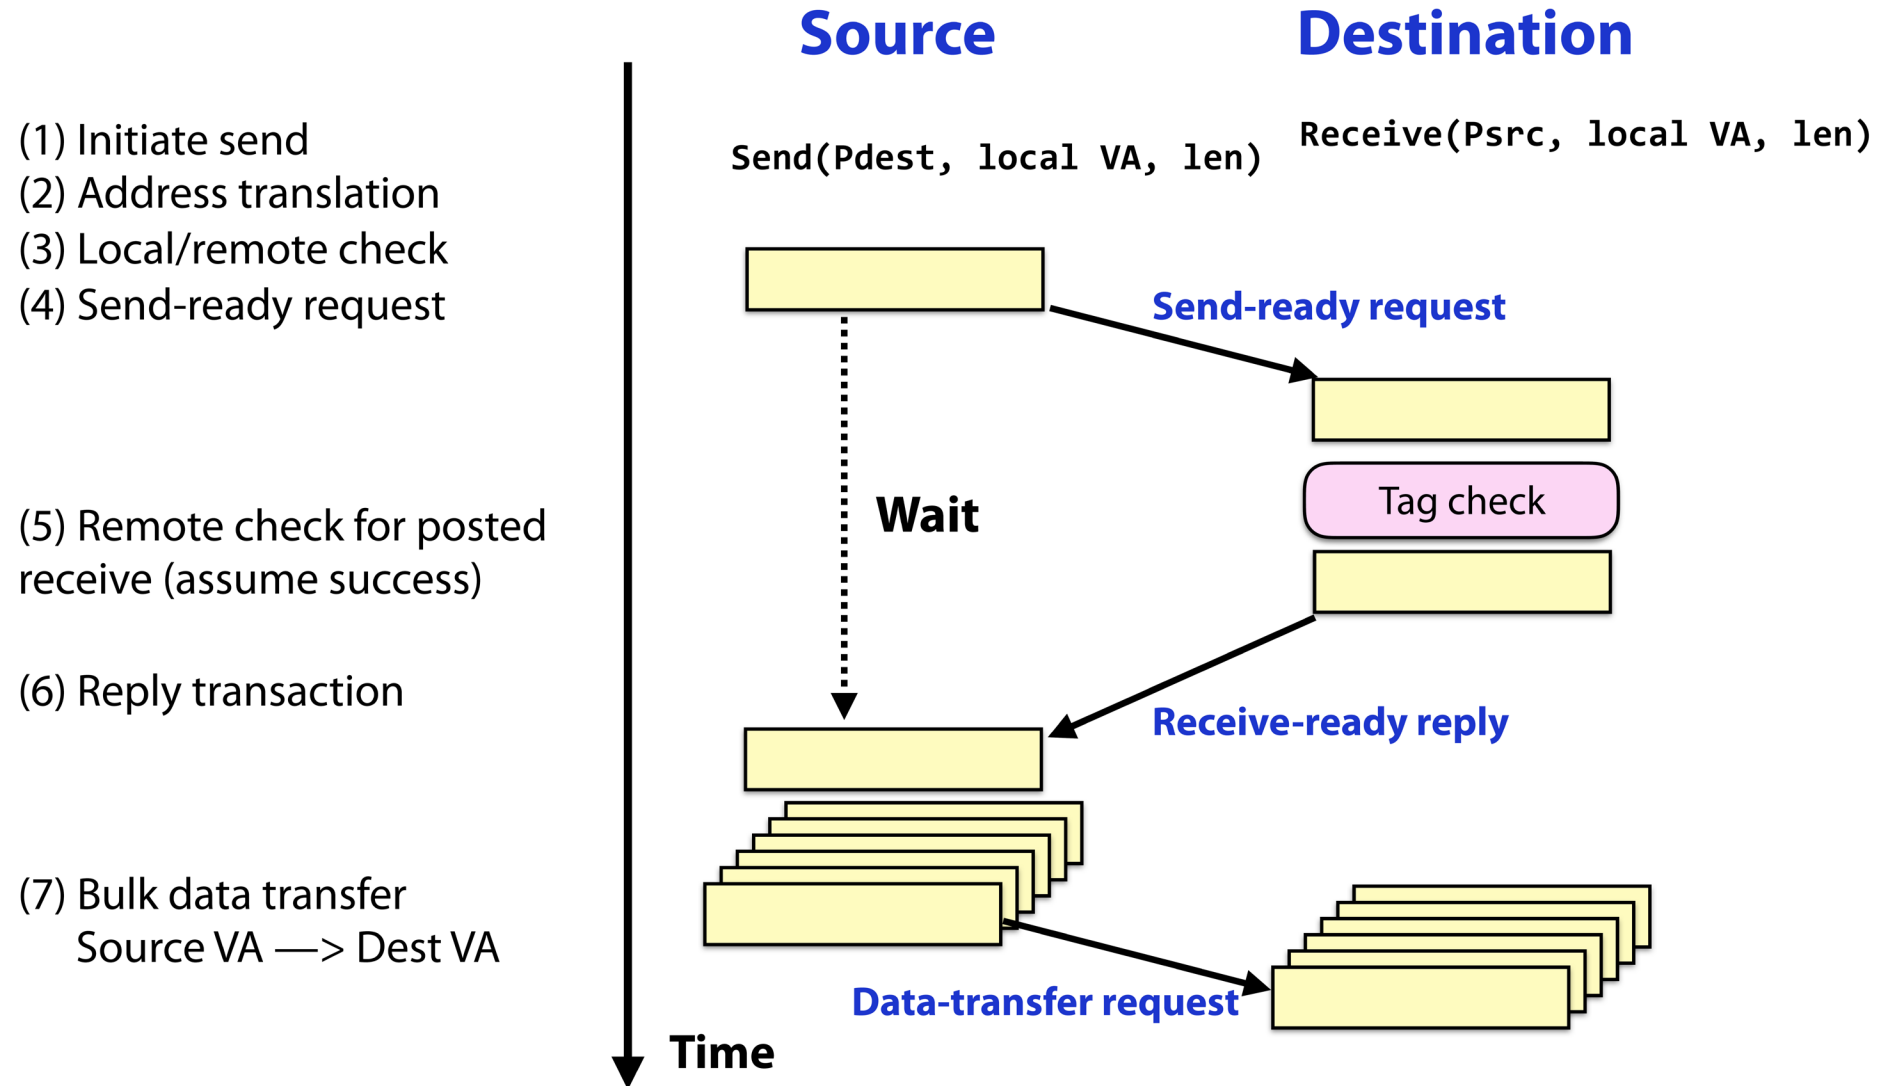
Asynchronous

Sender can execute its communication operation without any coordination with the receiver

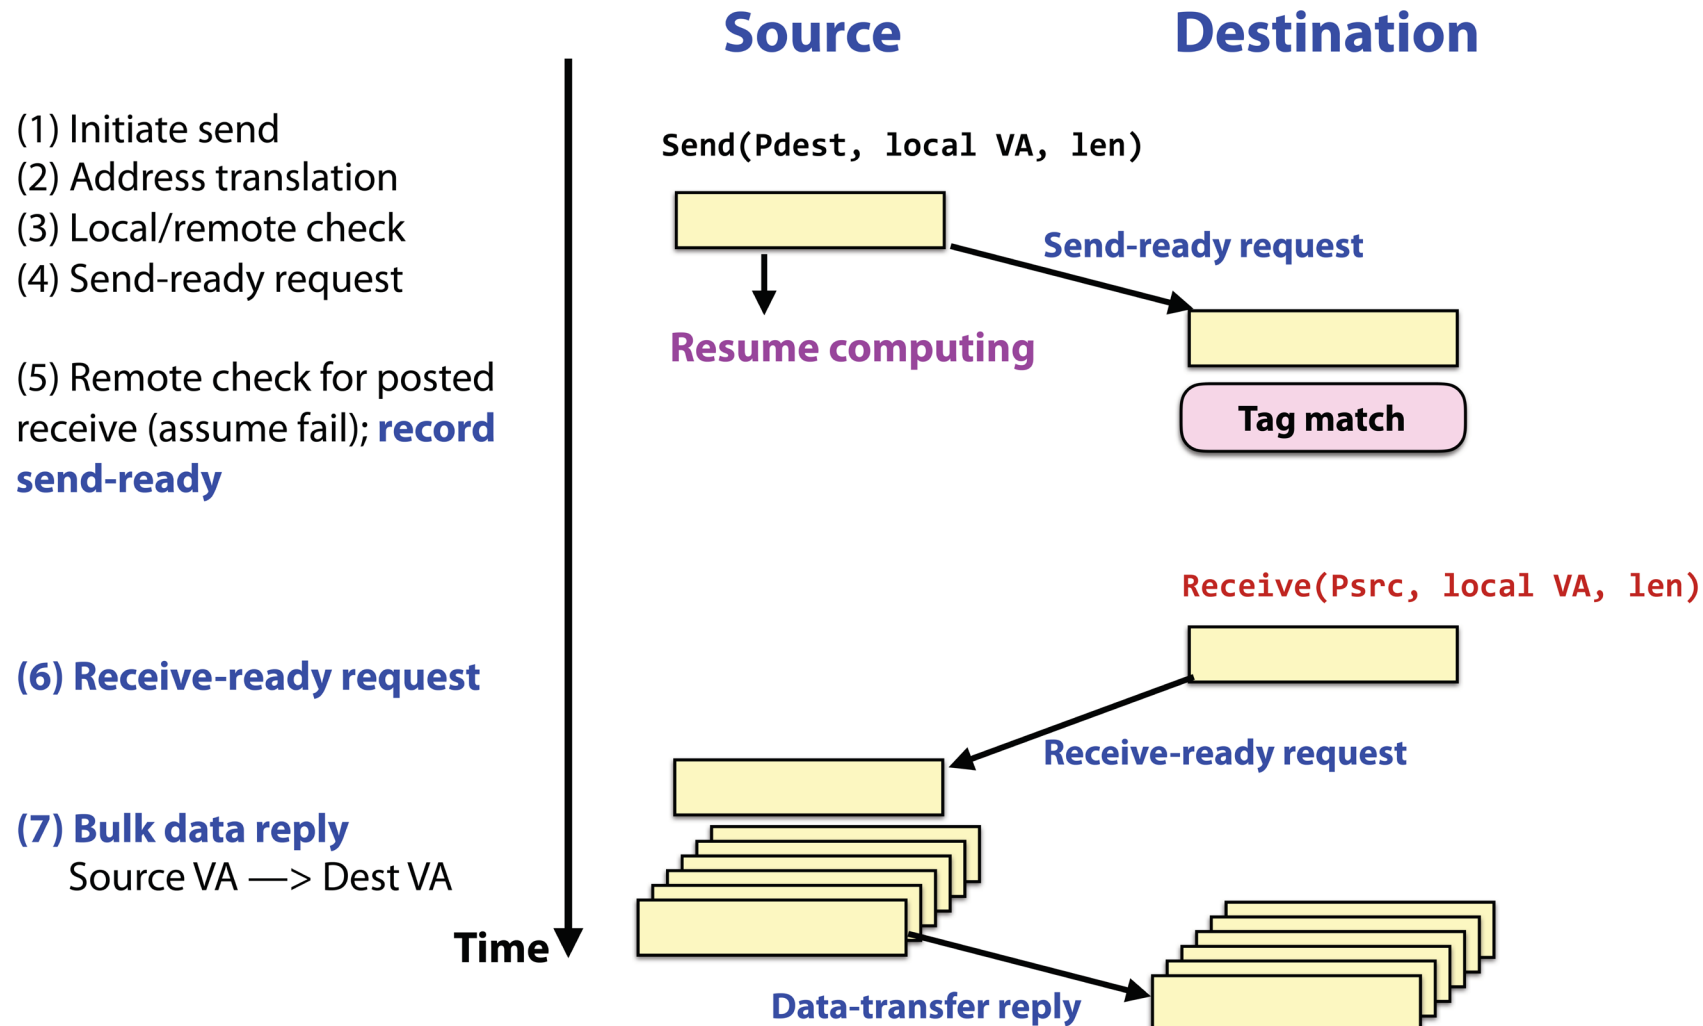
Implementation Options

- Synchronous:
 - Send completes after matching receive and source data sent
 - Receive completes after data transfer complete from matching send
- Asynchronous:
 - Send completes after send buffer may be reused

Synchronous Communication



Asynchronous Communication



Summary of Communication Protocols

1. **Overhead** of guaranteeing semantic correctness:

- ❑ **Idling** (non-buffered)
- ❑ **Buffer management** (buffered)

2. **Side effect:**

- ❑ Safe and easier programming (**blocking**)
- ❑ Hide communication overhead (**non-blocking**)

3. **Local or global view:**

- ❑ Synchronous vs. asynchronous communication

Summary

- Data distribution
- Information exchange for shared and distributed address space
- Reading
 - Main textbook: chapter 3