

CS4231  
Parallel and Distributed Algorithms

Lecture 1  
Instructor: YU Haifeng

## A bit of self-introduction...

- <https://www.comp.nus.edu.sg/~yuhf>

# Module Overview

- Module homepage ready on LumiNUS
  - Check module homepage often!
- Prerequisite:
  - CS3230 Design and Analysis of Algorithms or CS3210 Parallel Computing --- CS3230 is more important
  - Can **NOT** be waived – will not honor request for waiver
- What is the module about:
  - Designing parallel/distributed algorithms, and proving their correctness/properties
  - This is a theory module, **involving mostly with proofs and theorems**

# Is This Module Hard?

- This is an **elective module catering to students with strong interests and background in the subject**
- While having the same theoretical nature as the **compulsory module CS3230**, this module will be much harder than CS3230
  - Analogy:  
English modules taken by English major students  
vs.  
English modules taken by non-English-major students
- This module will be taught in very different ways from CS3230

# Is This Module Hard?

- This is a 4000-level module:
  - Not recommended for 1<sup>st</sup>, 2<sup>nd</sup>, or 3<sup>rd</sup> year students – but if you think you can handle this challenging module, you can still take it but you are encouraged to chat with me before enrolling
- Overall, this is designed to be a challenging module
  - If you did not enjoy CS3230, you will not enjoy this module
  - If you enjoyed CS3230, you may still not enjoy this module if your formal/mathematical skill and abstract thinking ability are not strong
  - In previous offerings, it was not unusual that some students who did well in CS3230 had difficulty handling this module

# Why This Module Is Important

- This module is one of the most theoretical modules in SoC
  - Theoretical = impractical = useless?
  - No. **Actually it directly relates to your career...**
- Reason #1: Deep understanding of distributed algorithms **distinguishes you from others**
  - (Unfortunately) Everyone's career has a lot to do with competing with other people
  - Good programming skills ensure that you don't lose to others; deep understanding helps you to win

# Why This Module Is Important


- Reason #2: Understanding foundational concepts **enables you to learn new material in the future**
  - Giving you some meat vs. giving you a hunting gun
  - Computer science is particularly fast-changing – don't expect what you learn today can be directly applicable 10 years later
- Reason #3: Foundational concepts are harder to grasp
  - This can be **your only opportunity** to learn these concepts in your whole career

## An Example

- Boss asks you to solve the following simple problem
  - Two nodes A and B, each has a starting value of 0 or 1
  - Each needs to output a single value 0 or 1
  - They can communicate but messages can be lost
  - Goal: A and B should output the same value. Specifically,
    - If A and B both start with 0, they should both output 0.
    - If A and B both start with 1 and if no messages are lost, they should both output 1.
    - If A and B both start with 1 and if some message is lost, they should **either** both output 1 **or** both output 0.
    - If A and B start with different values, they should **either** both output 1 **or** both output 0.



To be, or not to be, that is the question



The decision is up to you...

## What students in previous years say

- “A very difficult module and I found it very hard to understand”
- “An interesting module that touch on parallel algorithms and distributed system, but concepts are a little too difficult sometimes”
- “The concepts are interesting, but really difficult to digest.”
- “Better to prove the algorithm in a more straight forward way. Too many lemma used, which make student hard to follow when using the lemma.”

## Relation with CS3211

- CS3211
  - PARALLEL AND CONCURRENT PROGRAMMING
- CS3211 is more programming oriented
- CS4231 is more about designing algorithms and proving their correctness/properties
- Overlap is small

# Teaching Format

- Wednesday 6:30pm-9:30pm (not 6:30pm-8:30pm) in every lecturing week: Lecturing + tutorial (break in the middle)
  - Over Zoom: For better security, Zoom link will only be published in Luminus announcement immediately before the class
- By NUS university guideline
  1. One-hour lesson: 45-minute teaching + concludes 15 minutes before the end of the hour
  2. Two-hour lesson: 90-minute teaching + 5-minute break in the middle + concludes 25 minutes before the end of the second hour
  3. Three-hour lesson: 135-minute teaching + 20-minute break in the middle (potentially split into two smaller breaks) + concludes 25 minutes before the end of the third hour
- You are required to read the materials to be covered before each lecture – Otherwise you won't follow

# Module Format

- Office hours:
  - Wednesday 2:00pm to 4:00pm every lecturing week, online via skype -- my skype id is “live:.cid.84882ce4ec7e3e6c”
  - You are also welcome to approach me at other times or email me
- Weekly homework
- No systematic programming homework/exercise
  - To avoid overlapping with CS3211
  - To avoid excessive workload in this module
  - But you are still encouraged to implement algorithms learnt

## 2 Compulsory Textbooks + 2 Reference Textbooks

- "Distributed Algorithms for Message-Passing Systems" by Michel Raynal
  - 2013, 1<sup>st</sup> edition, compulsory. No newer editions available.
  - Newest textbook I can find that is suitable for this module. This textbook is solid and well-written.
- "Concurrent and Distributed Computing in Java" by Vijay Garg
  - 2004, 1<sup>st</sup> edition, compulsory. No newer editions available.
  - This textbook is easier to understand, has good exercises, but old.
  - Erratum: <http://www.ece.utexas.edu/~garg/dist/jbk-corrections.txt>
- I will supplement these 2 compulsory textbooks with newer developments whenever possible.
- Why there are not many newer textbooks?
  - Theoretical and foundational materials do not actually change so often, so people do not write a lot of new textbooks on such materials...

## 2 Compulsory Textbooks + 2 Reference Textbooks

- 2 additional textbooks as “references”
  - Only for students who want to learn more materials, beyond the requirements of this module
- “Distributed Algorithms: An Intuitive Approach” by Wan Fokkink
  - 2018, 2<sup>nd</sup> edition, as reference book only.
  - This book is a bit too hard for this module, but you can refer to if you want.
  - The book has a 2018 2<sup>nd</sup> edition, which however is not very different from the 2013 1<sup>st</sup> edition, in terms of the topics that are relevant to this module.
- “Distributed Algorithms” by Nancy Lynch
  - 1996, 1st edition, as reference book only.
  - This is the "bible" on distributed algorithms. It is very good, but is also much harder to understand. If you want, you can use this book as a reference.

# Grading Policy

- 38% mid-term exam and 60% final exam
  - Both exams cover whatever have been taught by the time of the exam
  - Mid-term: Closed book
  - Final: Open book
- 2% mock exam
- Homework assignments do not directly contribute to final score
  - BUT some exam questions will be variants of homework questions
- Cheating ABSOLUTELY not tolerated and will be reported to dept and school



## Mid-term Exam Date

- 6:30pm-9:30pm on Wed 3 Mar 2021 in class
  - Will be invigilated as an E-exam, and will follow School's SOP on E-exams (<https://mysoc.nus.edu.sg/academic/e-exam-sop-for-students/>)
- Same policy as the final exam:
  - Not showing up for the mid-term exam = zero mark
  - Showing up later for the mid-term exam = less time to work on the exam (no extra time will be given)

## Mock Exam Date

- **Wed 17 Feb 2021** in class (likely to be second part of the class)
  - Same format, policy, and invigilation as the mid-term exam

# You MUST attend mid-term exam and mock exam

- **Being able to be present for both the mid-term exam and the mock exam is prerequisite for taking this module.**
- If you feel you have trouble showing up for the mid-term exam or the mock exam, let me know now...

# Class Participation and Penalty for Non-participation

- The School and also myself encourage class participation
  - But hard to incorporate into assessment (fairness issues)
- My approach to encourage participation:
  - No recording
  - If you miss any important discussions/announcements I made during lecture – you pay the price yourself
  - There will be important things that are only discussed during lecture, and not via email or LumiNUS or other venues
- Examples of questions that I will not answer:
  - I was at a party last night so I didn't attend the lecture, did you say anything important in the lecture?

# Material Covered Today and Next Week

- Today: Chapter 2 “Mutual Exclusion Problem”
  - No tutorial today
- Next week:
  - Chapter 3 “Synchronization Primitives”
  - Read before you come to class next week

Break



# Mutual Exclusion Problem

- Context: Shared memory systems

- Multi-processor computers
- Multi-threaded programs

- Shared variable  $x$

- Initial value 0

- Program  $x = x+1$

<i>process 0</i>	<i>process 1</i>
read $x$ into a register (value read: 0)	
increment the register (1)	
write value in register back to $x$ (1)	
	read $x$ into a register (value read: 1)
	increment the register (2)
	write value in register back to $x$ (2)

# Mutual Exclusion Problem

- Context: Shared memory systems
  - Multi-processor computers
  - Multi-threaded programs
- Shared variable  $x$ 
  - Initial value 0
- Program  $x = x+1$

<i>process 0</i>	<i>process 1</i>
read $x$ into a register (value read: 0)	
increment the register (1)	
	read $x$ into a register (value read: 0)
	increment the register (1)
write value in register back to $x$ (1)	
	write value in register back to $x$ (1)



# Critical Section

Critical Section  
(also called  
Critical Region)

RequestCS(int processId)

Read x into a register

Increment the register

Write the register value back to x

ReleaseCS(int processId)

# Roadmap

- Software solutions
  - Unsuccessful attempts
  - Peterson's algorithm
  - Bakery algorithm
- Hardware solutions
  - Disabling interrupts to prevent context switch
  - Special machine-level instructions

## Attempt 1

Shared boolean variable openDoor;

//whether door is open

RequestCS(int processId) {

while (openDoor == false) {};

openDoor = false; // close door behind me

}

ReleaseCS(int processId) {

openDoor = true; // open door to let other people in

}

Both process may  
see openDoor as  
true



**Violate mutual exclusion: Two processes in critical region**

## Attempt 2

Shared boolean variable wantCS[0], wantCS[1] initialized to false

**wantCS[0] = wantCS[1] = true**

*Process 0*

```
RequestCS(0) {  
    wantCS[0] = true;  
    while (wantCS[1] == true) {};  
}
```

```
ReleaseCS(0) {  
    wantCS[0] = false;  
}
```

*Process 1*

```
RequestCS(1) {  
    wantCS[1] = true;  
    while (wantCS[0] == true) {};  
}
```

```
ReleaseCS(1) {  
    wantCS[1] = false;  
}
```

**No progress: No one can enter critical region**

## Attempt 3

Shared int turn initialized to 0

*Process 0*

```
RequestCS(0) {  
    while (turn == 1) {};  
}
```

```
ReleaseCS(0) {  
    turn = 1;  
}
```

*Process 1*

```
RequestCS(1) {  
    while (turn == 0) {};  
}
```

```
ReleaseCS(1) {  
    turn = 0;  
}
```

Starvation: Process 0 may never enter critical region again  
(There are other kinds of starvation...)

## Properties Needed

- **Mutual exclusion**: No more than one process in the critical section
- **Progress**: If one or more process wants to enter and if no one is in the critical section, then one of them can enter the critical section
- **No starvation**: If a process wants to enter, it eventually can always enter
  - Need to consider the worst-case schedule

# Peterson's Algorithm

Shared bool wantCS[0] = false, bool wantCS[1] = false, int turn = 0;

*Process 0*

```
RequestCS(0) {  
    wantCS[0] = true;  
    turn = 1;  
    while (wantCS[1] == true &&  
           turn == 1) {};  
}
```

```
ReleaseCS(0) {  
    wantCS[0] = false;  
}
```

*Process 1*

```
RequestCS(1) {  
    wantCS[1] = true;  
    turn = 0;  
    while (wantCS[0] == true &&  
           turn == 0) {};  
}
```

```
ReleaseCS(1) {  
    wantCS[1] = false;  
}
```

# Correctness Proof for Peterson's Alg.

*Process 0*

```
RequestCS(0) {  
    wantCS[0] = true;  
    turn = 1;  
    while (wantCS[1] == true &&  
           turn == 1) {};  
}
```

*Process 1*

```
RequestCS(1) {  
    wantCS[1] = true;  
    turn = 0;  
    while (wantCS[0] == true &&  
           turn == 0) {};  
}
```

**Mutual exclusion:** Proof by contradiction. (The textbook's proof is vague.)

Case 1:  $\text{turn} == 0$  when P0 and P1 are both in critical section.

Then P0 executed “ $\text{turn} = 1$ ” before P1 executed “ $\text{turn} = 0$ ”.

Hence  $\text{wantCS}[0] == \text{false}$  as seen by P1.

But  $\text{wantCS}[0]$  set to true by Process 0.

Case 2:  $\text{turn} == 1$ . Symmetric – complete yourself...



# Correctness Proof for Peterson's Alg.

*Process 0*

```
RequestCS(0) {  
    wantCS[0] = true;  
    turn = 1;  
    while (wantCS[1] == true &&  
           turn == 1) {};  
}
```

*Process 1*

```
RequestCS(1) {  
    wantCS[1] = true;  
    turn = 0;  
    while (wantCS[0] == true &&  
           turn == 0) {};  
}
```

**Progress:** Proof by contradiction and consider the value of turn when both P0 and P1 are waiting.

Case 1: turn == 0. Then P0 can enter.

Case 2: turn == 1. Symmetric – complete yourself...

# Correctness Proof for Peterson's Alg.

*Process 0*

```
RequestCS(0) {  
    wantCS[0] = true;  
    turn = 1;  
    while (wantCS[1] == true &&  
           turn == 1) {};  
}
```

*Process 1*

```
RequestCS(1) {  
    wantCS[1] = true;  
    turn = 0;  
    while (wantCS[0] == true &&  
           turn == 0) {};  
}  
ReleaseCS(1) {  
    wantCS[1] = false;  
}
```

**No starvation:** Proof by contradiction.

Case 1: If P0 waiting, then wantCS[1] = true and turn = 1.

P1 in critical region -- will exit and set wantCS[1] to false.

What if P1 wants to enter again immediately?

Case 2: P1 is waiting. Symmetric – complete yourself...

# Lamport's Bakery Algorithm

- For  $n$  processes
  - Get a number first
  - Get served when all people with lower number have been served
- Two shared arrays of  $n$  elements
  - `boolean choosing[i] = false; // process  $i$  is trying to get a number`
  - `int number[i] = 0; // the number got by process  $i$ ;`  
    // “0” means process  $i$  not interested in being served

```
ReleaseCS(int myid) {  
    number[myid] = 0;  
}
```

// a utility function

```
boolean Smaller(int number1, int id1, int number2, int id2) {  
    if (number1 < number 2) return true;  
    if (number1 == number2) {  
        if (id1 < id2) return true; else return false;  
    }  
    if (number 1 > number2) return false;  
}
```

```

RequestCS(int myid) {
    choosing[myid] = true;
    for (int  $j = 0$ ;  $j < n$ ;  $j++$ )
        if (number[ $j$ ] > number[myid]) number[myid] = number[ $j$ ];
    number[myid]++;
    choosing[myid] = false;

    for (int  $j = 0$ ;  $j < n$ ;  $j++$ ) {
        while (choosing[ $j$ ] == true);
        while (number[ $j$ ] != 0 &&
            Smaller(number[ $j$ ],  $j$ , number[myid], myid));
    }
}

```

get a number

wait for people ahead of me

```

choosing[myid] = true;
for (int j = 0; j < n; j++)
    if (number[j] > number[myid])
        number[myid] = number[j];
number[myid]++;
choosing[myid] = false;

for (int j = 0; j < n; j++) {
    while (choosing[j] == true);
    while (number[j] != 0 &&
        Smaller(number[j], j,
            number[myid], myid));
}

```

- **Progress:** Proof by contradiction. Consider any set of processes that wants to enter the CS but no one can make progress. Each process is guaranteed to get a queue #. Let process  $i$  be the one with the smallest queue number. Consider where process  $i$  can be blocked:

- Case 1:

Process  $j$  will eventually set choosing[ $j$ ] to false

Process  $j$  will then block (otherwise there is progress already!)

- Case 2: Impossible since process  $i$  has the smallest queue number

- **No starvation:** Can be similarly shown...work it out yourself...

**Mutual exclusion:** Suppose  $i$  and  $k$  both in critical section.

At T1, process  $i$  is here

```
choosing[myid] = true;
for (int j = 0; j < n; j++)
    if (number[j] > number[myid])
        number[myid] = number[j];
number[myid]++;
choosing[myid] = false;
```

```
for (int j = 0; j < n; j++) {
    while (choosing[j] == true);
    while (number[j] != 0 &&
        Smaller(number[j], j,
        number[myid], myid));
}
```

At T1, process  $k$  is here

- W.l.o.g, assume  $\text{Smaller}(\text{number}[i], i, \text{number}[k], k)$  after they are in the critical sec
- Process  $k$  must see  $\text{number}[i] == 0$  at that time T1: We want to know where process  $i$  is at time T1.
  - Case 1: Process  $i$  has not executed “if ( $\text{number}[k] > \text{number}[i]$ )”. Then eventually  $\text{number}[i] > \text{number}[k]$ . Impossible.
  - Case 2: Has executed “if ()”
    - Subcase 2.1: Process  $i$  has executed “ $\text{number}[\text{myid}]++$ ”; -- impossible since  $\text{number}[i] == 0$
    - Subcase 2.2: Has not executed “ $\text{number}[\text{myid}]++$ ”; -- This is the only possible case.

Case 2: At T2, process *i* is here

=====

At T1, process *i* is here

```
choosing[myid] = true;
for (int j = 0; j < n; j++)
    if (number[j] > number[myid])
        number[myid] = number[j];
number[myid]++;
choosing[myid] = false;
```

At T2, process *k* is here

```
for (int j = 0; j < n; j++) {
    while (choosing[j] == true);
    while (number[j] != 0 &&
           Smaller(number[j], j,
                  number[myid], myid));
}
```

At T1, process *k* is here

- Now continue and consider the time T2 when process *k* invoked “while (choosing[ *i* ] = true);” and passed that statement.
- We want to see where process *i* is at T2. Since choosing[ *i* ] = false, process *i* must either have finished choosing its queue number or have not started choosing:
  - Case 1: process *i* has finished choosing and has executed choosing[ *i* ] = false; Impossible since T2 < T1.
  - Case 2: process *i* has not started choosing and has not executed choosing[ *i* ] = true. But then number[ *i* ] will be larger than number[ *k* ]. Contradiction.

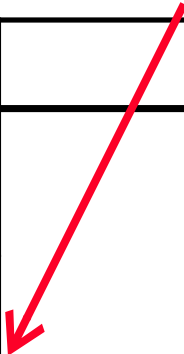


# Roadmap

- Software solutions
  - Unsuccessful attempts
  - Peterson's algorithm
  - Bakery algorithm
- Hardware solutions
  - Disabling interrupts to prevent context switch
  - Special machine-level instructions

# Disable Interrupts

Do not allow context switch here



process 1	process 2
read x into a register (value read: 0)	
increment the register (1)	
	read x into a register (value read: 0)
	increment the register (1)
write value in register back to x (1)	
	write value in register back to x (1)

# Special Machine-level Instructions

```
boolean TestAndSet(Boolean openDoor, boolean newValue) {  
    boolean tmp = openDoor.getValue();  
    openDoor.setValue(newValue);  
    return tmp;  
}
```

Executed  
atomically –  
cannot be  
interrupted

---

```
shared Boolean variable openDoor initialized to true;  
RequestCS(process_id) {  
    while (TestAndSet(openDoor, false) == false) {}  
}  
ReleaseCS(process_id) { openDoor.setValue(true); }
```

# Summary

- Module overview & policy
- Mutual exclusion problem in shared-memory systems
- Software solutions
  - Unsuccessful attempts
  - Peterson's algorithm
  - Bakery algorithm
- Hardware solutions
  - Disabling interrupts to prevent context switch
  - Special machine-level instructions

# Homework Assignment

- Devise a mutual exclusion algorithm for  $n$  processes by using Peterson's 2-mutual-exclusion algorithm **as a black-box**
- Page 28:
  - Problem 2.1 (clearly write out a scenario for 2 processes where problem occur as on slide 22)
  - Problem 2.3 (same as above)
  - Problem 2.4 (either give a proof or construct a problematic scenario as above. Do this for all three properties.)
- Bring you **completed** homework to class next week