

CS4231  
Parallel and Distributed Algorithms

Lecture 2

Instructor: Haifeng YU

## Review of Last Lecture

- Mutual exclusion problem in shared-memory systems
- Software solutions
  - Unsuccessful attempts
  - Peterson's algorithm
  - Bakery algorithm
- Hardware solutions
  - Disabling interrupts to prevent context switch
  - Special machine-level instructions

# Today's Roadmap

- Chapter 3 “Synchronization Primitives”
- Why do we need synchronization primitives?
- Synchronization primitive: Semaphore
- Synchronization primitive: Monitor

# The Busy Wait Problem

- Solutions developed in last lecture has a common busy wait problem
  - Wastes CPU cycles
  - We want to release the CPU to other processes
  - **Need OS support**
- Synchronization primitives
  - OS-level APIs that the program may call
  - Don't worry about how they are implemented
- Semaphores
- Monitors

# Semaphore Semantics

- (Figure 3.1 and 3.2 in the textbook can be confusing, you can ignore that.)

- Internally, each semaphore has

- A boolean *value* – initially true
- A *queue* of blocked processes – initially empty

P():

```
if (value == false) {  
    add myself to queue  
    and block;  
}
```

}

*value* = false;

Executed  
atomically  
(e.g.,  
interrupt  
disabled)

(if blocks, will context switch  
to some other process)

V():

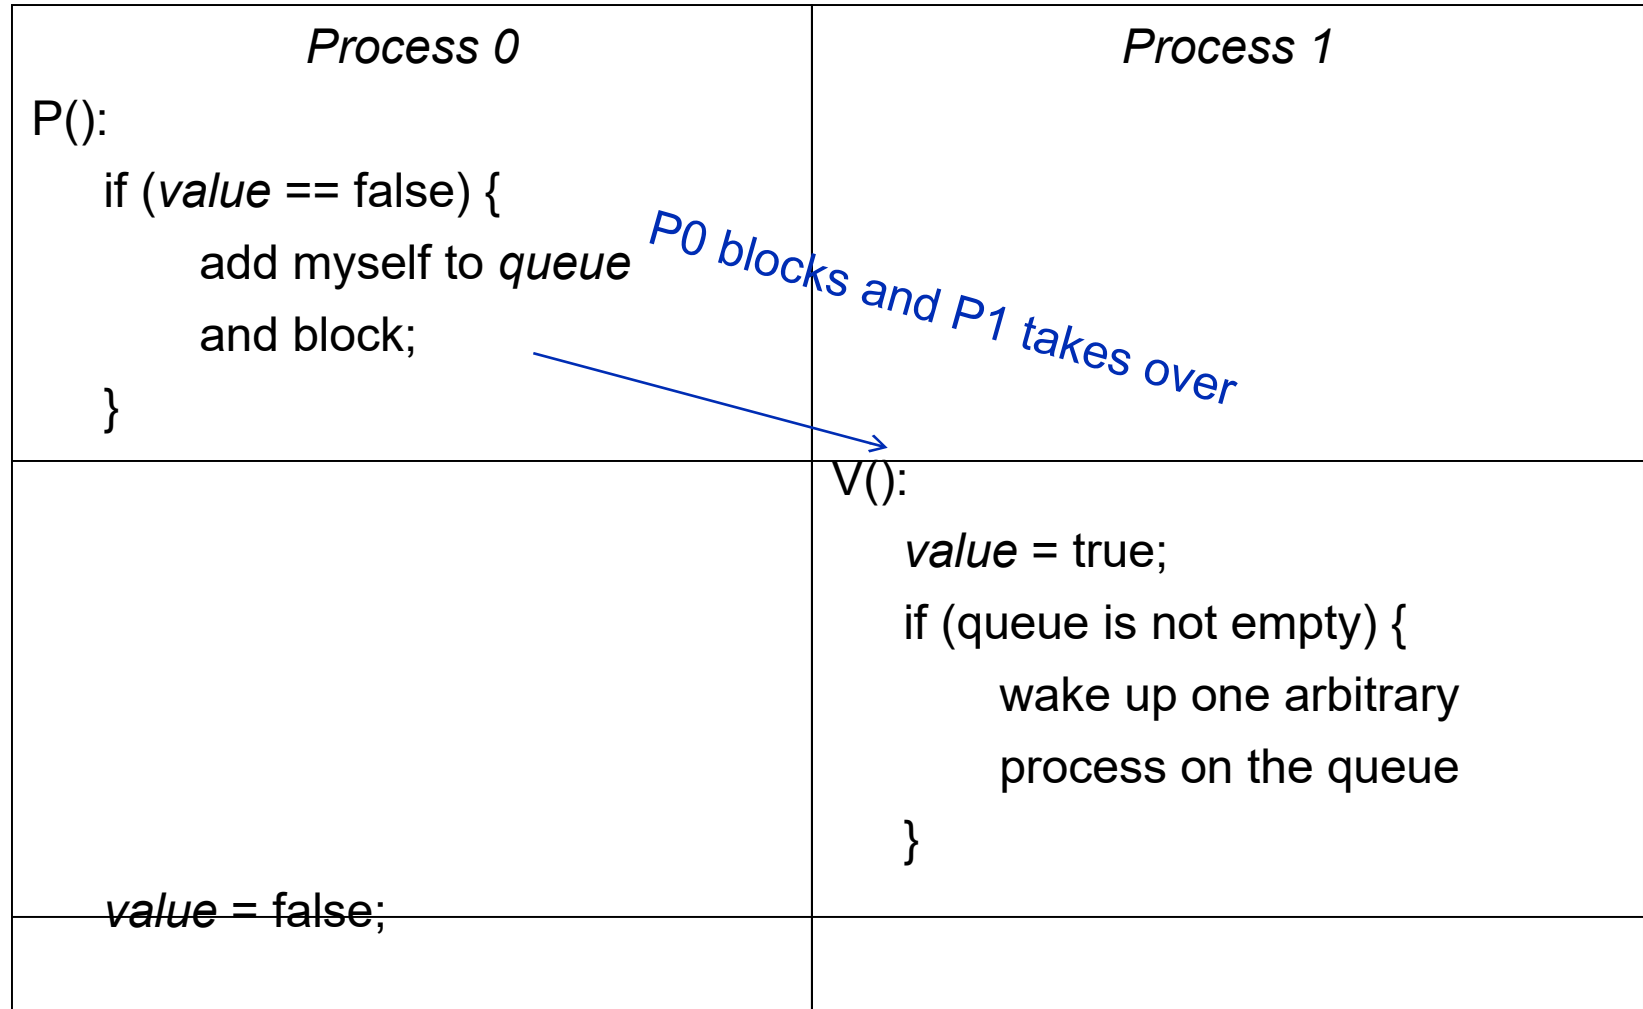
*value* = true;

```
if (queue is not empty) {  
    wake up one arbitrary  
    process on the queue  
}
```

}

Executed  
atomically  
(e.g.,  
interrupt  
disabled)

## Example: P0 invokes P() when value is false, and then P1 invokes V()



# Semaphore Semantics

- Exactly one process is waken up in  $V()$ 
  - The process waken up is **chosen arbitrarily**
  - Some implementations choose the first process on the queue – Should always check the API semantics for the system you are using

# Using Semaphore for Mutual Exclusion

- RequestCS() { P(); }
- ReleaseCS() { V(); }
- The nice thing about this mutual exclusion design is no busy waiting



# Dining Philosopher Problem (Dijkstra'65)

- 5 philosophers,  
5 chopsticks

- Use 5  
semaphores,  
one for each  
chopstick:

- Chopstick[1..5]


*Philosopher i*

```
while (true) {  
    //think for a while, getting hungry  
    chopstick[i].P();  
    chopstick[(i+1) % 5].P();  
    //eat now; (critical section)  
    chopstick[i].V();  
    chopstick[(i+1) % 5].V();  
}
```

# The Danger of Deadlock

<i>philosopher 1</i>	<i>philosopher 2</i>	<i>philosopher 3</i>	<i>philosopher 4</i>	<i>philosopher 5</i>
<code>cstick[1].P();</code> <code>cstick[2].P();</code>	<code>cstick[2].P();</code> <code>cstick[3].P();</code>	<code>cstick[3].P();</code> <code>cstick[4].P();</code>	<code>cstick[4].P();</code> <code>cstick[5].P();</code>	<code>cstick[5].P();</code> <code>cstick[1].P();</code>

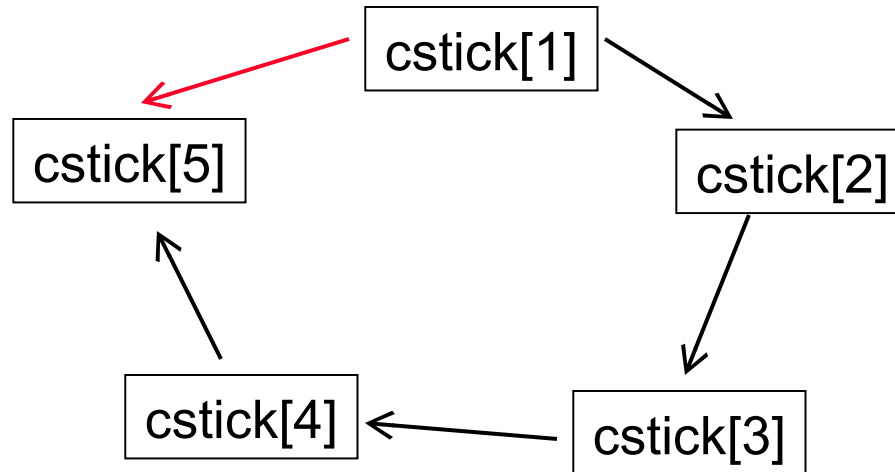
- It is possible for all processes to block
  - **Deadlock !**



# Avoiding Deadlock

<i>philosopher 1</i>	<i>philosopher 2</i>	<i>philosopher 3</i>	<i>philosopher 4</i>	<i>philosopher 5</i>
<code>cstick[1].P();</code> <code>cstick[2].P();</code>	<code>cstick[2].P();</code> <code>cstick[3].P();</code>	<code>cstick[3].P();</code> <code>cstick[4].P();</code>	<code>cstick[4].P();</code> <code>cstick[5].P();</code>	<code>cstick[1].P();</code> <code>cstick[5].P();</code>

- Avoid cycles (or have a total ordering of the chopsticks)



# Today's Roadmap

- Chapter 3 “Synchronization Primitives”
- Why do we need synchronization primitives?
- Synchronization primitive: Semaphore
- Synchronization primitive: Monitor
  - Monitor semantics
  - Using monitors to solve synchronization problems

# Monitors

- Semaphore are quite low-level
  - Monitors are higher-level and easier to use
  - You can always use one to implement the other
- Java only has monitors

# Monitor Semantics

- Monitor object M
- Every object in Java is a monitor

enterMonitor():

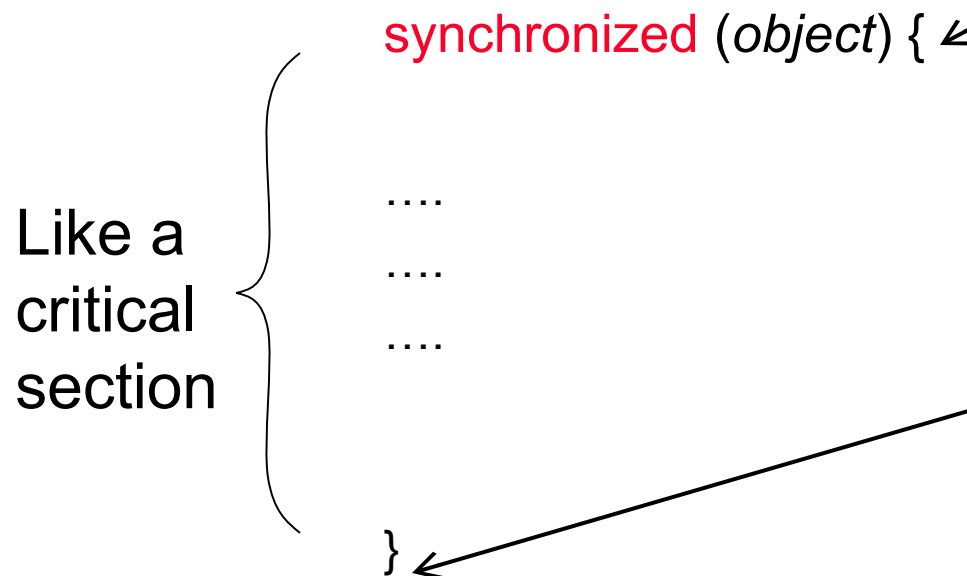
enter monitor if no one is in; otherwise add myself to the **monitor-queue** and block;

exitMonitor():

If **monitor-queue** is not empty, pick one arbitrary process from **monitor-queue** and unblock it

Like a  
critical  
section

**synchronized** (*object*) {  
....  
....  
....  
}



sometime we also say that a monitor has a **monitor lock**

# Monitor Semantics

- Each monitor has two queues of blocked processes
  - **monitor-queue** and **wait-queue**
- Three special methods for using when inside a monitor

**synchronized** (*object*) {

....

*object*.wait();

....

*object*.notify();

....

*object*.notifyAll();

....

}

Add myself to the **wait-queue**,  
**exit monitor**, and then block ---  
all done atomically

If **wait-queue** is not empty, pick  
one arbitrary process from **wait-queue**  
and unblock it

If **wait-queue** is not empty, unblock  
all processes in **wait-queue**

# Two Kinds of Monitors

<i>Process 0</i>	<i>Process 1</i>
synchronized (object) { .... object.wait();	
	synchronized (object) { .... object.notify();
which process should continue execution at this point? }	}

Only one process can be inside the monitor: Two possibilities....



## First kind: Hoare-style Monitor

<i>Process 0</i>	<i>Process 1</i>
synchronized (object) { .... object.wait();	
	synchronized (object) { .... object.notify();
.... }	
	.... }

process 0 takes over the execution

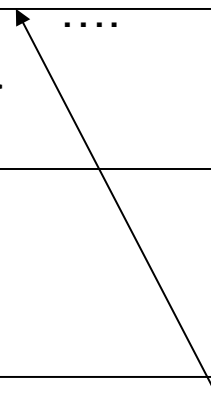
## First kind: Hoare-style Monitor

<i>Process 0</i>	<i>Process 1</i>
synchronized (object) { if (x != 1) object.wait();	
	synchronized (object) { x=1; object.notify();
assert(x == 1); // x must be 1 x = 2; }	
	// x may no longer be 1 here }

process 0 takes over the execution

## Second kind: Java-style Monitor

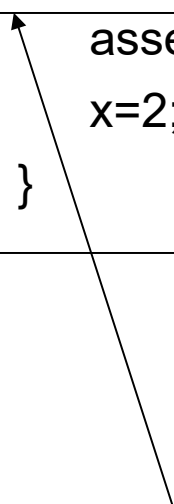
<i>Process 0</i>	<i>Process 1</i>
synchronized (object) { .... object.wait();	
	synchronized (object) { .... object.notify();
	.... }
.... }	



process 1 continues execution

## Second kind: Java-style Monitor

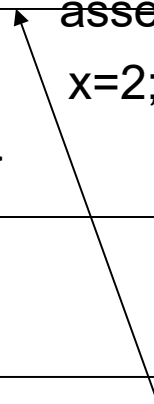
<i>Process 0</i>	<i>Process 1</i>
synchronized (object) { if (x != 1) object.wait();	
	synchronized (object) { x=1; object.notify();
	assert(x == 1); // x must be 1 x=2; }
// needs to acquire monitor lock // x may not be 1 here }	



process 1 continues execution

## Second kind: Java-style Monitor

<i>Process 0</i>	<i>Process 1</i>
synchronized (object) { while (x != 1) object.wait();	
	synchronized (object) { x=1; object.notify();
	assert(x == 1); // x must be 1 x=2; }
assert(x == 1); // if P0 gets here }	



process 1 continues execution

## Two kinds of monitors: More

- Java-style monitor is more popular
  - But you should always check the semantics of the monitor if you are not using Java
- Synchronization code is extremely difficult (if not impossible) to debug
  - There are bugs that people fail to debug after many years
  - Need to get it right the first time

## Other Ways of Using Monitor in Java

```
public synchronized void  
  myMethod () {  
    ....  
    ....  
    ....  
}
```

=

```
public void myMethod () {  
    synchronized (this) {  
        ....  
        ....  
        ....  
    }  
}
```

static methods can also be synchronized –  
the monitor lock is class wide

# Nested Monitor in Java

```
synchronized (ObjA) {  
    synchronized (ObjB) {  
        ObjB.wait();  
    }  
}
```

- Here Java only releases the monitor lock on ObjB and not the monitor lock on ObjA.

- Hence this piece of code will block and will not reach ObjB.notify() – deadlock!

```
synchronized (ObjA) {  
    synchronized (ObjB) {  
        ObjB.notify();  
    }  
}
```

- Different monitor implementations may differ in how nested monitors are implemented – check the spec to tell...



# Today's Roadmap

- Chapter 3 “Synchronization Primitives”
- Why do we need synchronization primitives?
- Synchronization primitive: Semaphore
- Synchronization primitive: Monitor
  - Monitor semantics
  - Using monitors to solve synchronization problems

# The Producer-Consumer Problem

- A circular buffer of size  $n$
- A single producer and a single consumer
- Producer places item to the end of the buffer if
  - Buffer is not full
- Consumer removes item from the head of the buffer if
  - Buffer is not empty
- Example: hard drive as producer and printer as consumer

# Using Monitor to Solve the Producer/Consumer problem

object sharedBuffer;

```
void produce() {  
    synchronized (sharedBuffer) {  
        if (sharedBuffer is full)  
            sharedBuffer.wait();  
        add an item to sharedBuffer;  
        if (sharedBuffer *was* empty)  
            sharedBuffer.notify();  
    }  
}
```

```
void consume() {  
    synchronized (sharedBuffer) {  
        if (sharedBuffer is empty)  
            sharedBuffer.wait();  
        remove item from sharedBuffer;  
        if (sharedBuffer *was* full)  
            sharedBuffer.notify();  
    }  
}
```

# Using Monitor to Solve the Producer/Consumer problem

object sharedBuffer;

```
void produce() {  
    synchronized (sharedBuffer) {  
        if (sharedBuffer is full)  
            sharedBuffer.wait();  
        add an item to sharedBuffer;  
        if (sharedBuffer *was* empty)  
            sharedBuffer.notify();  
    }  
}
```

notification is lost if no  
process is waiting – very  
different semantics from  
V()

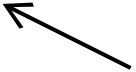
# The Reader-Writer Problem

- Multiple readers and writers are accessing a file
  - A writer must have exclusive access
  - But readers may simultaneously access the file

```
int numReader, numWriter; Object object;
```

```
void writeFile() {  
    synchronized (object) {  
        while (numReader > 0 ||  
            numWriter > 0)  
            object.wait();  
        numWriter = 1;  
    }  
    // write to file;  
    synchronized (object) {  
        numWriter = 0;  
        object.notifyAll();  
    }  
}
```

```
void readFile() {  
    synchronized (object) {  
        while (numWriter > 0)  
            object.wait();  
        numReader++;  
    }  
    // read from file;  
    synchronized (object) {  
        numReader--;  
        object.notify();  
    }  
}
```

you can prove that it must  
be a writer who is notified

# The Starvation Problem

- Writers may get starved if there is a continuous stream of readers
  - There's a way to avoid that....will be your homework...

# Summary

- Why do we need synchronization primitives
  - Busy waiting waste CPU
  - But synchronization primitives need OS support
- Semaphore:
  - Using semaphore to solve dining philosophers problem
  - Avoiding deadlocks
- Monitor:
  - Easier to use than semaphores / more popular
  - Two kinds of monitors
  - Using monitor to solve producer-consumer and reader-writer problem



# Homework Assignment

- Page 51:
  - Problem 3.3, 3.4, 3.6
  - All using monitors and not semaphores
- Homework due a week from today
- Read Chapter 4, Chapter 5.1, 5.2