

CS4231
Parallel and Distributed Algorithms

Lecture 5

Instructor: Haifeng YU

Review of Last Lecture

- Chapter 7 “Models and Clocks”
 - Goal: Define “time” in a distributed system
- Logical clock
 - “happened before” \Rightarrow smaller clock value
- Vector clock
 - “happened before” \Leftrightarrow smaller clock value
- Matrix clock
 - Gives a process knowledge about what other processes know


Today's Roadmap

- Chapter 9 “Global Snapshot”
- What is our goal?
- Formalizing consistent global snapshot
- Protocol for capturing a consistent global snapshot

Motivation


- Goal: Take a picture of the global computation
 - A snapshot of local states on n processes, together all messages on the fly, **taken at exactly the same time**
 - Note: The book uses two terms “global state” and “global snapshot” – we will stick to “global snapshot”
- Useful for debugging
- Useful for backup/check-pointing
- Useful for calculating global predicate
 - E.g., Exactly how much currency do we have in the country (notice that money flows among people constantly)?

Motivation



- If each process have access to completely accurate physical time – then trivial
 - All processes record local state at 1:00:00pm
 - We always get \$200 total regardless when we take the snapshot

Motivation




- But we don't have completely accurate physical time
- user1 and user2 may record states in slightly different time
- How much inaccuracy can we tolerate?
 - Inaccuracy needs to be below minimum message propagation delay
- Hard!

Weaken Our Goal

- Goal: Take a snapshot of the global computation
 - A snapshot of local states on n processes ~~taken at exactly the same time~~
 - A snapshot of local states on n processes such that the global snapshot ~~could have happened sometime in the past~~
- “In the past”
 - We don’t know when it occurred
- “Could have happened” – User cannot tell the difference
 - But in reality, may not have happened
- Somewhat surprisingly, such a weaker goal is still useful (e.g., debugging, backup, computing global predicates)


Review of “Happened-Before” Relation

- A partial order among events (i.e., local computation, send, receive)
 - Process order, send-receive order, transitivity



Consistent Snapshot Examples

- Consistent snapshot:
 - A snapshot of local states on n processes such that the global snapshot **could have happened sometime in the past**
 - The green cut could have happened in the past
 - The red cut could never happen in the past – why?




Formalizing Consistent Snapshot

- **Global snapshot:**

- A set of events such that if e_2 is in the set and e_1 is before e_2 in process order, then e_1 must be in the set
- How does this compare to your intuition?

- **Consistent global snapshot:**

- A global snapshot such that if e_2 is in the set and e_1 is before e_2 in send-receive order, then e_1 must be in the set
- Do we need to include transitive relations as in “Happened-before”?




Existence of Consistent Global Snapshot

- It is a good habit to prove the existence of something newly defined (why?)
 - The proof will also later lead to an important characterization in our case...
- Consider the ordered events e_1, e_2, \dots on any given process
- Theorem: For any positive integer m , there exists a consistent global snapshot S such that


$$\begin{cases} e_i \in S & \text{for } i \leq m \\ e_i \notin S & \text{for } i > m \end{cases}$$

Proof Intuition



- Example: On process2, we want all blue events to be in S , and all red events not to be in S
 - Find all events that happened before **any of** the blue events and include them in S
 - Claim: S must be a consistent global snapshot

Characterization of Consistent Global Snapshot

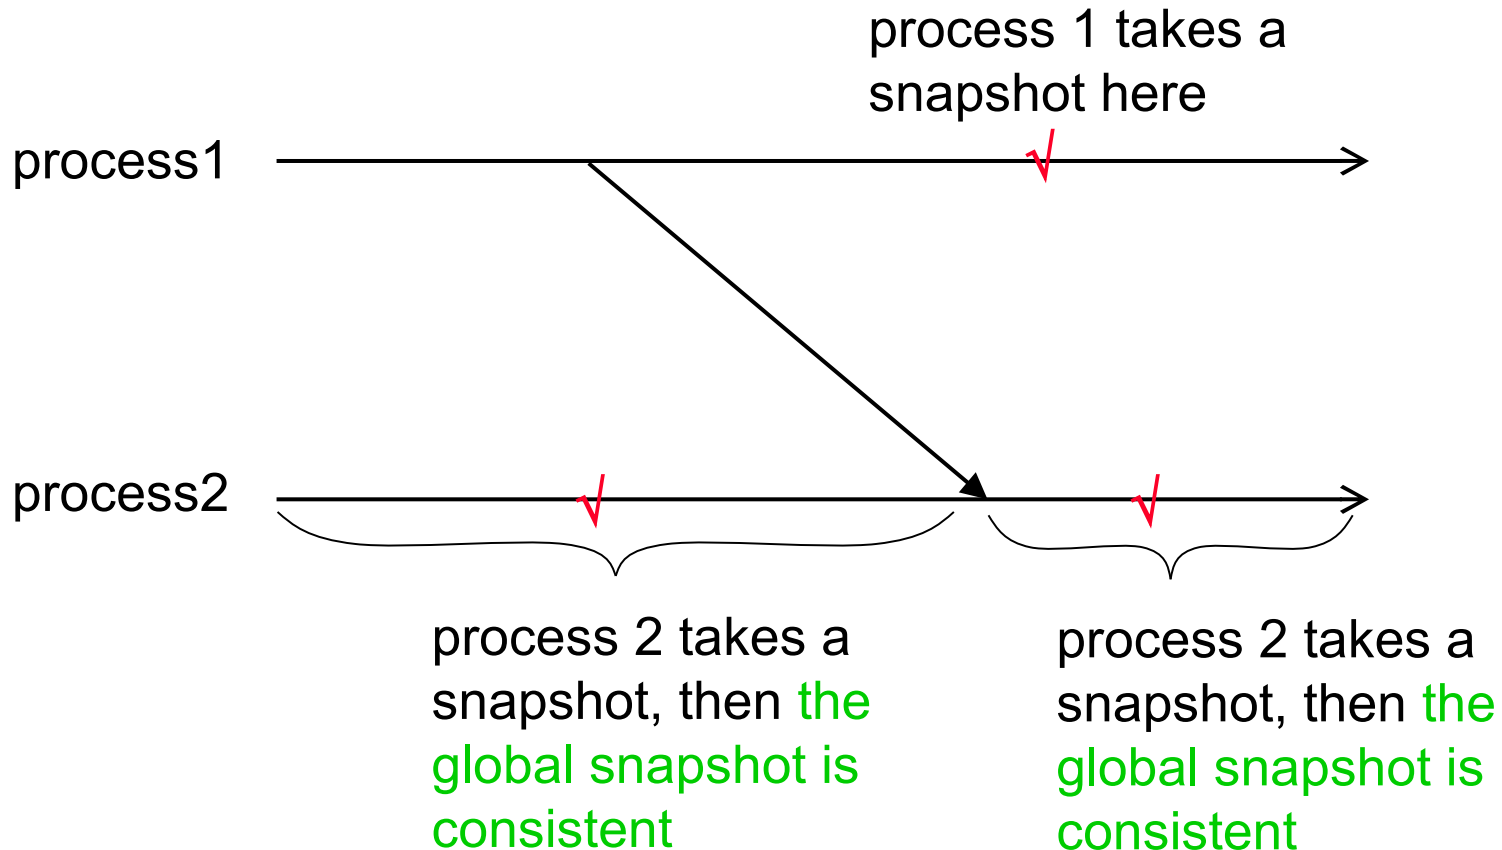


- There are other valid choices for S
- A key characterization: 3 kinds of events
 - Must be in S : Those happened before blue
 - Cannot be S : Those happened after red
 - May or may not be S : Other events

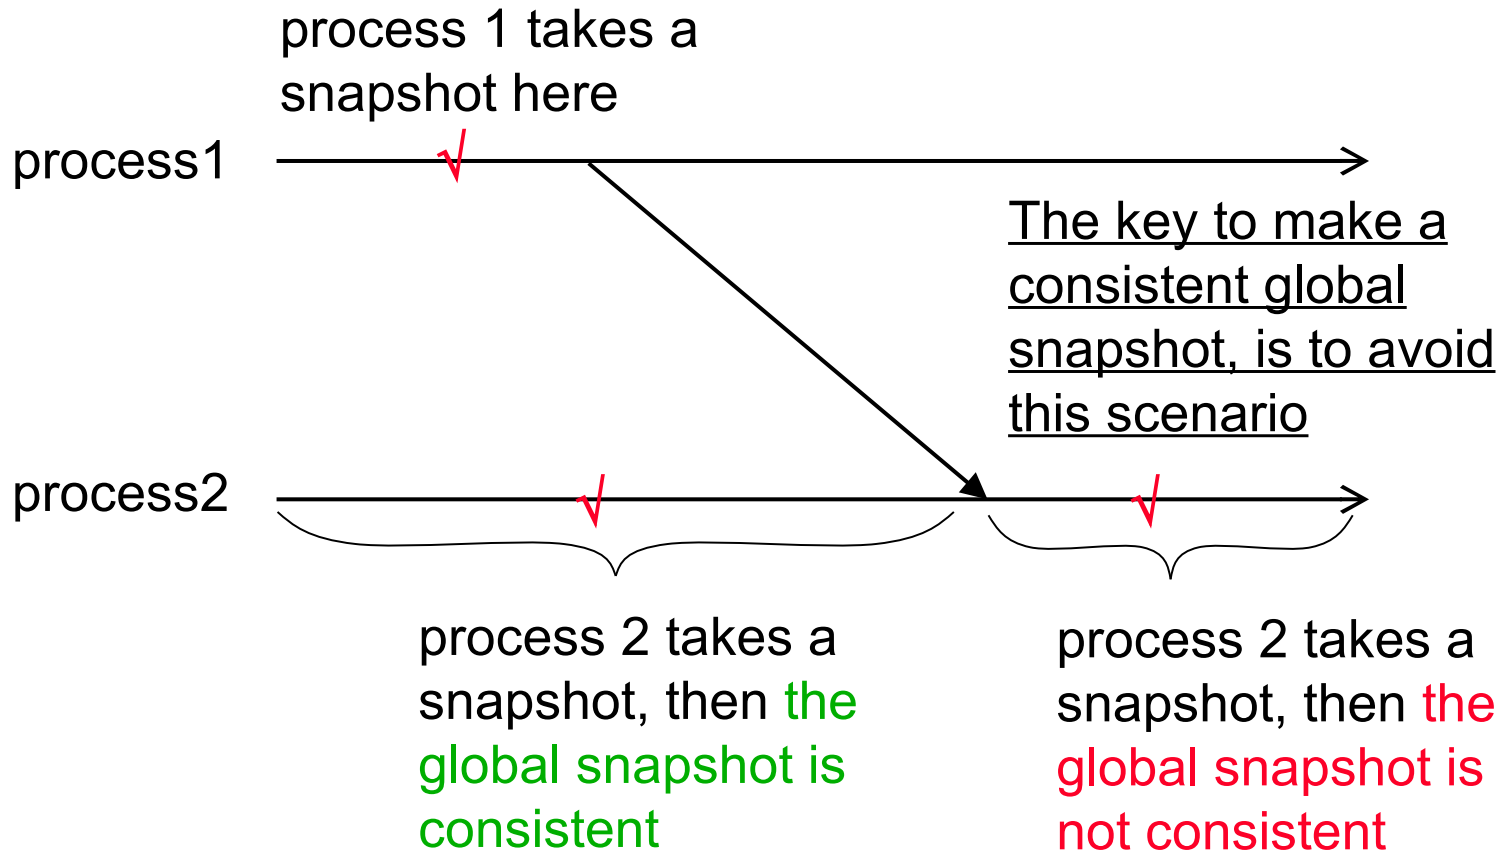
Capturing a Consistent Global Snapshot

- Communication model
 - No message loss
 - Communication channels are unidirectional
 - FIFO delivery on each channel
- Ensuring FIFO:
 - Each process maintains a message number counter for each channel and stamps each message sent
 - Receiver will only deliver messages in order

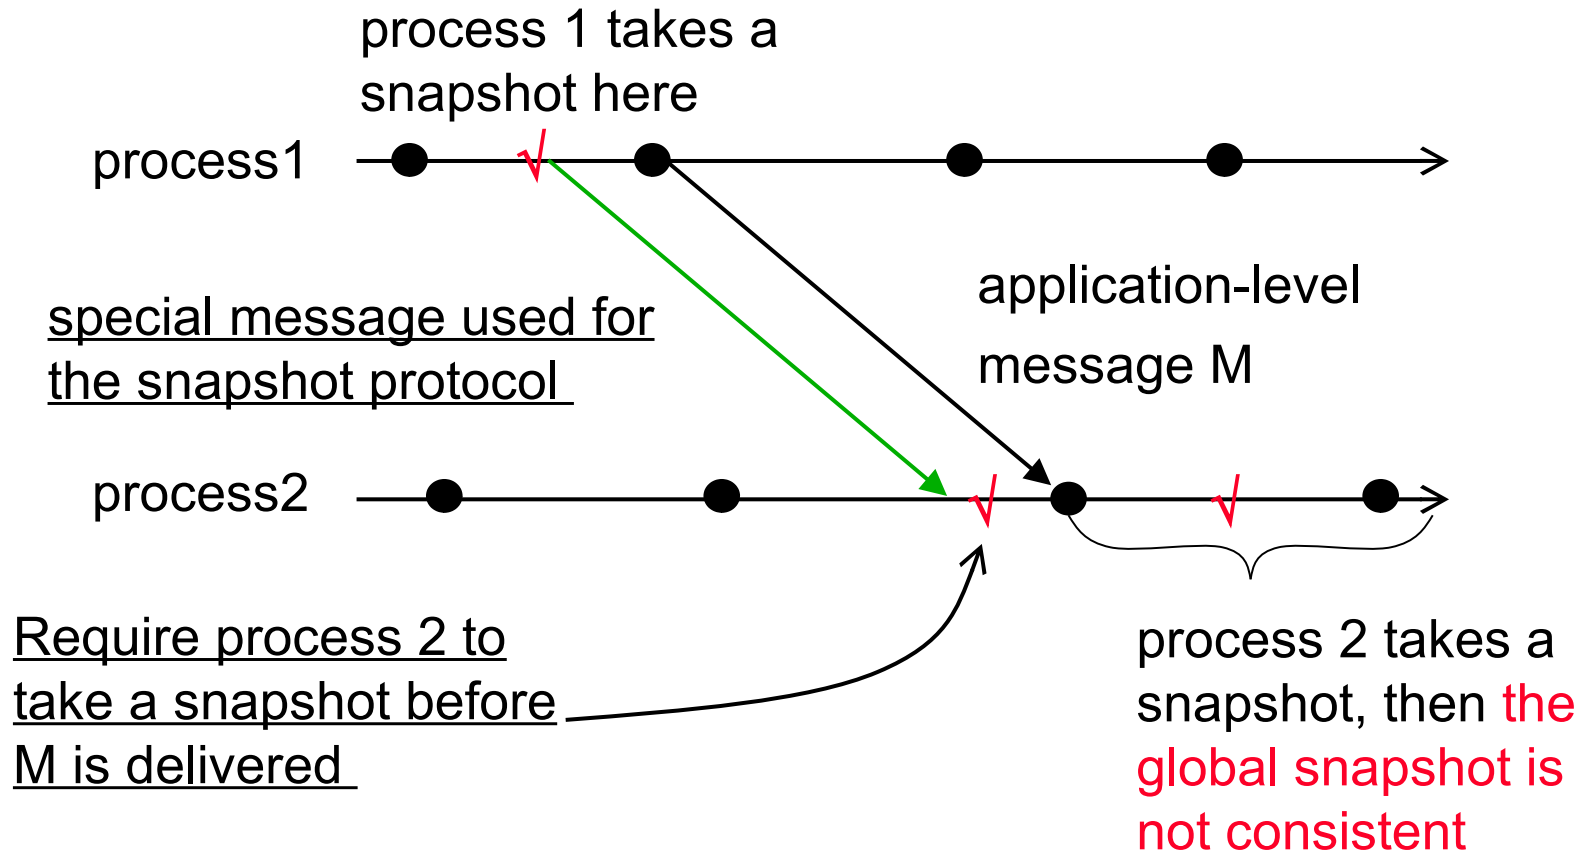
Protocol Intuition



Protocol Intuition



Protocol Intuition




Chandy&Lamport's Protocol:

Taking local snapshots


- Each process is either
 - Red (it has taken the local snapshot), OR
 - White (it has not taken the local snapshot)
- Protocol initiated by a single process by turning itself from White to Red
 - Once a process turns to Red, it immediately send out Marker messages to all other processes
 - Upon receiving Marker, a process turns Red
- Next, we would like also to capture messages that are “on-the-fly” when the snapshot is taken

Characterization of Application Messages


Given a message M, there are 4 possibilities




- M is sent **before** the local snapshot (on the sender) and received **before** the local snapshot (on the receiver): Not “on-the-fly” (why?)



- M is sent **after** the local snapshot and received **after** the local snapshot: Not “on-the-fly”




- M is sent **after** the local snapshot and received **before** the local snapshot: “on-the-fly”




- M is sent **before** the local snapshot and received **after** the local snapshot: “on-the-fly”

Chandy&Lamport's Protocol: Taking snapshots for messages on the fly




(Note: Figure 9.2 on the textbook cannot really happen with FIFO channel – ignore the figure please.)

So the only remaining case is:



- M is sent **before** the local snapshot and received **after** the local snapshot: “**on-the-fly**”

Chandy&Lamport's Protocol: Taking snapshots for messages on the fly



process2 records all messages received in this window – These are the exact set of messages that are only the “fly”

Summary

- Chapter 9 “Global Snapshot”
- What is our goal?
- Formalizing consistent global snapshot
- Protocol for capturing a consistent global snapshot

Homework Assignment

- Page 161

- Problem 9.5

How do you use Lamport's logical clock to compute a consistent global snapshot?

(You should assume that you are given the set of ALL events. You should prove that the global snapshot you get is consistent.)

- Prove the following

- If G and H are both consistent global snapshot, then $G \cap H$ and $G \cup H$ are also consistent global snapshots

- Homework due a week from today

- Read Chapter 12